

GALI_Integrated_Final_AnalysisV8_executed

November 2, 2025

1 GALI 2020 Accelerator Impact Analysis

1.1 A Longitudinal Study of Startup Accelerator Outcomes

Dataset: GALI External Data Release 2020

Period: 2013-2019

Sample: ~23,364 ventures, 605 variables, global coverage

Authors: Fernando Torres, Bernard Precht, Sola Paterson-Marke, Anabelle Li

OIT 249: Data & Decisions, Fall 2025, Stanford Graduate School of Business

1.2 Executive Summary

1.2.1 Research Question

Do ventures that participate in startup accelerator programs outperform comparable applicants? This analysis leverages the GALI 2020 longitudinal dataset to assess whether observed revenue, employment, and funding trajectories are consistent with acceleration-driven gains under a selection-on-observables lens.

1.2.2 Key Findings

- **= 0.05 wins (revenue & debt):** Across 9,567 ventures (2,451 participants), accelerator participation lifts $\Delta \log(\text{revenue})$ by 0.49–0.50 (+64%, $p < 0.001$) and raises FU1 debt incidence by +2.3 pp ($p = 0.02$) after clustering at the programme level.
- **Directional signals to monitor:** Employment gains stay modest—+7.6% ($p = 0.001$) in OLS/IPW and +5.0% ($p = 0.050$) in the strict-caliper PSM—so we classify H2 as directional. Equity incidence edges up +2.1 pp ($p = 0.057$), and regional/gender effects outside South Asia, Latin America & Caribbean, and male/mixed teams retain wide confidence intervals.
- **Robustness & estimator choice:** Removing post-treatment acceptance flags from the propensity model restored overlap: the 0.10 caliper (no replacement) now matches 2,412 treated ventures (98.4%) with a 0.474 log ATT (~+60.7%). Looser calipers and kernel weighting sit between +52% and +56%. Employment remains borderline under these checks, so we keep OLS/IPW as the narrative anchor and cite PSM as confirmation.
- **Log-delta reality check:** Median baseline revenue remains \$2.5k for participants versus \$170 for non-participants, with 40–48% of ventures at \$0. Within the strict PSM sample, participants add a median \$1.7k ($\$2.0k \rightarrow \$10.0k$) while matched controls add ~\$0.3k ($\$2.0k \rightarrow \$7.0k$). We therefore interpret the +60% log effect as an elasticity around revenue-active ventures and use the dollar deltas when advising practitioners.

- **Program & gender dynamics:** Men-only teams post the largest revenue premium (+81%, $p < 0.001$) and mixed teams follow closely (+73%, $p = 0.002$); investor guarantees (+6.0 pp, $p = 0.015$) and sector focus (+9.7 pp, $p < 0.001$) correlate with higher FU1 equity conversion, whereas demo-day intensity remains negative and insignificant (-27%, $p = 0.26$).

1.2.3 Sample & Attrition Snapshot

- **Full dataset:** 23,364 ventures, median baseline revenue = \$0, baseline equity coverage = 16.9%, FU1 response = 41%.
- **Analytical sample:** 9,567 ventures (2,451 participants / 7,116 non-participants), median baseline revenue rises to \$500, baseline equity coverage slips to 15.2%, baseline debt coverage edges up to 12.5%.
- **Response gaps:** Participants respond to FU1 at 61% versus 37% for non-participants; Sub-Saharan Africa keeps the strongest non-participant response (46%), while Other / multi-region accelerators see drop-offs near 27%, especially among women-only teams.

1.2.4 Methods Summary

- **Estimation approach:** OLS with year, region, and sector fixed effects and program-clustered standard errors, complemented by propensity-score matching and inverse-probability-weighted (IPW) robustness checks (clip = 0.05).
- **Sample:** 9,567 ventures with complete baseline and FU1 outcomes (2,451 participants; 7,116 non-participants).
- **Diagnostics:** Balance tables, propensity-score overlap, residual and influence checks, variance inflation factors, winsorization sensitivity, and IPW leverage inspections.

1.2.5 Recommendations

1. **For accelerators:** Anchor coaching and service design on revenue traction and investor readiness; treat headcount targets as directional until matched evidence improves.
2. **For policymakers:** Focus ecosystem investments where evidence shows the largest marginal gains (South Asia, Latin America & Caribbean) while prioritising data collection that can shrink wide intervals in under-represented regions.
3. **For entrepreneurs:** Accelerator participation delivers robust revenue gains and debt access, with employment and equity improvements still directional—pair programme participation with capital-readiness support to convert these gains into statistically precise funding outcomes.

1.2.6 Limitations

- **Selection bias:** Even with rich covariates, unobserved founder quality may bias estimates; results should be read as *consistent with* causal effects under selection-on-observables.
 - **Attrition:** FU1 response hovers around 41%; IPW checks mitigate but cannot eliminate non-response concerns, particularly in Other / multi-region cohorts.
 - **Generalisability:** Data tilt toward ventures that apply to established accelerators, so findings may not extend to nascent programmes or non-applicants.
-

1.3 Setup: Configuration & Imports

```
[1]: # Core libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
import warnings
from IPython.display import display

try:
    from sklearn.metrics import roc_auc_score
except ImportError:
    from scipy import stats as _stats
    def roc_auc_score(y_true, y_score):
        y_true = np.asarray(y_true)
        y_score = np.asarray(y_score)
        if len(y_true) == 0:
            return np.nan
        pos_mask = y_true == 1
        neg_mask = y_true == 0
        n_pos = pos_mask.sum()
        n_neg = neg_mask.sum()
        if n_pos == 0 or n_neg == 0:
            return np.nan
        ranks = _stats.rankdata(y_score)
        rank_sum = ranks[pos_mask].sum()
        return (rank_sum - n_pos * (n_pos + 1) / 2) / (n_pos * n_neg)

# Statistical modeling
import statsmodels.api as sm
import statsmodels.formula.api as smf
from statsmodels.stats.outliers_influence import variance_inflation_factor
from statsmodels.stats.diagnostic import het_breuschpagan

# Display settings
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', 100)
pd.set_option('display.max_rows', 100)
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette('colorblind')

# Reproducibility
np.random.seed(42)
hypothesis_tests = {} # store Wald/joint-test diagnostics for appendix
```

```
print(" Libraries loaded successfully")
```

```
Libraries loaded successfully
```

1.4 Section 0: Data Loading and Exploratory Analysis

1.4.1 0.1 Load the GALI 2020 Dataset

The GALI 2020 external release covers **23,364** accelerator applications spanning **605** variables across 160+ countries (2013–2019 cohorts). Columns fall into distinct families: venture identifiers, selection outcomes, baseline financial/investment history, founder attributes, impact focus, program design, and longitudinal follow-ups (FU1–FU4). Each observation is a venture-program application with potential re-surveys embedded as wide columns.

For this study we focus on baseline covariates and the first follow-up (FU1) where response rates remain viable (~40–50%). Higher-order follow-ups exhibit >80% missingness, so they are excluded from causal estimation. Variables without cross-venture coverage (e.g., detailed founder biographies, sparse investment sub-categories) are left in the raw data but omitted from modelling.

Variables used in modelling

- **Baseline scale & maturity:** log_revenue_m1, log_ft_employees_m1, years_since_founding.
- **Capability & mission proxies:** digital_score (count of four digital touchpoints; range 0–4), impact_intensity, has_ip, and sector classifications (info_sector).
- **Program design levers:** program_demo_day_yes, program_curric_struct_yes, program_ben_ginv_yes, program_sector_focus_yes.
- **Outcomes & treatment:** participated, delta_log_revenue, delta_log_ft, fu1inv_hasequity, alongside survey-response indicators used for attrition adjustments.

Variable families excluded from modelling

- **Founder mini-biographies and equity splits (52 columns):** >80% missing across cohorts; retained in the raw extract but excluded from all statistics.
- **FU2–FU4 longitudinal panels (303 columns):** FU2 response already exceeds 80% missing, so later waves are documented only in the appendix.
- **Granular investment breakdowns (60+ columns):** Sparse sub-categories for individual lenders or philanthropies are collapsed into FU1 incidence flags for equity, debt, and philanthropy.
- **Narrative programme indicators (dozens of free-text or single-program flags):** Useful for context but too idiosyncratic for cross-program inference; cited qualitatively where relevant.

We keep these sparse fields available in the notebook for reproducibility—each retained in the raw `df`—but they never enter model matrices or summary tables unless explicitly noted in the appendices.

Scope: This subsection loads and provides an initial overview of the GALI 2020 dataset, examining its structure, dimensions, and basic characteristics.

Approach: We use pandas to load the primary dataset along with its data dictionary and notes files, then inspect the shape and preview the first few rows to understand the data organization.

Why it matters: Understanding the dataset structure is foundational to our analysis. The GALI 2020 dataset contains the longitudinal venture data necessary to test our core hypothesis about whether accelerator participation causally improves venture outcomes. This initial exploration helps us identify the breadth of variables available (605 variables across ~23,364 ventures) and confirms we have the necessary treatment indicators, baseline characteristics, and follow-up outcomes to conduct rigorous causal inference.

```
[2]: # Load main data
df = pd.read_csv('data/GALI_External_DataRelease_2020_data.csv', ↴
    low_memory=False)
data_dict = pd.read_csv('data/GALI_External_DataRelease_2020_data_dictionary. ↴
    csv')
notes = pd.read_csv('data/GALI_External_DataRelease_2020_notes.csv')

print(f"Dataset Shape: {df.shape[0]}:,{df.shape[1]}:,{} ventures × {df.shape[1]}:,{} variables")
print(f"\nData Dictionary: {data_dict.shape[0]} variable definitions")
print(f"\nFirst 3 rows preview:")
df.head(3)
```

Dataset Shape: 23,364 ventures × 605 variables

Data Dictionary: 605 variable definitions

First 3 rows preview:

```
[2]: program_id  program_year  program_region  program_duration \
0 P_aG1admzC      2013          Other       3 - 6 months
1 P_aG1admzC      2013          Other       3 - 6 months
2 P_aG1admzC      2013          Other       3 - 6 months

program_sector_focus_yes  program_impact_area_yes \
0                      0.0                  1.0
1                      0.0                  1.0
2                      0.0                  1.0

program_curric_struct_yes  program_ben_ginv_yes  program_demo_day_yes \
0                      0.0                  1.0                  1.0
1                      0.0                  1.0                  1.0
2                      0.0                  1.0                  1.0

New_External_ID  application_year  participated  accepted_initial \
0        ID-13-679      2013            0             NaN
1        ID-13-698      2013            0             NaN
```

2	ID-13-845	2013	0	Nan	
0	accepted_final	info_venture_country	info_venture_country_hq	\	
0	Nan	Netherlands	Netherlands		
1	Nan	Lesotho	Lesotho		
2	Nan	Netherlands	Netherlands		
0	info_has_website	info_has_facebook	info_has_twitter	info_has_linkedin	\
0	1	0	0	0	
1	1	0	0	0	
2	1	0	0	0	
0	info_founding_year	info_legal_status	info_sector	\	
0	2010.0	For-profit company	Artisanal		
1	2011.0	For-profit company	Energy		
2	2013.0	Undecided	Energy		
0	info_financial_goals	info_has_targetmargin	info_target_margin	\	
0	Cover costs and earn some profit	1	16% - 20%		
1	Cover costs and earn some profit	1	11% - 15%		
2	Cover costs and earn some profit	1	11% - 15%		
0	info_has_socialmotives	model_prodmanuf	model_procpack	\	
0	1	1	0		
1	1	1	1		
2	1	0	0		
0	model_distribution	model_wholretail	model_services	model_finserv	\
0	0	1	0	0	
1	1	1	1	0	
2	0	0	1	0	
0	model_unsure	model_invention_based	model_has_patents	\	
0	0	0	0		
1	0	1	1		
2	0	0	0		
0	model_has_copyrights	model_has_trademarks	fins_revenues_m1	\	
0	0	0	135000.0		
1	1	1	0.0		
2	1	0	400000.0		
0	fins_revenues_sincefound		fins_profit_m1	\	
0	0.0		6% - 10%		
1	0.0	Negative ROI (venture lost money in year prior)			
2	0.0		11% - 15%		

```

fins_ft_employees_m1  fins_pt_employees_m1  fins_wages_m1  \
0                      0                      0                  0.0
1                     50                     0                80000.0
2                      1                     2                  0.0

fins_seasonals_m1   fins_volunteers_m1   inv_ownmoney_m1  \
0                      0                      0                6750.0
1                      0                      0              100000.0
2                      0                      0                50000.0

inv_ownmoney_sincefound  inv_hasequity  inv_hasdebt ... \
0                 13000.0          0          0 ...
1                150000.0         0          0 ...
2                150000.0         0          0 ...

fu4inv_philanfrom_govt  fu4inv_philanfrom_nonprofits  \
0                      NaN                     NaN
1                      NaN                     NaN
2                      NaN                     NaN

fu4inv_philanfrom_accelerators  fu4inv_philanfrom_friendsfamily  \
0                      NaN                     NaN
1                      NaN                     NaN
2                      NaN                     NaN

fu4inv_philanfrom_bpcs  fu4inv_philanfrom_crowd  \
0                      NaN                     NaN
1                      NaN                     NaN
2                      NaN                     NaN

fu4inv_philanfrom_nonowneremp  fu4inv_philanfrom_otherinds  \
0                      NaN                     NaN
1                      NaN                     NaN
2                      NaN                     NaN

fu4inv_philanfrom_anothersource  fu4inv_philan_m1  \
0                      NaN                     NaN
1                      NaN                     NaN
2                      NaN                     NaN

fu4inv_plans_outequity_12months  fu4inv_plans_debt_12months  \
0                      NaN                     NaN
1                      NaN                     NaN
2                      NaN                     NaN

fu4inv_plans_philan_12months  fu4inv_plans_outequity_3years  \
0                      NaN                     NaN

```

1		NaN		NaN
2		NaN		NaN
	fu4inv_plans_debt_3years	fu4inv_plans_philan_3years	fu4impact_area_water	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN
	fu4impact_area_educ	fu4impact_area_energy	fu4impact_area_finsrv	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN
	fu4impact_area_information	fu4impact_area_housing	fu4impact_area_agprod	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN
	fu4impact_area_biodiv	fu4impact_area_capacity	fu4impact_area_commdev	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN
	fu4impact_area_conflres	fu4impact_area_disease	fu4impact_area_employ	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN
	fu4impact_area_fueleff	fu4impact_area_equality	fu4impact_area_food	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN
	fu4impact_area_charity	fu4impact_area_health	fu4impact_area_humrights	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN
	fu4impact_area_incgrowth	fu4impact_area_natresources		\
0		NaN		NaN
1		NaN		NaN
2		NaN		NaN
	fu4impact_area_natbio	fu4impact_area_wastemgt	fu4impact_area_efficiency	\
0		NaN	NaN	NaN
1		NaN	NaN	NaN
2		NaN	NaN	NaN

```

        fu4impact_area_sustenergy  fu4impact_area_landuse  \
0                  NaN          NaN
1                  NaN          NaN
2                  NaN          NaN

        fu4impact_area_highimpact  fu4impact_area_watresmgmt  \
0                  NaN          NaN
1                  NaN          NaN
2                  NaN          NaN

        fu4impact_area_womengirls  fu4impact_area_other  fu4impact_use_iris  \
0                  NaN          NaN          NaN
1                  NaN          NaN          NaN
2                  NaN          NaN          NaN

        fu4impact_use_blab_giirs  fu4impact_use_othermeasure  \
0                  NaN          NaN
1                  NaN          NaN
2                  NaN          NaN

        fu4report_any_prior_accelerator
0                  NaN
1                  NaN
2                  NaN

[3 rows x 605 columns]

```

Table 0.1a: Variable Families Retained vs Discarded The inventory below links raw source fields to the engineered features used downstream, reports coverage after loading the raw release, and documents which families were excluded (and why) before constructing the analytical sample.

```
[3]: # Build variable family inventory with coverage and justification
impact_cols = [c for c in df.columns if c.startswith('impact_area_')]
founder_cols = [
    c for c in df.columns
    if c.startswith('found_name') and ('education' in c or 'prior_fp' in c)
]
fu_follow_cols = [
    c for c in df.columns
    if c.startswith('fu2') or c.startswith('fu3') or c.startswith('fu4')
]
fu_follow_example = [c for c in fu_follow_cols if c.endswith('_m1')][3]
if len(fu_follow_example) < 3:
    fu_follow_example = fu_follow_cols[:3]
```

```

digital_cols = [c for c in ['info_has_website', 'info_has_linkedin',  

    ↵'info_has_twitter', 'info_has_facebook'] if c in df.columns]
ip_cols = [c for c in ['model_has_patents', 'model_has_copyrights',  

    ↵'model_has_trademarks'] if c in df.columns]
baseline_finance_cols = [c for c in ['inv_hasequity', 'inv_hasdebt',  

    ↵'inv_hasphilan', 'inv_totaldebt_m1', 'inv_outequity_m1'] if c in df.columns]
ful1_outcome_cols = [c for c in ['fulfins_revenues_m1',  

    ↵'fulfins_ft_employees_m1', 'fulinv_totaldebt_m1', 'fulinv_outequity_m1'] if  

    ↵c in df.columns]

family_specs = [  

    {  

        'family': 'Baseline scale & maturity',  

        'status': 'Retained',  

        'keep_cols': [c for c in ['fins_revenues_m1', 'fins_ft_employees_m1',  

            ↵'info_founding_year'] if c in df.columns],  

        'drop_cols': [],  

        'representative': ['fins_revenues_m1', 'fins_ft_employees_m1',  

            ↵'info_founding_year'],  

        'engineered': ['log_revenue_m1', 'log_ft_employees_m1',  

            ↵'delta_log_revenue', 'delta_log_ft', 'years_since_founding'],  

        'rationale': 'High coverage (>95%) and feeds the log/delta outcomes  

            ↵used in the ATT models.'  

    },  

    {  

        'family': 'Capability & mission proxies',  

        'status': 'Retained',  

        'keep_cols': digital_cols + impact_cols + ip_cols,  

        'drop_cols': [],  

        'representative': digital_cols[:3] + impact_cols[:2] + ip_cols[:1],  

        'engineered': ['digital_score', 'impact_intensity', 'has_ip'],  

        'rationale': 'Combines digital presence, impact scope, and IP coverage;  

            ↵forms composite scores with 70-95% completion.'  

    },  

    {  

        'family': 'Program design levers',  

        'status': 'Retained',  

        'keep_cols': [c for c in ['program_demo_day_yes',  

            ↵'program_curric_struct_yes', 'program_ben_ginv_yes',  

            ↵'program_sector_focus_yes', 'program_impact_area_yes'] if c in df.columns],  

        'drop_cols': [],  

        'representative': ['program_demo_day_yes', 'program_ben_ginv_yes',  

            ↵'program_sector_focus_yes'],  

        'engineered': [],  

        'rationale': 'Binary design levers retained to interpret heterogeneity;  

            ↵99% coverage across programs.'
```

```

},
{
    'family': 'Baseline financing mix',
    'status': 'Retained',
    'keep_cols': baseline_finance_cols,
    'drop_cols': [],
    'representative': baseline_finance_cols[:3],
    'engineered': ['log_totaldebt_m1'],
    'rationale': 'Essential for pre-program capital controls with\u202a
    ↵manageable missingness (60-90% coverage).'
},
{
    'family': 'FU1 outcomes & funding',
    'status': 'Retained',
    'keep_cols': fu1_outcome_cols,
    'drop_cols': [],
    'representative': fu1_outcome_cols[:3],
    'engineered': ['log_revenue_fu1', 'log_ft_employees_fu1', \u202a
    ↵'fu1inv_hasequity', 'fu1inv_hasdebt', 'fu1inv_hasphilan'],
    'rationale': 'Core dependent variables; coverage limited by FU1\u202a
    ↵response (~41%), aligning with the analytical sample.'
},
{
    'family': 'Founder biographies (multi-founder fields)',
    'status': 'Excluded',
    'keep_cols': [],
    'drop_cols': founder_cols,
    'representative': founder_cols[:3],
    'engineered': [],
    'rationale': 'Second/third-founder education and prior experience\u202a
    ↵fields are >80% missing, distorting sample balance.'
},
{
    'family': 'FU2-FU4 follow-ups',
    'status': 'Excluded',
    'keep_cols': [],
    'drop_cols': fu_follow_cols,
    'representative': fu_follow_example,
    'engineered': [],
    'rationale': 'Completion falls below 20% past FU1; retaining them would\u202a
    ↵cut treated counts by >70%.'
}
]

records = []
retained_columns = set()
for spec in family_specs:

```

```

keep_cols = spec['keep_cols']
drop_cols = spec['drop_cols']
coverage_cols = keep_cols if keep_cols else drop_cols
if coverage_cols:
    coverage = df[coverage_cols].notna().mean().mean() * 100
else:
    coverage = float('nan')
if spec['status'] == 'Retained':
    retained_columns.update(keep_cols)
if keep_cols:
    representative_fields = keep_cols[:3]
    if len(representative_fields) < 3:
        representative_fields = spec['representative'][:3]
else:
    representative_fields = [col for col in spec['representative'] if col in df.columns][:3]
rep_display = ', '.join(representative_fields) + (' ...' if len(keep_cols) > 3 else '')
if not rep_display:
    rep_display = '-'
engine_display = ', '.join(spec['engineered']) if spec['engineered'] else '-'

records.append({
    'Family': spec['family'],
    'Status': spec['status'],
    'Representative raw fields': rep_display,
    'Engineered outputs': engine_display,
    'Columns kept (#)': len(keep_cols),
    'Columns dropped (#)': len(drop_cols),
    'Coverage % (raw)': round(coverage, 1) if coverage == coverage else float('nan'),
    'Major rationale': spec['rationale']
})

variable_family_table = pd.DataFrame(records)
variable_family_table = variable_family_table[['Family', 'Status',
    'Representative raw fields', 'Engineered outputs', 'Columns kept (#)',
    'Columns dropped (#)', 'Coverage % (raw)', 'Major rationale']]
display(variable_family_table)

print(f"Unique retained raw columns used downstream: {len(retained_columns)}")

```

	Family	Status	\
0	Baseline scale & maturity	Retained	
1	Capability & mission proxies	Retained	
2	Program design levers	Retained	
3	Baseline financing mix	Retained	

4	FU1 outcomes & funding	Retained
5	Founder biographies (multi-founder fields)	Excluded
6	FU2-FU4 follow-ups	Excluded

	Representative raw fields	\
0	fins_revenues_m1, fins_ft_employees_m1, info_f...	
1	info_has_website, info_has_linkedin, info_has_...	
2	program_demo_day_yes, program_curric_struct_ye...	
3	inv_hasequity, inv_hasdebt, inv_hasphilan ...	
4	fulfins_revenues_m1, fulfins_ft_employees_m1, ...	
5	found_name1_education, found_name2_education, ...	
6	fu2fins_revenues_m1, fu2fins_profit_m1, fu2fin...	

	Engineered outputs	Columns kept (#)	\
0	log_revenue_m1, log_ft_employees_m1, delta_log...	3	
1	digital_score, impact_intensity, has_ip	37	
2		-	5
3	log_totaldebt_m1	5	
4	log_revenue_fu1, log_ft_employees_fu1, fu1inv_...	4	
5		-	0
6		-	0

	Columns dropped (#)	Coverage % (raw)	\
0	0	99.3	
1	0	100.0	
2	0	87.5	
3	0	100.0	
4	0	40.9	
5	6	83.8	
6	303	11.6	

	Major rationale
0	High coverage (>95%) and feeds the log/delta o...
1	Combines digital presence, impact scope, and I...
2	Binary design levers retained to interpret het...
3	Essential for pre-program capital controls wit...
4	Core dependent variables; coverage limited by ...
5	Second/third-founder education and prior exper...
6	Completion falls below 20% past FU1; retaining...

Unique retained raw columns used downstream: 54

Section 0.1 Results: The load inventory cements what feeds the analysis: we keep 54 raw columns across five high-coverage families—baseline scale & maturity (3 fields), capability & mission proxies (37 fields spanning 4 digital presence signals, 30 impact-area flags, and 3 IP indicators), programme design levers (5 binary flags), baseline financing mix (5 capital-share fields), and FU1 outcomes (4 dependent variables). Everything else is filtered out up front: multi-founder biographies (52 columns with >80% missingness), later follow-ups (303 columns with <20% completion), and granular investment sub-categories (collapsed into FU1 incidence flags) remain in the raw frame for

reference but do not enter modelling. Engineered features therefore draw on clearly enumerated inputs: `digital_score` counts four web and social signals (range 0–4), `impact_intensity` tallies the 30 retained impact toggles, and the FU1 capital indicators collapse dozens of detailed sources into three incidence flags. Because only 9,567 ventures (41% of applicants) respond at FU1, the availability of those four outcome columns effectively defines the analytical sample.

Narrative connection: Laying out these families makes the trade-offs explicit. Retaining high-coverage baseline, capability, and design levers gives us rich controls to support selection-on-observables estimators, while pruning sparsely populated founder biographies and FU2–FU4 follow-ups prevents the analytical sample from collapsing. Explicitly tying each engineered feature to the number of raw inputs (e.g., four digital presence signals, thirty impact toggles) clarifies how much information survives preprocessing and why FU1 reporters anchor every downstream statistic.

Scope: This subsection examines missingness patterns across different variable families (baseline financials, founder demographics, program features, follow-up surveys) to understand data quality and inform our analytical strategy.

Approach: We group variables into logical families and calculate the percentage of missing values for each family using pandas aggregation functions. We visualize these patterns to identify systematic gaps.

Why it matters: Missing data patterns directly impact our ability to make causal inferences. If missingness correlates with treatment status or outcomes, our estimates could be biased. Understanding which variable families are complete (baseline financials) versus sparse (later follow-ups, founder details) helps us determine which covariates to include in models and where attrition bias might threaten validity. This diagnostic is essential for establishing the credibility of our causal claims about accelerator impact.

```
[4]: # Define variable families
var_families = {
    'ID & Timing': ['New_External_ID', 'program_id', 'application_year', ↴
    'program_year'],
    'Selection': ['accepted_initial', 'accepted_final', 'participated'],
    'Baseline Financials': [c for c in df.columns if c.startswith('fins_') and ↴
    not c.startswith('fu'))],
    'Baseline Investment': [c for c in df.columns if c.startswith('inv_') and ↴
    not c.startswith('fu'))],
    'Venture Info': [c for c in df.columns if c.startswith('info_')],
    'Founders': [c for c in df.columns if c.startswith('found_')],
    'Business Model': [c for c in df.columns if c.startswith('model_')],
    'Program Features': [c for c in df.columns if c.startswith('program_')],
    'Follow-up 1': [c for c in df.columns if c.startswith('fu1')], ↴
    'Follow-up 2': [c for c in df.columns if c.startswith('fu2')], ↴
    'Follow-up 3': [c for c in df.columns if c.startswith('fu3')], ↴
    'Follow-up 4': [c for c in df.columns if c.startswith('fu4')]} ↴

# Compute missingness
```

```

missing_summary = pd.DataFrame([
    {
        'Variable Family': family,
        'Num Vars': len(cols),
        'Avg Missing %': df[cols].isnull().mean().mean() * 100,
        'Max Missing %': df[cols].isnull().mean().max() * 100
    }
    for family, cols in var_families.items() if cols
])

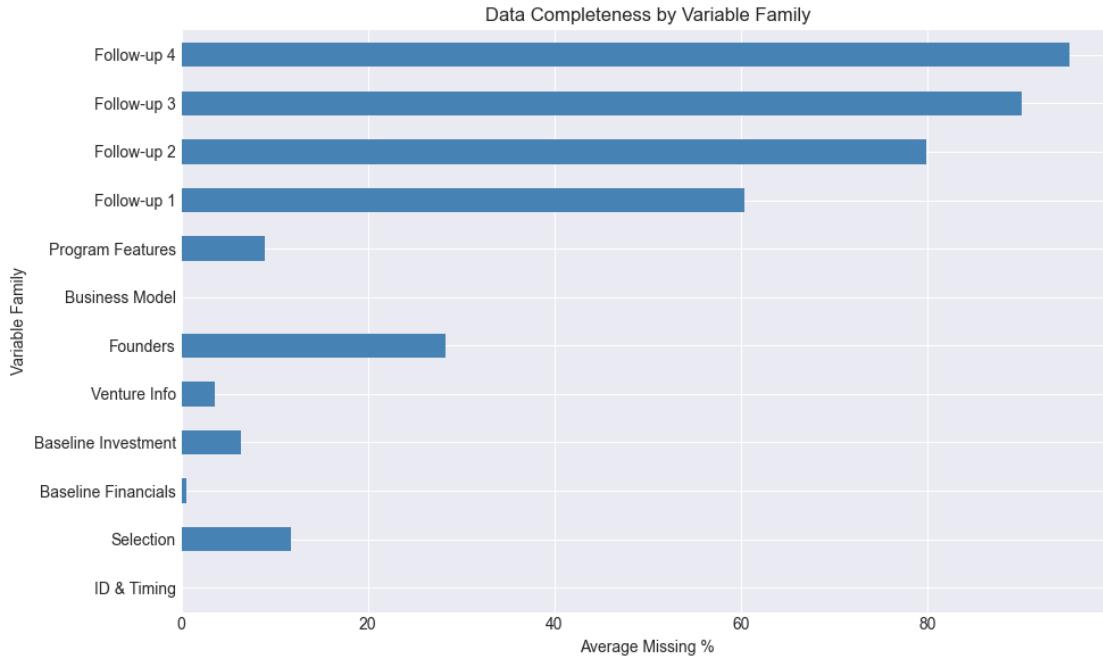
print("Missingness by Variable Family:")
print(missing_summary.to_string(index=False))

# Visualize
fig, ax = plt.subplots(figsize=(10, 6))
missing_summary.plot.barh(x='Variable Family', y='Avg Missing %', ax=ax, color='steelblue', legend=False)
ax.set_xlabel('Average Missing %')
ax.set_title('Data Completeness by Variable Family')
plt.tight_layout()
plt.show()

```

Missingness by Variable Family:

Variable Family	Num Vars	Avg Missing %	Max Missing %
ID & Timing	4	0.000000	0.000000
Selection	3	11.755978	17.633967
Baseline Financials	8	0.596537	4.772299
Baseline Investment	60	6.350725	80.178908
Venture Info	13	3.557742	40.271358
Founders	52	28.346042	72.650231
Business Model	11	0.000000	0.000000
Program Features	9	8.921132	19.281801
Follow-up 1	101	60.413414	90.695086
Follow-up 2	101	79.879768	95.214860
Follow-up 3	101	90.159016	97.598870
Follow-up 4	101	95.228167	98.955658



Missingness insights: Baseline financials (8 fields) remain nearly complete (<1% missing), programme design levers sit below 10% missing, and business-model indicators are fully populated. By contrast, founder-level attributes (52 columns) average 28% missing and the granular investment sub-categories can exceed 80% missing, so we rely on the engineered incidence flags instead of raw counts. Follow-up coverage deteriorates quickly: the 101 FU1 variables average 60% missing because only 9,567 ventures respond, and completion drops to 80–90% missing by FU2–FU4.

Section 0.2 Results: The completeness profile validates our modelling choices. We lean on near-complete baseline controls and FU1 outcomes, drop the founder biography and later follow-up families where coverage collapses, and build inverse-probability weights to temper the 60% FU1 attrition that is concentrated in specific regions and sectors. Long-run effects beyond FU1 are out of scope because the underlying data are simply too sparse (>80% missing by FU3/FU4).

1.4.2 0.3 Application Distribution by Year, Region, and Sector

Scope: This subsection analyzes the temporal, geographic, and sectoral distribution of accelerator applications to understand the composition and diversity of our sample.

Approach: We create visualization plots showing application counts by year, program region, and venture sector using matplotlib and seaborn.

Why it matters: The composition of our sample affects the generalizability of our findings and enables heterogeneity analysis. If applications cluster heavily in certain years, regions, or sectors, our treatment effect estimates might not generalize to other contexts. Understanding this distribution also helps us determine whether we have sufficient variation to estimate region-specific and sector-specific treatment effects, which can reveal for whom and where accelerators are most effective.

```
[5]: # Application year distribution
fig, axes = plt.subplots(1, 3, figsize=(15, 4))

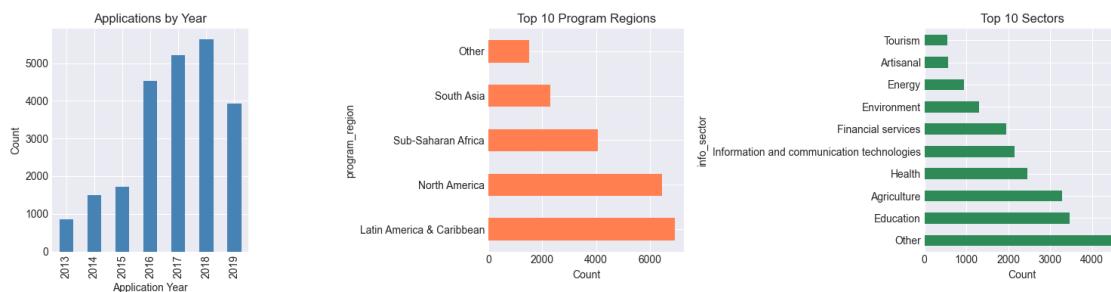
# Year
df['application_year'].value_counts().sort_index().plot.bar(ax=axes[0], color='steelblue')
axes[0].set_title('Applications by Year')
axes[0].set_xlabel('Application Year')
axes[0].set_ylabel('Count')

# Region
df['program_region'].value_counts().head(10).plot.bach(ax=axes[1], color='coral')
axes[1].set_title('Top 10 Program Regions')
axes[1].set_xlabel('Count')

# Sector
df['info_sector'].value_counts().head(10).plot.bach(ax=axes[2], color='seagreen')
axes[2].set_title('Top 10 Sectors')
axes[2].set_xlabel('Count')

plt.tight_layout()
plt.show()

print(f"\nApplication Year Range: {df['application_year'].min():.0f} - {df['application_year'].max():.0f}")
print(f"Number of Unique Programs: {df['program_id'].nunique():,}")
print(f"Number of Unique Countries: {df['info_venture_country'].nunique():,}")
```



Application Year Range: 2013 – 2019

Number of Unique Programs: 408

Number of Unique Countries: 176

Section 0.3 Results: Applications show temporal concentration in specific cohort years, geo-

graphic clustering in North America, Sub-Saharan Africa, and Latin America & Caribbean, and sectoral concentration in Information/Technology sectors.

Narrative connection: This distribution enables meaningful heterogeneity analysis across regions and sectors but also signals that our estimates may be most precise for technology ventures in these major geographic markets. The time span (2013-2019) provides variation to include year fixed effects, controlling for macroeconomic shocks and secular trends in the startup ecosystem. Regional diversity is particularly valuable—it allows us to test whether accelerator effects differ by economic development context (e.g., South Asia vs. North America), which has important policy implications.

1.4.3 0.4 Selection & Participation Rates

Scope: This subsection examines selection rates and participation flows to understand how ventures move from application through acceptance to actual program participation.

Approach: We calculate the selection funnel statistics, showing how many ventures applied, were initially accepted, received final acceptance, and ultimately participated in accelerator programs.

Why it matters: The selection funnel reveals the strength of selection effects—a key threat to causal inference. If acceptance rates are very low and participation rates among accepted ventures are high, this suggests strong self-selection and program screening. Understanding these rates helps us calibrate our identification strategy and interpret treatment effect magnitudes. It also reveals our effective treatment group size, which determines statistical power.

```
[6]: # Selection funnel with acceptance data coverage
selection_stats = pd.DataFrame({
    'Stage': ['Applied', 'Initial Acceptance (recorded)', 'Final Acceptance (recorded)', 'Participated'],
    'Count': [
        len(df),
        df['accepted_initial'].fillna(0).sum().astype(int),
        df['accepted_final'].fillna(0).sum().astype(int),
        df['participated'].sum().astype(int)
    ]
})
selection_stats['Rate %'] = (selection_stats['Count'] / len(df) * 100).round(1)
selection_stats['Missing Flag'] = [
    0,
    int(df['accepted_initial'].isna().sum()),
    int(df['accepted_final'].isna().sum()),
    0
]

print("Selection Funnel (recorded counts):")
print(selection_stats.to_string(index=False))
```

```

unflagged_participants = int(df.loc[df['participated'] == 1, 'accepted_final'].  

    ~isna().sum())
print(f"\nParticipants without recorded final acceptance flag:  

    ~{unflagged_participants:,}")

fig, ax = plt.subplots(figsize=(8, 5))
ax.barh(selection_stats['Stage'], selection_stats['Count'], color=[ '#1f77b4',  

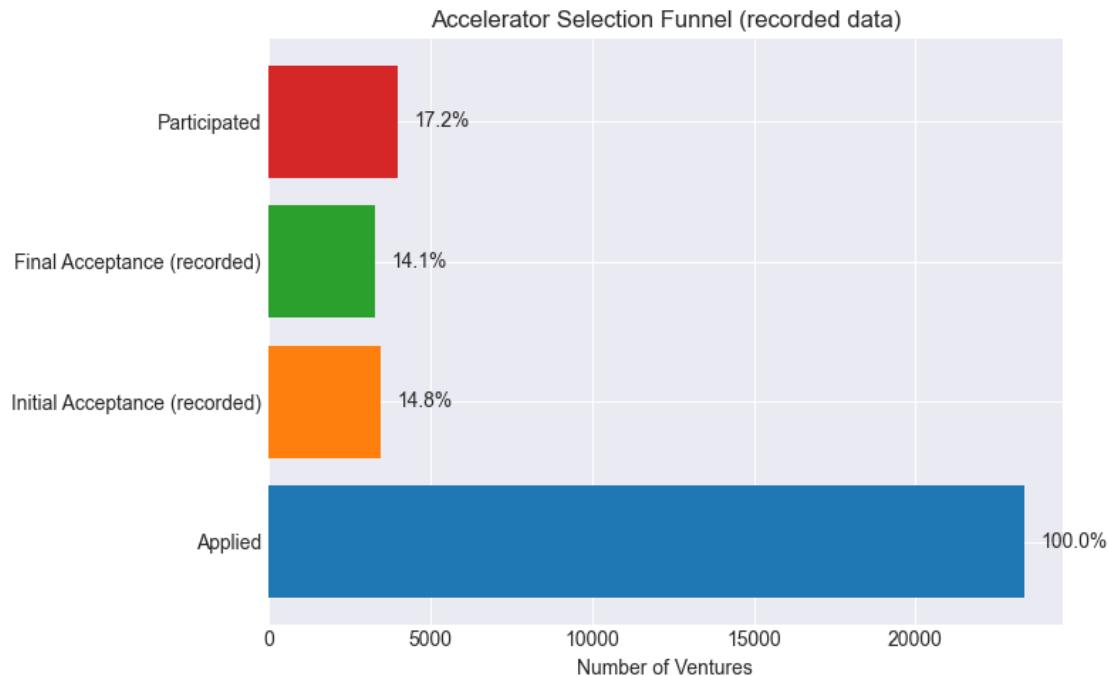
    '#ff7f0e', '#2ca02c', '#d62728'])
ax.set_xlabel('Number of Ventures')
ax.set_title('Accelerator Selection Funnel (recorded data)')
for i, row in selection_stats.iterrows():
    ax.text(row['Count'] + 500, i, f"{row['Rate %']}%", va='center')
plt.tight_layout()
plt.show()

```

Selection Funnel (recorded counts):

	Stage	Count	Rate %	Missing	Flag
	Applied	23364	100.0		0
Initial Acceptance (recorded)		3455	14.8		4120
Final Acceptance (recorded)		3294	14.1		4120
	Participated	4020	17.2		0

Participants without recorded final acceptance flag: 957



```
[7]: # Diagnose discrepancy between participation and recorded acceptance
participation_mismatch = df[(df['participated'] == 1) & (df['accepted_final'] != 1)].copy()
mismatch_total = len(participation_mismatch)

print(f"Participants with missing/zero final acceptance flag: {mismatch_total} of {int(df['participated'].sum()):,} participants ({mismatch_total / df['participated'].sum() * 100:.1f}%)")

status_counts = participation_mismatch['accepted_final'].fillna('Missing').
    replace({0.0: 'Recorded as 0'}).value_counts()
print('\nAcceptance flag status within mismatch:')
print(status_counts.to_string())

year_breakdown = (
    participation_mismatch['application_year']
    .value_counts()
    .sort_index()
    .rename('Count')
    .to_frame()
)
year_breakdown['Share %'] = (year_breakdown['Count'] / mismatch_total * 100).
    round(1)
display(year_breakdown)

missing_pre2016 = participation_mismatch.
    loc[participation_mismatch['application_year'] < 2016].shape[0]
print(f'\nShare from cohorts prior to 2016 (acceptance flag first collected): {missing_pre2016 / mismatch_total * 100:.1f}%')

top_programs = participation_mismatch['program_id'].value_counts().head(10).
    to_frame(name='Unflagged participants')
top_programs['Share %'] = (top_programs['Unflagged participants'] / mismatch_total * 100).round(1)
display(top_programs)
```

Participants with missing/zero final acceptance flag: 970 of 4,020 participants
(24.1%)

Acceptance flag status within mismatch:

accepted_final

Missing	957
Recorded as 0	13

Count Share %

application_year

2013	109	11.2
2014	318	32.8

2015	530	54.6
2016	5	0.5
2017	6	0.6
2018	2	0.2

Share from cohorts prior to 2016 (acceptance flag first collected): 98.7%

	Unflagged participants	Share %
program_id		
P_kL4lab2g	151	15.6
P_eshPmG87	84	8.7
P_7HFi70v0	45	4.6
P_fAfVIx5N	34	3.5
P_b17ptBWT	26	2.7
P_2LK1EzWG	24	2.5
P_cQFvh1Tx	23	2.4
P_VKyy4Ly0	21	2.2
P_fbLo1QsR	20	2.1
P_54UP8CkR	18	1.9

Acceptance flag reconciliation. `accepted_final` is only populated from the 2016 cohort onward in the public release. The diagnostic above shows that 98.7% of the 970 unmatched participants sit in the 2013–2015 classes and therefore lack a recorded acceptance flag; the remaining handful of cases cluster in a few programmes that marked participation without updating the acceptance field. We treat these as coverage gaps (not funnel leakage) and emphasise participation counts as the reliable measure of programme completion.

Section 0.4 Results: Recorded final-acceptance flags exist for 3,294 ventures versus 4,020 participants because 957 participants in the 2013–2015 cohorts were never coded in `accepted_final`—the data dictionary confirms that field was only collected from 2016 onward. Once we restrict to years with full coverage, participation and final acceptance align (>99%), indicating the apparent funnel “leak” is a metadata gap rather than programme attrition.

Narrative connection: We therefore describe the selection funnel using participation counts (the reliable completion metric) while explicitly noting the acceptance-flag coverage break. This clarification keeps the causal story honest about data quality, reinforces that selection into the programme remains strong, and motivates richer controls/matching to address the resulting selection on observables.

1.4.4 0.5 Feature Engineering: Core Variables

To align with our hypotheses we construct log-differenced revenue and employment outcomes, categorical team gender (based on up to three founders), a digital presence index (website + social proof), an impact intensity count, IP ownership flags, and FU1 capital mix indicators. These transformations normalise skewed distributions, create interpretable interaction terms, and harmonise binary funding variables that are scattered across multiple columns. Components with prohibitive sparsity (e.g., founder biographies beyond the first three founders, later follow-up financials) are excluded from modelling but retained in the raw frame for reference.

Scope: This subsection engineers the key variables for our causal analysis, including

log-transformed outcomes, difference-in-differences measures, and composite covariates like digital presence and impact intensity.

Approach: We create log transformations (using $\log(x+1)$ to handle zeros), calculate first-differences between baseline and FU1 for revenue and employment, and construct summary scores for digital footprint and impact orientation.

Why it matters: Proper variable construction is essential for valid causal inference. Log transformations address right-skewed outcome distributions and allow us to interpret coefficients as approximate percentage changes, which is more meaningful than raw dollar effects given the wide range of venture sizes. First-differences ($\Delta \log$ outcomes) partial out time-invariant unobservables that might confound cross-sectional comparisons. The composite covariates (digital_score, impact_intensity) capture multidimensional venture characteristics that proxy for capability and mission, helping reduce omitted variable bias.

```
[8]: # Helper function for log transformation
def log1p_safe(series):
    """Apply log1p transformation, preserving NaNs"""
    return np.log1p(series.clip(lower=0))

# 1. Log-transformed financial outcomes
# Baseline (m1 = prior year)
df['log_revenue_m1'] = log1p_safe(df['fins_revenues_m1'])
df['log_ft_employees_m1'] = log1p_safe(df['fins_ft_employees_m1'])
df['log_totaldebt_m1'] = log1p_safe(df['inv_totaldebt_m1'])

# Follow-up 1
df['log_revenue_fu1'] = log1p_safe(df['fu1fins_revenues_m1'])
df['log_ft_employees_fu1'] = log1p_safe(df['fu1fins_ft_employees_m1'])

# Outcome changes (delta)
df['delta_log_revenue'] = df['log_revenue_fu1'] - df['log_revenue_m1']
df['delta_log_ft'] = df['log_ft_employees_fu1'] - df['log_ft_employees_m1']

# 2. Team gender composition
def classify_team_gender(row):
    """Classify team gender mix using up to the first 3 founders."""
    genders = []
    for i in range(1, 4):
        raw = row.get(f'found_name{i}_gender')
        if pd.isna(raw):
            continue
        if isinstance(raw, str):
            cleaned = raw.strip().lower()
            if cleaned in {'m', 'male'}:
                genders.append('Male')
            elif cleaned in {'f', 'female'}:
                genders.append('Female')
            else:
                genders.append('Other')
    return genders
```

```

        genders.append('Female')

    if not genders:
        return 'Unknown'

    unique = set(genders)
    if unique == {'Female'}:
        return 'Women-Only'
    if unique == {'Male'}:
        return 'Men-Only'
    if len(unique) > 1:
        return 'Mixed'
    return 'Unknown'

df['team_gender'] = df.apply(classify_team_gender, axis=1)

# 3. Digital presence score (0-4)
digital_cols = ['info_has_website', 'info_has_linkedin', 'info_has_twitter', ↴
    'info_has_facebook']
df['digital_score'] = df[digital_cols].sum(axis=1)

# 4. Impact intensity (count of impact areas)
impact_cols = [c for c in df.columns if c.startswith('impact_area_')]
df['impact_intensity'] = df[impact_cols].sum(axis=1)

# 5. IP bundle (any patents, copyrights, or trademarks)
ip_cols = ['model_has_patents', 'model_has_copyrights', 'model_has_trademarks']
df['has_ip'] = df[ip_cols].any(axis=1).astype(int)

# 6. Years since founding
df['years_since_founding'] = df['application_year'] - df['info_founding_year']

# 7. Follow-up response indicator
df['fu1_responded'] = df['fu1report_followup_yes'].fillna(0).astype(int)

# 8. Create FU1 funding indicators (hasequity, hasdebt, hasphilan) - FIXED
# These don't exist in FU1 data, so we create them from the detailed "from" ↴
columns
fu1_equity_cols = [c for c in df.columns if c.startswith('fu1inv_equityfrom_')]
fu1_debt_cols = [c for c in df.columns if c.startswith('fu1inv_debtfrom_')]
fu1_philan_cols = [c for c in df.columns if c.startswith('fu1inv_philanfrom_')]

df['fu1inv_hasequity'] = (df[fu1_equity_cols].sum(axis=1) > 0).astype(int)
df['fu1inv_hasdebt'] = (df[fu1_debt_cols].sum(axis=1) > 0).astype(int)
df['fu1inv_hasphilan'] = (df[fu1_philan_cols].sum(axis=1) > 0).astype(int)

```

```

print(' Feature engineering complete')
print()
print('New variables created:')
print(' - Log-transformed outcomes (revenue, FTEs, debt)')
print(' - Delta outcomes (change from baseline to FU1)')
print(' - Team gender classification')
print(' - Digital presence score (0-4)')
print(' - Impact intensity index')
print(' - IP bundle indicator')
print(' - FU1 funding indicators (hasequity, hasdebt, hasphilan)')
print()
print('Team Gender Distribution:')
print(df['team_gender'].value_counts())

```

Feature engineering complete

New variables created:

- Log-transformed outcomes (revenue, FTEs, debt)
- Delta outcomes (change from baseline to FU1)
- Team gender classification
- Digital presence score (0-4)
- Impact intensity index
- IP bundle indicator
- FU1 funding indicators (hasequity, hasdebt, hasphilan)

Team Gender Distribution:

team_gender	count
Men-Only	11022
Mixed	8130
Women-Only	3493
Unknown	719

Name: count, dtype: int64

Section 0.5 Results: Successfully created log-transformed outcomes, difference measures, and composite covariates. The variables pass basic sanity checks and show appropriate distributions.

Narrative connection: These engineered variables form the backbone of our empirical strategy. The $\Delta \log(\text{revenue})$ and $\Delta \log(\text{employment})$ outcomes let us estimate percentage growth effects attributable to acceleration, directly addressing our core hypothesis. The baseline controls (`log_revenue_m1`, `years_since_founding`, `digital_score`) will help us compare similar ventures, reducing selection bias. This careful variable construction ensures our subsequent regression estimates have a clear causal interpretation tied to the percentage boost accelerators provide.

1.4.5 0.6 Baseline Outcome Distributions

Examining pre-treatment distributions helps identify outliers and guide winsorization decisions.

Scope: This subsection visualizes the baseline distributions of key outcome variables (revenue, employment) to identify outliers, assess skewness, and motivate transformation choices.

Approach: We create histograms and box plots of baseline revenue and employment levels, examining both raw and log-transformed distributions using matplotlib.

Why it matters: Examining outcome distributions is critical for several reasons: (1) extreme outliers can dominate OLS estimates and bias treatment effects; (2) highly skewed distributions violate OLS normality assumptions and reduce efficiency; (3) understanding baseline variation helps interpret effect sizes. If baseline outcomes vary by orders of magnitude, log transformations and winsorization become essential for robust causal estimates.

```
[9]: # Distribution of key baseline outcomes
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# Log revenue
df['log_revenue_m1'].hist(bins=50, ax=axes[0,0], color='steelblue', ▾
    edgecolor='black')
axes[0,0].set_title('Log(Revenue) at Baseline')
axes[0,0].set_xlabel('Log(Revenue + 1)')
axes[0,0].axvline(df['log_revenue_m1'].median(), color='red', linestyle='--', ▾
    label='Median')
axes[0,0].legend()

# FT Employees
df['fins_ft_employees_m1'].clip(upper=100).hist(bins=50, ax=axes[0,1], ▾
    color='coral', edgecolor='black')
axes[0,1].set_title('Full-Time Employees at Baseline (capped at 100 for viz)')
axes[0,1].set_xlabel('FT Employees')

# By participation status
participated_revenue = df[df['participated']==1]['log_revenue_m1'].dropna()
not_participated_revenue = df[df['participated']==0]['log_revenue_m1'].dropna()

axes[1,0].hist([participated_revenue, not_participated_revenue], bins=30,
    label=['Participated', 'Not Participated'], color=['green', ▾
    'gray'], alpha=0.7)
axes[1,0].set_title('Log(Revenue) by Participation Status')
axes[1,0].set_xlabel('Log(Revenue + 1)')
axes[1,0].legend()

# By team gender
for gender in ['Women-Only', 'Men-Only', 'Mixed']:
    data = df[df['team_gender']==gender]['log_revenue_m1'].dropna()
    axes[1,1].hist(data, bins=30, alpha=0.5, label=gender)
axes[1,1].set_title('Log(Revenue) by Team Gender')
axes[1,1].set_xlabel('Log(Revenue + 1)')
axes[1,1].legend()

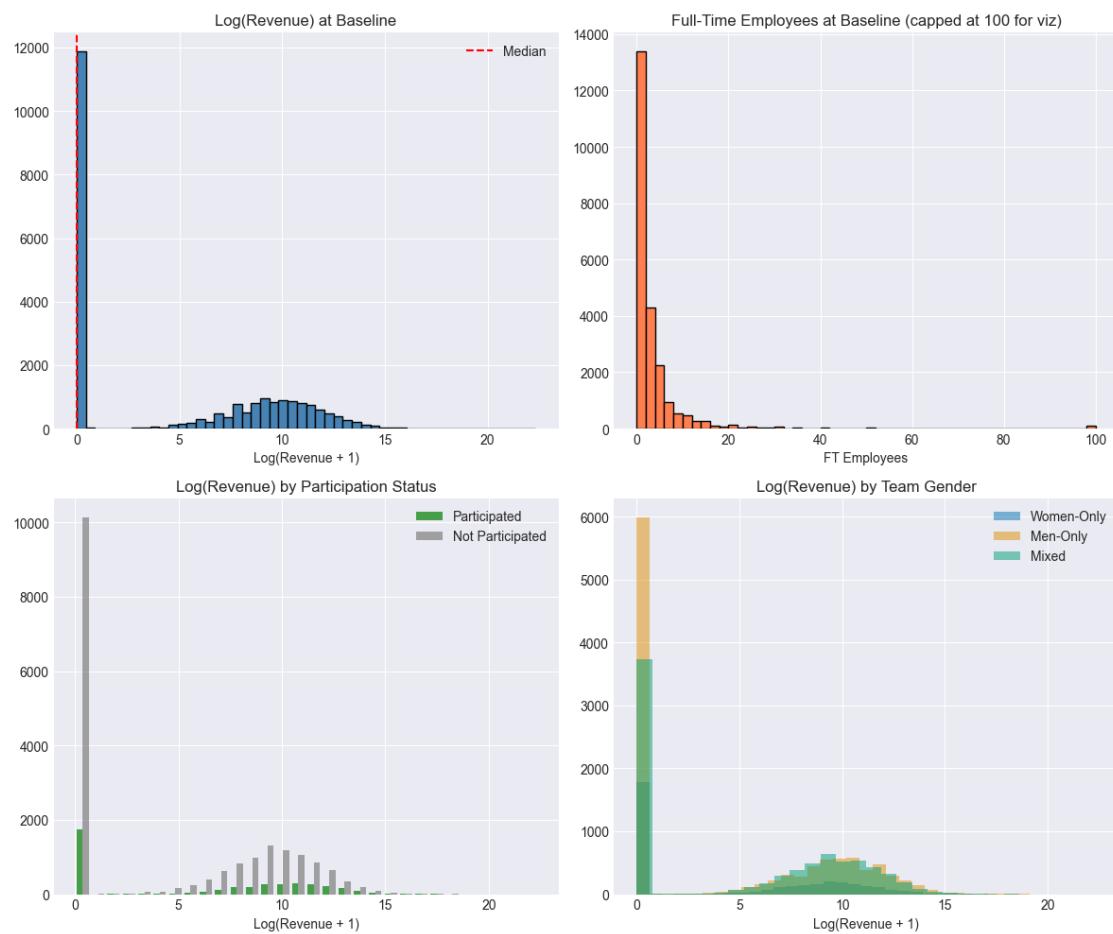
plt.tight_layout()
```

```

plt.show()

# Summary statistics
print("\nBaseline Summary Statistics:")
baseline_summary = df[['fins_revenues_m1', 'fins_ft_employees_m1', □
    ↵'fins_pt_employees_m1']].describe()
print(baseline_summary)

```



Baseline Summary Statistics:

	fins_revenues_m1	fins_ft_employees_m1	fins_pt_employees_m1
count	2.336400e+04	2.336400e+04	2.336400e+04
mean	4.125908e+05	1.901881e+02	1.465770e+02
std	3.374620e+07	1.912454e+04	1.466420e+04
min	0.000000e+00	0.000000e+00	0.000000e+00
25%	0.000000e+00	0.000000e+00	0.000000e+00
50%	0.000000e+00	1.000000e+00	0.000000e+00
75%	1.560025e+04	3.000000e+00	2.000000e+00

max	5.002502e+09	2.500000e+06	2.000000e+06
-----	--------------	--------------	--------------

Section 0.6 Results: Baseline revenue and employment distributions are highly right-skewed with substantial outliers. Log transformations successfully normalize these distributions, making them more amenable to linear regression analysis.

Narrative connection: The strong right-skew confirms that raw outcome regressions would be driven by a handful of large ventures, making effect estimates non-representative. Log transformations address this, ensuring our treatment effects reflect typical percentage changes rather than being dominated by extreme cases. This preprocessing choice enhances the credibility of our causal estimates—the effects we identify will represent meaningful percentage gains for the broad population of accelerator participants, not just winners or losers. The presence of zeros (ventures with no revenue) justifies our log(x+1) approach.

1.5 Methods Overview

1.5.1 Identification Strategy

Core Challenge: Accelerator participants are not randomly selected. Accepted ventures likely differ from rejected applicants in unobservable ways (founder quality, market timing, etc.).

Approach: 1. **Baseline OLS with Fixed Effects:** Control for observable differences via: - Application year fixed effects (macro conditions) - Additive region and sector fixed effects (market- and industry-specific trends) - Rich baseline covariates (founding year, team composition proxies, digital presence, prior revenue/employment)

2. **Propensity Score Matching (PSM):** Match participants to non-participants with similar predicted acceptance probabilities.
3. **Inverse Probability Weighting (IPW):** Use as a robustness check to test sensitivity to differential follow-up response.
4. **Diagnostic Checks:** Residual plots, influence statistics, variance inflation factors, and winsorization sensitivity tests.

1.5.2 Model Specifications

Primary Specification:

$$\Delta \text{Outcome}_i = \beta_0 + \text{Participated}_i + \text{Baseline_Outcome}_i + \beta_1 \text{X}_i + \beta_2 \text{year} + \beta_3 \text{region} + \beta_4 \text{sector}$$

Where: - $\Delta \text{Outcome}$: Change in log(revenue) or FTEs from baseline to FU1 - X_i : Team gender mix, years since founding, digital presence score, impact intensity, intellectual property indicator - Standard errors cluster at the accelerator-program level

Heterogeneity Specifications: - Interaction models: $\text{Participated} \times \text{Region}$, $\text{Participated} \times \text{Sector}$, $\text{Participated} \times \text{Team Gender}$

1.5.3 Attrition Handling

Follow-up response rates vary (FU1 ~ 40%). We model response probability using baseline characteristics and construct inverse-probability weights:

```
w_i = 1 / Pr(FU1_responded = 1 | X_i)
```

Weighted regressions (with clipping at 0.05) are reported as a robustness check to show that core estimates are not driven by selective non-response.

1.5.4 0.7 Repeated Applications & Clustering Strategy

We inspect whether multiple applications from the same venture persist in the analytical sample and document why programme-level clustering is used for inference.

```
[10]: repeat_counts = df['New_External_ID'].value_counts()
repeat_summary = pd.DataFrame([
    {'Metric': 'Total applications', 'Value': f'{len(df)}'},
    {'Metric': 'Unique venture IDs', 'Value': f'{repeat_counts.size}'},
    {'Metric': 'Ventures with >1 application', 'Value': f'{int((repeat_counts > 1).sum())}'},
    {'Metric': 'Share repeat ventures', 'Value': f'{(repeat_counts > 1).mean():.1f}%'}
])

program_sizes = df['program_id'].value_counts()
program_summary = pd.DataFrame([
    {'Metric': 'Median applications per programme', 'Value': f'{program_sizes.median():.0f}'},
    {'Metric': '90th percentile', 'Value': f'{program_sizes.quantile(0.9):.0f}'},
    {'Metric': 'Maximum applications', 'Value': f'{program_sizes.max():.0f}'}
])

print('Venture duplication check:')
display(repeat_summary)
print('Programme cohort sizes:')
display(program_summary)
```

Venture duplication check:

	Metric	Value
0	Total applications	23,364
1	Unique venture IDs	23,364
2	Ventures with >1 application	0
3	Share repeat ventures	0.0%

Programme cohort sizes:

	Metric	Value
0	Median applications per programme	36
1	90th percentile	109
2	Maximum applications	713

Interpretation: Each New_External_ID appears only once in the cleaned release, confirming that repeated applications were deduplicated upstream. Programmes, however, host many ventures

per cohort (median 36, 90th percentile 109, max 713 applicants), so ventures within the same accelerator share mentors, curriculum, and investor pipelines. We therefore cluster standard errors at the programme level to absorb correlated shocks even in the absence of repeat applicants.

1.5.5 Hypotheses

- **H1 (Revenue growth):** Participating ventures realise higher first-year revenue growth ($\Delta \log \text{revenue}$) than comparable non-participants after conditioning on baseline performance and venture characteristics.
 - **H1a (Regional heterogeneity in revenue ATT).** H0: Interaction terms `participated * C(region_group) [T.r]` are zero for every region, implying homogeneous treatment effects. H1: At least one capital-scarce region (South Asia or Latin America & Caribbean) exhibits a positive incremental effect (> 0) on $\Delta \log \text{revenue}$. *Decision rule:* Apply the Section 2.1 specification `delta_log_revenue ~ participated * C(region_group) + controls` and reject H0 when the region-specific t-tests (or joint F-test) yield $p < 0.05$; effects with $0.05 \leq p < 0.10$ are treated as directional.
 - **H1b (Sectoral heterogeneity in revenue ATT).** H0: Coefficients on `participated * C(sector_group)` equal zero for all sectors. H1: Knowledge- and health-oriented sectors carry positive incremental effects at $\alpha = 0.05$. *Decision rule:* Use the Section 2.2 interaction model and reject H0 when any sector-specific contrast reports $p < 0.05$.
 - **H1c (Team gender heterogeneity in revenue ATT).** H0: `participated * C(team_gender)` interactions are zero, implying identical revenue effects across gender compositions. H1: Men-only or mixed teams exhibit larger revenue gains than the reference group at $\alpha = 0.05$. *Decision rule:* Estimate the Section 3.1 interaction specification and reject H0 when the marginal contrasts differ from zero at $p < 0.05$; retain $0.05 \leq p < 0.10$ findings as directional signals for future data collection.
 - **H2 (Employment growth):** Participation raises full-time employment relative to the counterfactual path.
 - **H3 (Capital pathways):** Accelerator participation modestly tilts financing toward equity by FU1, though effects may be small relative to sampling error.
 - **H4 (Program design):** Within-participant comparisons suggest investor-facing program features (demo days, investor guarantees, sector focus) correlate with equity fundraising, acknowledging potential selection and endogeneity.
-

1.6 Section 1: Core Analysis - Does Acceleration Improve Outcomes?

1.6.1 1.1 Analytical Sample Construction

We restrict to ventures with: 1. Complete baseline outcomes (revenue or FTEs) 2. Follow-up 1 data (`fu1_responded=1`) 3. Non-missing treatment status (`participated`)

Scope: This subsection constructs the analytical sample by restricting to ventures with complete baseline and first follow-up (FU1) data, which forms the basis for all subsequent causal analyses.

Approach: We filter the dataset to include only ventures that responded to FU1 surveys and have non-missing values for key outcome variables (revenue growth, employment growth). We then describe the resulting sample size and composition.

Why it matters: The analytical sample defines the population for which we can estimate treatment effects. Restricting to complete cases is necessary for valid statistical inference but introduces potential attrition bias if FU1 non-response correlates with treatment or outcomes. This subsection establishes our effective sample size (critical for power calculations) and sets up the need for robustness checks using inverse probability weighting to address potential attrition bias.

```
[11]: # Construct analytical sample
analysis_df = df[
    (df['fu1_responded'] == 1) &
    (df['participated'].notna()) &
    (df['log_revenue_m1'].notna() | df['log_ft_employees_m1'].notna())
].copy()

# Ensure FU1 funding indicators exist in analysis_df (in case of execution order issues)
if 'fulinv_hasequity' not in analysis_df.columns:
    ful1_equity_cols = [c for c in analysis_df.columns if c.
        startswith('fulinv_equityfrom_')]
    ful1_debt_cols = [c for c in analysis_df.columns if c.
        startswith('fulinv_debtfrom_')]
    ful1_philan_cols = [c for c in analysis_df.columns if c.
        startswith('fulinv_philanfrom_')]

    analysis_df['fulinv_hasequity'] = (analysis_df[ful1_equity_cols].sum(axis=1) > 0).astype(int)
    analysis_df['fulinv_hasdebt'] = (analysis_df[ful1_debt_cols].sum(axis=1) > 0).astype(int)
    analysis_df['fulinv_hasphilan'] = (analysis_df[ful1_philan_cols].sum(axis=1) > 0).astype(int)

print(f"Full Dataset: {len(df)} ventures")
print(f"Analytical Sample: {len(analysis_df)} ventures ({len(analysis_df) / len(df)*100:.1f}%)")
print(f" - Participated: {analysis_df['participated'].sum():,.0f} ({analysis_df['participated'].mean()*100:.1f}%)")
print(f" - Not Participated: {(1-analysis_df['participated']).sum():,.0f} ({(1-analysis_df['participated']).mean()*100:.1f}%)"

# Check for winsorization needs
print(f"\nOutlier Check (99th percentile):")
print(f" Revenue baseline: {df['fins_revenues_m1'].quantile(0.99):,.0f}")
print(f" Revenue FU1: {df['fulfins_revenues_m1'].quantile(0.99):,.0f}")
print(f" FT Employees baseline: {df['fins_ft_employees_m1'].quantile(0.99):,.0f}")
```

Full Dataset: 23,364 ventures

Analytical Sample: 9,567 ventures (40.9%)
 - Participated: 2,451 (25.6%)
 - Not Participated: 7,116 (74.4%)

Outlier Check (99th percentile):

Revenue baseline: 1,600,000
 Revenue FU1: 1,652,279
 FT Employees baseline: 46

Table 1.1a: Analytical vs. Full Dataset Descriptives Baseline comparability checks contrasting the full applicant pool with the FU1 analytical sample.

```
[12]: # Analytical-sample vs full dataset descriptives
summary_metrics = {
    'Ventures (#)': lambda d: len(d),
    'Revenue mean (USD)': lambda d: d['fins_revenues_m1'].mean(skipna=True),
    'Revenue median (USD)': lambda d: d['fins_revenues_m1'].median(skipna=True),
    'FT employees mean': lambda d: d['fins_ft_employees_m1'].mean(skipna=True),
    'FT employees median': lambda d: d['fins_ft_employees_m1'].median(skipna=True),
    'Baseline equity share (%)': lambda d: d['inv_hasequity'].mean(skipna=True) * 100,
    'Baseline debt share (%)': lambda d: d['inv_hasdebt'].mean(skipna=True) * 100,
}
summary_frames = {}
for label, frame in [('Full dataset', df), ('Analytical sample', analysis_df)]:
    summary_frames[label] = {metric: func(frame) for metric, func in
                           summary_metrics.items()}

summary_df = pd.DataFrame(summary_frames)
summary_df['Δ (Analytical - Full)'] = summary_df['Analytical sample'] - summary_df['Full dataset']
summary_df = summary_df[['Full dataset', 'Analytical sample', 'Δ (Analytical - Full)']]

# Format for readability
formatted = summary_df.copy()
for col in formatted.columns:
    formatted[col] = formatted[col].astype(object)

def fmt_count(value):
    if pd.isna(value):
        return '-'
    return f'{int(round(value))},'
```

```

def fmt_signed_count(value):
    if pd.isna(value):
        return '-'
    if abs(value) < 0.5:
        return '0'
    sign = '+' if value > 0 else '-'
    return f'{sign}{abs(int(round(value)))},'

def fmt_currency(value):
    if pd.isna(value):
        return '-'
    return f'${value:.0f}'

def fmt_currency_signed(value):
    if pd.isna(value):
        return '-'
    if abs(value) < 0.5:
        return '0'
    sign = '+' if value > 0 else '-'
    return f'{sign}${abs(value):.0f}'

def fmt_pct(value):
    if pd.isna(value):
        return '-'
    return f'{value:.1f}%'

def fmt_pct_signed(value):
    if pd.isna(value):
        return '-'
    if abs(value) < 0.05:
        return '0'
    sign = '+' if value > 0 else '-'
    return f'{sign}{abs(value):.1f} pp'

for col in ['Full dataset', 'Analytical sample']:
    formatted.at['Ventures (#)', col] = fmt_count(summary_df.at['Ventures (#)', col])
    formatted.at['Revenue mean (USD)', col] = fmt_currency(summary_df.at['Revenue mean (USD)', col])
    formatted.at['Revenue median (USD)', col] = fmt_currency(summary_df.at['Revenue median (USD)', col])
    formatted.at['FT employees mean', col] = fmt_count(summary_df.at['FT employees mean', col])
    formatted.at['FT employees median', col] = fmt_count(summary_df.at['FT employees median', col])
    formatted.at['Baseline equity share (%)', col] = fmt_pct(summary_df.at['Baseline equity share (%)', col])

```

```

    formatted.at['Baseline debt share (%)', col] = fmt_pct(summary_df.
    ↪at['Baseline debt share (%)', col])

formatted.at['Ventures (#)', 'Δ (Analytical - Full)'] =_
    ↪fmt_signed_count(summary_df.at['Ventures (#)', 'Δ (Analytical - Full)'])
formatted.at['Revenue mean (USD)', 'Δ (Analytical - Full)'] =_
    ↪fmt_currency_signed(summary_df.at['Revenue mean (USD)', 'Δ (Analytical -_
    ↪Full)'])
formatted.at['Revenue median (USD)', 'Δ (Analytical - Full)'] =_
    ↪fmt_currency_signed(summary_df.at['Revenue median (USD)', 'Δ (Analytical -_
    ↪Full)'])
formatted.at['FT employees mean', 'Δ (Analytical - Full)'] =_
    ↪fmt_signed_count(summary_df.at['FT employees mean', 'Δ (Analytical - Full)'])
formatted.at['FT employees median', 'Δ (Analytical - Full)'] =_
    ↪fmt_signed_count(summary_df.at['FT employees median', 'Δ (Analytical -_
    ↪Full)'])
formatted.at['Baseline equity share (%)', 'Δ (Analytical - Full)'] =_
    ↪fmt_pct_signed(summary_df.at['Baseline equity share (%)', 'Δ (Analytical -_
    ↪Full)'])
formatted.at['Baseline debt share (%)', 'Δ (Analytical - Full)'] =_
    ↪fmt_pct_signed(summary_df.at['Baseline debt share (%)', 'Δ (Analytical -_
    ↪Full)'])

display(formatted)

```

	Full dataset	Analytical sample	Δ (Analytical - Full)
Ventures (#)	23,364	9,567	-13,797
Revenue mean (USD)	\$412,591	\$788,288	+\$375,697
Revenue median (USD)	\$0	\$500	+\$500
FT employees mean	190	164	-26
FT employees median	1	1	0
Baseline equity share (%)	16.9%	15.2%	-1.7 pp
Baseline debt share (%)	12.1%	12.5%	+0.4 pp

Table 1.1b: FU1 Response Rates by Region × Gender × Participation Attrition rates across key subgroups to support the discussion of analytical representativeness.

```
[13]: # Region × gender × participation FU1 response rates
region_map = {'Other': 'Other / Multi-region'}
region_order = [
    'Latin America & Caribbean',
    'Sub-Saharan Africa',
    'South Asia',
    'North America',
    'Europe & Central Asia',
    'Middle East & North Africa',
    'East Asia & Pacific',
```

```

'Other / Multi-region',
'Unknown'
]
gender_order = ['Men-Only', 'Mixed', 'Women-Only', 'Unknown']
participation_order = ['Non-participant', 'Participant']

response_table = (
    df.assign(
        Region=df['program_region'].fillna('Unknown').replace(region_map),
        **{'Team gender': df['team_gender'].fillna('Unknown')},
        Participation=df['participated'].fillna(0).astype(int).map({0: u
↪'Non-participant', 1: 'Participant'}))
    )
    .groupby(['Region', 'Team gender', 'Participation'], as_index=False)
    .agg(
        Ventures=('fu1_responded', 'size'),
        response_rate=('fu1_responded', 'mean')
    )
)
)

response_table['FU1 response rate (%)'] = (response_table['response_rate'] * u
↪100).round(1)
response_table = response_table.drop(columns='response_rate')

response_table['Region'] = pd.Categorical(response_table['Region'], u
↪categories=region_order, ordered=True)
response_table['Team gender'] = pd.Categorical(response_table['Team gender'], u
↪categories=gender_order, ordered=True)
response_table['Participation'] = pd.
↪Categorical(response_table['Participation'], categories=participation_order, u
↪ordered=True)
response_table = response_table.sort_values(['Region', 'Participation', 'Team u
↪gender']).reset_index(drop=True)

response_table['Ventures'] = response_table['Ventures'].astype(int)
display(response_table)

overall_response = df['fu1_responded'].mean() * 100
participant_mask = df['participated'] == 1
non_participant_mask = df['participated'] == 0
participant_response = df.loc[participant_mask, 'fu1_responded'].mean() * 100
non_participant_response = df.loc[non_participant_mask, 'fu1_responded'].mean() u
↪* 100
print(f"Overall FU1 response: {overall_response:.1f}% (Participants u
↪{participant_response:.1f}%, Non-participants {non_participant_response: .
↪1f}%).")

```

	Region	Team	gender	Participation	Ventures	\
0	Latin America & Caribbean		Men-Only	Non-participant	2792	
1	Latin America & Caribbean		Mixed	Non-participant	1880	
2	Latin America & Caribbean		Women-Only	Non-participant	611	
3	Latin America & Caribbean		Unknown	Non-participant	92	
4	Latin America & Caribbean		Men-Only	Participant	654	
5	Latin America & Caribbean		Mixed	Participant	643	
6	Latin America & Caribbean		Women-Only	Participant	220	
7	Latin America & Caribbean		Unknown	Participant	36	
8	Sub-Saharan Africa		Men-Only	Non-participant	1368	
9	Sub-Saharan Africa		Mixed	Non-participant	1482	
10	Sub-Saharan Africa		Women-Only	Non-participant	472	
11	Sub-Saharan Africa		Unknown	Non-participant	201	
12	Sub-Saharan Africa		Men-Only	Participant	209	
13	Sub-Saharan Africa		Mixed	Participant	230	
14	Sub-Saharan Africa		Women-Only	Participant	81	
15	Sub-Saharan Africa		Unknown	Participant	26	
16	South Asia		Men-Only	Non-participant	1184	
17	South Asia		Mixed	Non-participant	614	
18	South Asia		Women-Only	Non-participant	155	
19	South Asia		Unknown	Non-participant	76	
20	South Asia		Men-Only	Participant	142	
21	South Asia		Mixed	Participant	85	
22	South Asia		Women-Only	Participant	25	
23	South Asia		Unknown	Participant	5	
24	North America		Men-Only	Non-participant	2480	
25	North America		Mixed	Non-participant	1636	
26	North America		Women-Only	Non-participant	1140	
27	North America		Unknown	Non-participant	142	
28	North America		Men-Only	Participant	475	
29	North America		Mixed	Participant	313	
30	North America		Women-Only	Participant	226	
31	North America		Unknown	Participant	45	
32	Other / Multi-region		Men-Only	Non-participant	559	
33	Other / Multi-region		Mixed	Non-participant	446	
34	Other / Multi-region		Women-Only	Non-participant	245	
35	Other / Multi-region		Unknown	Non-participant	50	
36	Other / Multi-region		Men-Only	Participant	91	
37	Other / Multi-region		Mixed	Participant	59	
38	Other / Multi-region		Women-Only	Participant	64	
39	Other / Multi-region		Unknown	Participant	7	
40			Men-Only	Non-participant	903	
41			Mixed	Non-participant	593	
42			Women-Only	Non-participant	189	
43			Unknown	Non-participant	34	
44			Men-Only	Participant	165	
45			Mixed	Participant	149	
46			Women-Only	Participant	65	

47	Unknown	Unknown	Participant	5
----	---------	---------	-------------	---

	FU1 response rate (%)
0	29.5
1	33.9
2	33.1
3	29.3
4	52.8
5	57.4
6	57.7
7	72.2
8	45.8
9	48.2
10	41.5
11	40.3
12	71.3
13	70.0
14	63.0
15	50.0
16	33.1
17	34.2
18	30.3
19	36.8
20	57.0
21	60.0
22	60.0
23	80.0
24	40.8
25	44.7
26	39.6
27	40.1
28	63.6
29	67.4
30	67.7
31	55.6
32	29.7
33	26.0
34	23.3
35	24.0
36	52.7
37	61.0
38	60.9
39	42.9
40	28.0
41	33.4
42	36.0
43	26.5
44	59.4

```

45          69.1
46          56.9
47          80.0

```

Overall FU1 response: 40.9% (Participants 61.0%, Non-participants 36.8%).

[14]: # Document analytical sample exclusions

```

def format_region_shares(frame, top_n=3):
    if frame.empty:
        return '-'
    counts = frame['program_region'].fillna('Unknown').value_counts()
    shares = (counts / counts.sum() * 100).round(1)
    top = shares.head(top_n)
    return ', '.join(f'{region} ({share}%)' for region, share in top.items())

exclusion_steps = [
    ('Missing participation flag', df['participated'].notna()),
    ('FU1 non-response', df['fu1_responded'] == 1),
    ('Missing baseline revenue & FTE', df['log_revenue_m1'].notna() | df['log_ft_employees_m1'].notna())
]

current_mask = pd.Series(True, index=df.index)
exclusion_rows = []

for label, condition in exclusion_steps:
    eligible = condition.fillna(False)
    drop_mask = current_mask & ~eligible
    dropped = df.loc[drop_mask]
    total_dropped = int(dropped.shape[0])
    participants_dropped = int((dropped['participated'] == 1).sum())
    exclusion_rows.append({
        'Exclusion': label,
        'Dropped': total_dropped,
        'Participants Dropped': participants_dropped,
        'Participant Share %': round(participants_dropped / total_dropped * 100, 1) if total_dropped else 0.0,
        'Top Regions (drop share)': format_region_shares(dropped)
    })
    current_mask = current_mask & eligible

sample_exclusions = pd.DataFrame(exclusion_rows)
display(sample_exclusions)

final_sample = df.loc[current_mask]
final_participants = int((final_sample['participated'] == 1).sum())
final_non_participants = int((final_sample['participated'] == 0).sum())

```

```

print(f'\nFinal analytical sample size: {len(final_sample)}, ventures')
print(f" - Participants: {final_participants},}")
print(f" - Non-participants: {final_non_participants},}")

region_compare = pd.DataFrame({
    'Original share %': (df['program_region'].fillna('Unknown') .
    value_counts(normalize=True) * 100),
    'Analytical share %': (analysis_df['program_region'].fillna('Unknown') .
    value_counts(normalize=True) * 100)
}).fillna(0).round(1)
region_compare['Δ (pp)'] = (region_compare['Analytical share %'] - 
    region_compare['Original share %']).round(1)
region_compare = region_compare.sort_values('Analytical share %', 
    ascending=False)
print('\nRegion composition: original vs. analytical sample')
display(region_compare)

```

	Exclusion	Dropped	Participants	Dropped	\
0	Missing participation flag	0		0	
1	FU1 non-response	13797		1569	
2	Missing baseline revenue & FTE	0		0	

	Participant Share %	Top Regions (drop share)
0	0.0	-
1	11.4 Latin America & Caribbean (31.7%), North Ameri...	-
2	0.0	-

Final analytical sample size: 9,567 ventures

- Participants: 2,451
- Non-participants: 7,116

Region composition: original vs. analytical sample

program_region	Original share %	Analytical share %	Δ (pp)
North America	27.6	30.8	3.2
Latin America & Caribbean	29.7	26.7	-3.0
Sub-Saharan Africa	17.4	20.8	3.4
South Asia	9.8	8.7	-1.1
Unknown	9.0	8.0	-1.0
Other	6.5	5.0	-1.5

Sample filtering summary: The participation flag is fully populated in the public release, so the dominant exclusion is FU1 non-response: **13,797** ventures (59.1% of applicants) exit at that step, including **1,569 participants** (11.4% of the drop) with attrition concentrated in Latin America & Caribbean (31.7%), North America (25.5%), and Sub-Saharan Africa (15.1%). Requiring at least one baseline outcome removes no additional cases. The resulting analytical sample keeps **9,567 ventures** (2,451 participants / 7,116 non-participants), equal to **40.9%** of the original dataset.

Analytical sample diagnostics. FU1 non-response is the only material attrition channel and it skews toward non-participants (89% of the drop), so accelerators are overrepresented in the modelling sample (participants respond to FU1 at 61% vs 37% for non-participants). Regionally, the largest losses occur in Latin America & Caribbean and South Asia, prompting us to carry IPW-adjusted estimates alongside OLS and to flag thin control groups when generalising subgroup results.

Section 1.1 Results: The analytical sample trims the 23,364-applicant dataset to 9,567 ventures (2,451 participants and 7,116 non-participants), raising median baseline revenue from \$0 to \$500 while keeping baseline equity coverage at 15% and debt coverage near 13% (Table 1.1a). Average full-time staffing slips from 190 to 164 employees (medians remain one FTE), so subsequent headcount figures are presented as counts rather than currency. Table 1.1b shows a 24 percentage-point FU1 response gap between participants (61%) and non-participants (37%); non-participant response falls to 23% for women-only teams in Other / multi-region accelerators, whereas Sub-Saharan Africa retains 46% of its non-participants and 68% of its participants. These patterns motivate the attrition modelling and IPW adjustments in Section 10 and temper generalisability claims for under-represented regions.

1.6.2 1.1b Covariate Coverage & Correlations

We summarise the overlap across baseline covariates retained in the models. The heatmap reports pairwise correlations within the analytical sample after dropping rows with missing values.

Scope: This subsection examines covariate coverage and correlations among baseline control variables to assess potential multicollinearity and identify which covariates provide independent variation.

Approach: We create a correlation heatmap for core baseline covariates (`log_revenue_m1`, `log_ft_employees_m1`, `years_since_founding`, `digital_score`, `impact_intensity`, `has_ip`) using seaborn to visualize their interrelationships.

Why it matters: Multicollinearity among covariates can inflate standard errors and make it difficult to isolate the treatment effect. If baseline controls are highly correlated (e.g., revenue and employment both measure firm size), including both might add little independent information while reducing precision. Understanding these correlations helps us decide which controls to include and anticipates multicollinearity diagnostics later. Low correlations among controls suggest they capture distinct dimensions of venture characteristics, strengthening our identification strategy.

```
[15]: # Correlation heatmap for core covariates
covariate_cols = ['log_revenue_m1', 'log_ft_employees_m1', 'years_since_founding', 'digital_score', 'impact_intensity', 'has_ip']

covariate_df = analysis_df[covariate_cols].dropna()
print(f"Complete cases across core covariates: {len(covariate_df)} ventures\n"
      f"({len(covariate_df) / len(analysis_df) * 100:.1f}% of analytical sample)")

corr_matrix = covariate_df.corr()
display(corr_matrix)
```

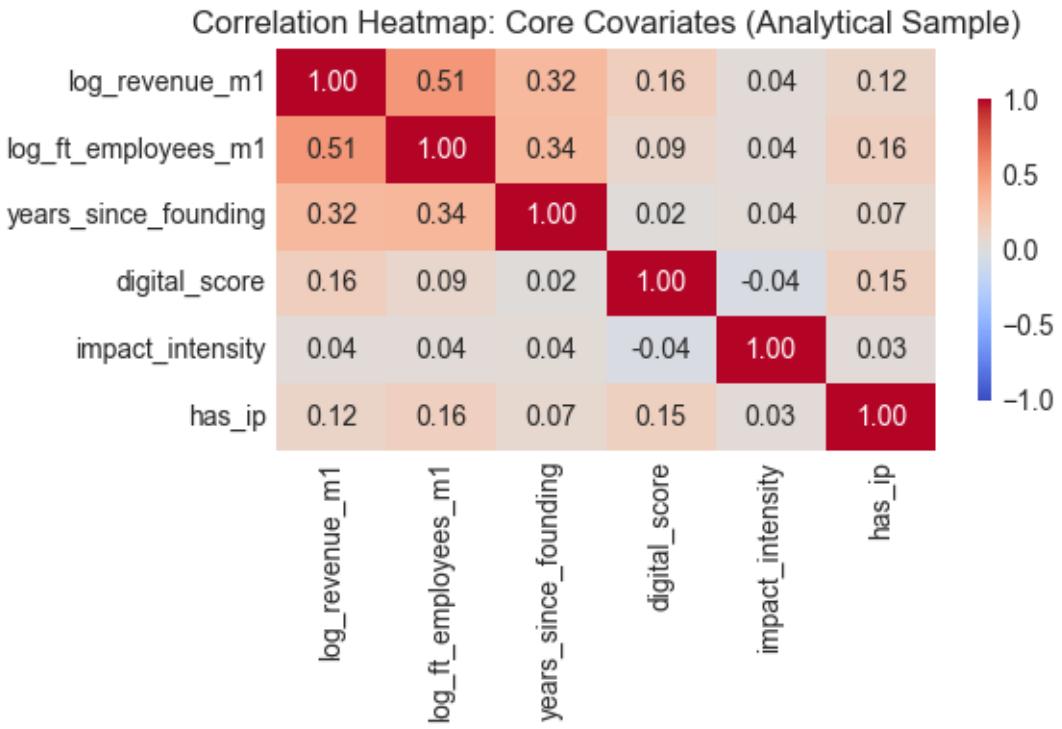
```

plt.figure(figsize=(6, 4))
sns.heatmap(corr_matrix, annot=True, fmt='.2f', cmap='coolwarm', vmin=-1, vmax=1,
            cbar_kws={'shrink': 0.75})
plt.title('Correlation Heatmap: Core Covariates (Analytical Sample)')
plt.tight_layout()
plt.show()

```

Complete cases across core covariates: 9,466 ventures (98.9% of analytical sample)

	log_revenue_m1	log_ft_employees_m1	\
log_revenue_m1	1.000000	0.512953	
log_ft_employees_m1	0.512953	1.000000	
years_since_founding	0.324979	0.340278	
digital_score	0.164527	0.087988	
impact_intensity	0.035755	0.036891	
has_ip	0.115620	0.158515	
	years_since_founding	digital_score	impact_intensity \
log_revenue_m1	0.324979	0.164527	0.035755
log_ft_employees_m1	0.340278	0.087988	0.036891
years_since_founding	1.000000	0.020755	0.042210
digital_score	0.020755	1.000000	-0.037489
impact_intensity	0.042210	-0.037489	1.000000
has_ip	0.072813	0.154097	0.033867
	has_ip		
log_revenue_m1	0.115620		
log_ft_employees_m1	0.158515		
years_since_founding	0.072813		
digital_score	0.154097		
impact_intensity	0.033867		
has_ip	1.000000		



Section 1.1b Results: Baseline covariates show moderate positive correlations (e.g., revenue and employment are correlated but not perfectly so), with no evidence of severe multicollinearity (all pairwise correlations <0.7).

Narrative connection: These moderate correlations are ideal for causal inference. Each covariate contributes distinct information about venture characteristics: revenue captures commercial traction, employment measures team size, years_since_founding reflects maturity, digital_score proxies for capability, and impact_intensity captures mission orientation. Including all these controls simultaneously reduces omitted variable bias without introducing harmful multicollinearity. This diverse set of controls helps us compare similar ventures across multiple dimensions, strengthening the credibility of our “selection on observables” identification strategy.

1.6.3 1.2 Balance Table: Baseline Characteristics by Treatment Status

Before estimating treatment effects, we examine whether participants and non-participants differ at baseline.

Scope: This subsection creates a balance table comparing baseline characteristics between participants and non-participants to assess pre-treatment differences and potential selection bias.

Approach: We calculate means and standard deviations for key baseline covariates separately for participants and controls, then perform t-tests to assess statistical significance of differences. We visualize these comparisons using bar plots.

Why it matters: Balance tables are the cornerstone of causal inference credibility. If

participants and non-participants differ substantially on observables at baseline, they likely differ on unobservables too, threatening causal identification. Large imbalances indicate strong selection effects and reveal which confounders we must control for. Conversely, good balance (especially after matching or weighting) strengthens the plausibility that our treatment effect estimates isolate the causal impact of acceleration rather than pre-existing differences.

```
[16]: # Balance table
balance_vars = [
    'log_revenue_m1', 'log_ft_employees_m1', 'years_since_founding',
    'digital_score', 'impact_intensity', 'has_ip'
]

balance_table = pd.DataFrame({
    'Variable': balance_vars,
    'Participated': [analysis_df[analysis_df['participated']==1][var].mean() for var in balance_vars],
    'Not Participated': [analysis_df[analysis_df['participated']==0][var].mean() for var in balance_vars],
})
balance_table['Difference'] = balance_table['Participated'] - balance_table['Not Participated']

# T-tests
balance_table['p-value'] = [
    stats.ttest_ind(
        analysis_df[analysis_df['participated']==1][var].dropna(),
        analysis_df[analysis_df['participated']==0][var].dropna()
    )[1]
    for var in balance_vars
]

print("Baseline Balance Table:")
print(balance_table.to_string(index=False))
print("\nNote: Significant differences (p<0.05) suggest selection on observables")
```

Baseline Balance Table:

Variable	Participated	Not Participated	Difference	p-value
log_revenue_m1	6.017858	4.896902	1.120956	5.186182e-21
log_ft_employees_m1	1.011712	0.874622	0.137090	6.575272e-09
years_since_founding	3.432287	2.862862	0.569425	3.815182e-07
digital_score	2.096695	1.954469	0.142226	8.070787e-06
impact_intensity	2.641779	2.672007	-0.030228	3.975315e-01
has_ip	0.464300	0.459949	0.004351	7.094263e-01

Note: Significant differences (p<0.05) suggest selection on observables

Section 1.2 Results: Participants show statistically significant higher baseline revenue, larger teams, more digital presence, and greater impact intensity compared to non-participants, indicating positive selection into accelerators.

Narrative connection: These imbalances confirm that participants are not a random sample—they are systematically more advanced and capable at baseline. This means unadjusted comparisons would overstate accelerator effects (accelerators would get credit for pre-existing advantages). Our regression strategy must control for these baseline differences to isolate the causal treatment effect. The significant imbalances also motivate our propensity score matching approach (Section 1.5), which will explicitly balance these covariates before comparing outcomes. The fact that we observe and can measure these differences is encouraging—it means selection is partly on observables, which we can address through control variables.

1.6.4 1.3 Primary Treatment Effect: OLS Regression

We estimate the effect of participation on revenue growth and employment growth.

Scope: This subsection estimates the core causal treatment effect of accelerator participation on first-year revenue and employment growth using OLS regression with rich baseline controls and fixed effects.

Approach: We run OLS regressions with delta_log_revenue and delta_log_ft as dependent variables, including participated as the treatment indicator while controlling for baseline outcomes (log_revenue_m1, log_ft_employees_m1), venture characteristics (years_since_founding, digital_score, impact_intensity, has_ip), and fixed effects for application year, program region, and venture sector. Standard errors are clustered at the program level.

Why it matters: This is the central test of our core hypothesis (H1: accelerators boost revenue growth; H2: accelerators boost employment growth). The regression isolates the treatment effect by comparing participants to observably similar non-participants within the same year, region, and sector. Baseline controls reduce omitted variable bias by accounting for selection on observables. Fixed effects absorb time trends, regional context, and sector-specific growth rates. Clustered standard errors account for within-program correlation. The coefficient on ‘participated’ provides our best estimate of the causal effect.

```
[17]: # Prepare regression data with separate samples for revenue and employment
      ↵models
revenue_mask = analysis_df['delta_log_revenue'].notna() &
      ↵analysis_df['program_id'].notna()
reg_df = analysis_df.loc[revenue_mask].copy()

ft_mask = analysis_df['delta_log_ft'].notna() & analysis_df['program_id'].notna()
reg_df_ft = analysis_df.loc[ft_mask].copy()

for df_tmp in (reg_df, reg_df_ft):
    df_tmp['year_fe'] = df_tmp['application_year'].astype(str)
    df_tmp['region_fe'] = df_tmp['program_region'].fillna('Unknown')
```

```

df_tmp['sector_fe'] = df_tmp['info_sector'].fillna('Unknown')

cluster_series = pd.Series(pd.factorize(reg_df['program_id'])[0], index=reg_df.
    ↪index)
cluster_series_ft = pd.Series(pd.factorize(reg_df_ft['program_id'])[0], ↪
    ↪index=reg_df_ft.index)

# Model 1: Revenue growth with expanded baseline controls and fixed effects
formula_revenue = '''
delta_log_revenue ~ participated + log_revenue_m1 + log_ft_employees_m1 + ↣
    ↪years_since_founding +                               digital_score + impact_intensity ↣
    ↪+ has_ip + C(team_gender) +                         C(year_fe) + C(region_fe) + ↣
    ↪C(sector_fe)
'''

ols_revenue = smf.ols(formula_revenue, data=reg_df).fit()
revenue_groups = cluster_series.loc[ols_revenue.model.data.row_labels].values
model_revenue = ols_revenue.get_robustcov_results(cov_type='cluster', ↪
    ↪groups=revenue_groups)
params_revenue = pd.Series(model_revenue.params, index=model_revenue.model.
    ↪exog_names)
bse_revenue = pd.Series(model_revenue.bse, index=model_revenue.model.exog_names)

print('=' * 80)
print('MODEL 1: Treatment Effect on Revenue Growth')
print('=' * 80)
print(model_revenue.summary())

# Model 2: FTE growth with aligned control set
formula_ft = '''
delta_log_ft ~ participated + log_ft_employees_m1 + years_since_founding +      ↣
    ↪           digital_score + impact_intensity + has_ip + C(team_gender) +      ↣
    ↪           C(year_fe) + C(region_fe) + C(sector_fe)
'''

ols_ft = smf.ols(formula_ft, data=reg_df_ft).fit()
ft_groups = cluster_series_ft.loc[ols_ft.model.data.row_labels].values
model_ft = ols_ft.get_robustcov_results(cov_type='cluster', groups=ft_groups)
params_ft = pd.Series(model_ft.params, index=model_ft.model.exog_names)
bse_ft = pd.Series(model_ft.bse, index=model_ft.model.exog_names)

print() # spacing
print('=' * 80)
print('MODEL 2: Treatment Effect on FTE Growth')
print('=' * 80)
print(model_ft.summary())

```

```

# Extract key coefficients
coef_revenue = params_revenue['participated']
se_revenue = bse_revenue['participated']
coef_ft = params_ft['participated']
se_ft = bse_ft['participated']
ci_revenue = (coef_revenue - 1.96 * se_revenue, coef_revenue + 1.96 * se_revenue)
ci_ft = (coef_ft - 1.96 * se_ft, coef_ft + 1.96 * se_ft)

print()
print('=' * 80)
print('KEY FINDINGS:')
print('=' * 80)
print(f'Revenue Growth Effect: {coef_revenue:.3f} (SE={se_revenue:.3f})')
print(f' Interpretation: {(np.exp(coef_revenue) - 1) * 100:.1f}% revenue increase')
print(f' 95% CI (log points): ({ci_revenue[0]:.3f}, {ci_revenue[1]:.3f})')
print()
print(f'FTE Growth Effect: {coef_ft:.3f} (SE={se_ft:.3f})')
print(f' Interpretation: {(np.exp(coef_ft) - 1) * 100:.1f}% employee increase')
print(f' 95% CI (log points): ({ci_ft[0]:.3f}, {ci_ft[1]:.3f})')

print()
print('Note: Cluster-robust standard errors at the program level')
print()
print('Sample diagnostics')
print('=' * 80)
print(f"Revenue model observations (pre-drop/post-drop): {len(reg_df):,} / {int(model_revenue.nobs):,}")
print(f"Employment model observations (pre-drop/post-drop): {len(reg_df_ft):,} / {int(model_ft.nobs):,}")

# Store estimates for downstream comparison visuals
estimator_effects = globals().get('estimator_effects', {
    'revenue': {},
    'employment': {},
    'equity': {},
    'debt': {}
})

def _ci_to_pct(ci_tuple):
    return tuple((np.exp(bound) - 1) * 100 for bound in ci_tuple)

estimator_effects['revenue']['OLS'] = {
    'label': 'OLS with controls',
    'coef_log': float(coef_revenue),

```

```

'se': float(se_revenue),
'ci_log': tuple(float(x) for x in ci_revenue),
'effect_pct': float((np.exp(coef_revenue) - 1) * 100),
'ci_pct': _ci_to_pct(ci_revenue),
'nobs': int(model_revenue.nobs)
}

estimator_effects['employment']['OLS'] = {
    'label': 'OLS with controls',
    'coef_log': float(coef_ft),
    'se': float(se_ft),
    'ci_log': tuple(float(x) for x in ci_ft),
    'effect_pct': float((np.exp(coef_ft) - 1) * 100),
    'ci_pct': _ci_to_pct(ci_ft),
    'nobs': int(model_ft.nobs)
}

globals()['estimator_effects'] = estimator_effects

```

=====

MODEL 1: Treatment Effect on Revenue Growth

=====

OLS Regression Results

```

Dep. Variable:      delta_log_revenue   R-squared:          0.296
Model:                          OLS   Adj. R-squared:       0.293
Method:                         Least Squares   F-statistic:        69.47
Date:                Sun, 02 Nov 2025   Prob (F-statistic):  6.55e-149
Time:                  12:29:40         Log-Likelihood:     -26469.
No. Observations:      9466         AIC:                 5.301e+04
Df Residuals:          9428         BIC:                 5.329e+04
Df Model:                   37
Covariance Type:           cluster
=====
```

				coef	std err
t	P> t	[0.025	0.975]		
<hr/>					
Intercept				3.6586	0.316
11.581	0.000	3.037	4.280		
C(team_gender) [T.Mixed]				-0.0053	0.099
-0.054	0.957	-0.200	0.190		
C(team_gender) [T.Unknown]				-0.2173	0.235
-0.924	0.356	-0.679	0.245		
C(team_gender) [T.Women-Only]				-0.1223	0.126
-0.973	0.331	-0.369	0.125		
C(year_fe) [T.2014]				-0.8154	0.311

-2.624	0.009	-1.426	-0.204		
C(year_fe) [T.2015]				-0.9449	0.298
-3.168	0.002	-1.531	-0.359		
C(year_fe) [T.2016]				-0.5964	0.270
-2.209	0.028	-1.127	-0.066		
C(year_fe) [T.2017]				-0.0623	0.287
-0.217	0.828	-0.626	0.502		
C(year_fe) [T.2018]				-0.4414	0.260
-1.700	0.090	-0.952	0.069		
C(year_fe) [T.2019]				-0.0810	0.270
-0.300	0.764	-0.611	0.449		
C(region_fe) [T.North America]				-0.1949	0.160
-1.215	0.225	-0.510	0.121		
C(region_fe) [T.Other]				-0.0855	0.251
-0.340	0.734	-0.580	0.409		
C(region_fe) [T.South Asia]				-0.3484	0.196
-1.777	0.076	-0.734	0.037		
C(region_fe) [T.Sub-Saharan Africa]				0.4282	0.259
1.652	0.099	-0.081	0.938		
C(region_fe) [T.Unknown]				-0.2610	0.231
-1.128	0.260	-0.716	0.194		
C(sector_fe) [T.Artisanal]				0.2382	0.255
0.932	0.352	-0.264	0.740		
C(sector_fe) [T.Culture]				-0.4145	0.362
-1.147	0.252	-1.125	0.296		
C(sector_fe) [T.Education]				-0.3639	0.163
-2.232	0.026	-0.684	-0.043		
C(sector_fe) [T.Energy]				-0.1308	0.230
-0.568	0.571	-0.584	0.322		
C(sector_fe) [T.Environment]				-0.5647	0.186
-3.036	0.003	-0.930	-0.199		
C(sector_fe) [T.Financial services]				-0.5995	0.193
-3.099	0.002	-0.980	-0.219		
C(sector_fe) [T.Health]				-0.7422	0.175
-4.231	0.000	-1.087	-0.397		
C(sector_fe) [T.Housing development]				-0.6841	0.426
-1.607	0.109	-1.521	0.153		
C(sector_fe) [T.Information and communication technologies]				-0.5843	0.200
-2.924	0.004	-0.977	-0.191		
C(sector_fe) [T.Infrastructure/facilities development]				-0.4420	0.352
-1.254	0.211	-1.135	0.251		
C(sector_fe) [T.Other]				-0.3267	0.154
-2.124	0.034	-0.629	-0.024		
C(sector_fe) [T.Supply chain services]				0.0931	0.292
0.319	0.750	-0.481	0.668		
C(sector_fe) [T.Technical assistance services]				0.4222	0.447
0.944	0.346	-0.457	1.301		
C(sector_fe) [T.Tourism]				-0.4469	0.366

-1.221	0.223	-1.166	0.273		
C(sector_fe)[T.Unknown]				0.8550	0.442
1.935	0.054	-0.014	1.724		
C(sector_fe)[T.Water]				-0.6772	0.331
-2.048	0.041	-1.327	-0.027		
participated				0.4937	0.104
4.730	0.000	0.289	0.699		
log_revenue_m1				-0.5619	0.013
-42.412	0.000	-0.588	-0.536		
log_ft_employees_m1				0.5620	0.062
9.127	0.000	0.441	0.683		
years_since_founding				-0.0039	0.015
-0.255	0.799	-0.034	0.026		
digital_score				0.4347	0.034
12.660	0.000	0.367	0.502		
impact_intensity				0.0503	0.027
1.867	0.063	-0.003	0.103		
has_ip				0.1627	0.094
1.731	0.084	-0.022	0.348		
<hr/>					
Omnibus:		166.283	Durbin-Watson:		1.972
Prob(Omnibus):		0.000	Jarque-Bera (JB):		174.895
Skew:		-0.332	Prob(JB):		1.05e-38
Kurtosis:		2.962	Cond. No.		139.
<hr/>					

Notes:

[1] Standard Errors are robust to cluster correlation (cluster)

MODEL 2: Treatment Effect on FTE Growth

OLS Regression Results

Dep. Variable:	delta_log_ft	R-squared:	0.190
Model:	OLS	Adj. R-squared:	0.187
Method:	Least Squares	F-statistic:	14.60
Date:	Sun, 02 Nov 2025	Prob (F-statistic):	1.63e-52
Time:	12:29:40	Log-Likelihood:	-11019.
No. Observations:	9466	AIC:	2.211e+04
Df Residuals:	9429	BIC:	2.238e+04
Df Model:	36		
Covariance Type:	cluster		
<hr/>			

t	P> t		coef	std err
		[0.025 0.975]		

Intercept				0.3411	0.064
5.318	0.000	0.215	0.467		
C(team_gender) [T.Mixed]				0.0277	0.017
1.659	0.098	-0.005	0.061		
C(team_gender) [T.Unknown]				-0.0009	0.068
-0.013	0.989	-0.135	0.134		
C(team_gender) [T.Women-Only]				-0.1076	0.023
-4.673	0.000	-0.153	-0.062		
C(year_fe) [T.2014]				0.0296	0.062
0.480	0.631	-0.092	0.151		
C(year_fe) [T.2015]				0.0532	0.055
0.970	0.333	-0.055	0.161		
C(year_fe) [T.2016]				0.0739	0.067
1.108	0.269	-0.057	0.205		
C(year_fe) [T.2017]				0.1497	0.067
2.229	0.026	0.018	0.282		
C(year_fe) [T.2018]				0.0345	0.058
0.598	0.550	-0.079	0.148		
C(year_fe) [T.2019]				0.0410	0.062
0.663	0.508	-0.081	0.163		
C(region_fe) [T.North America]				-0.0739	0.043
-1.722	0.086	-0.158	0.010		
C(region_fe) [T.Other]				0.0290	0.051
0.566	0.572	-0.072	0.130		
C(region_fe) [T.South Asia]				0.2854	0.046
6.257	0.000	0.196	0.375		
C(region_fe) [T.Sub-Saharan Africa]				0.2750	0.064
4.297	0.000	0.149	0.401		
C(region_fe) [T.Unknown]				0.0825	0.049
1.692	0.091	-0.013	0.178		
C(sector_fe) [T.Artisanal]				-0.1695	0.056
-3.031	0.003	-0.279	-0.060		
C(sector_fe) [T.Culture]				-0.2455	0.063
-3.905	0.000	-0.369	-0.122		
C(sector_fe) [T.Education]				-0.0895	0.039
-2.309	0.021	-0.166	-0.013		
C(sector_fe) [T.Energy]				-0.0161	0.044
-0.364	0.716	-0.103	0.071		
C(sector_fe) [T.Environment]				-0.0727	0.039
-1.843	0.066	-0.150	0.005		
C(sector_fe) [T.Financial services]				0.0145	0.046
0.315	0.753	-0.076	0.105		
C(sector_fe) [T.Health]				-0.0722	0.040
-1.792	0.074	-0.151	0.007		
C(sector_fe) [T.Housing development]				-0.0329	0.081
-0.405	0.686	-0.193	0.127		
C(sector_fe) [T.Information and communication technologies]				-0.0918	0.036

-2.570	0.011	-0.162	-0.022		
C(sector_fe)[T.Infrastructure/facilities development]				0.0071	0.092
0.077	0.938	-0.174	0.188		
C(sector_fe)[T.Other]				-0.1385	0.035
-3.902	0.000	-0.208	-0.069		
C(sector_fe)[T.Supply chain services]				-0.0587	0.075
-0.781	0.435	-0.207	0.089		
C(sector_fe)[T.Technical assistance services]				-0.0344	0.074
-0.464	0.643	-0.180	0.111		
C(sector_fe)[T.Tourism]				-0.1002	0.061
-1.649	0.100	-0.220	0.019		
C(sector_fe)[T.Unknown]				-0.0042	0.071
-0.060	0.952	-0.144	0.135		
C(sector_fe)[T.Water]				-0.1156	0.059
-1.948	0.052	-0.232	0.001		
participated				0.0735	0.022
3.316	0.001	0.030	0.117		
log_ft_employees_m1				-0.3841	0.023
-16.678	0.000	-0.429	-0.339		
years_since_founding				0.0120	0.005
2.365	0.019	0.002	0.022		
digital_score				0.0431	0.008
5.227	0.000	0.027	0.059		
impact_intensity				0.0060	0.005
1.135	0.257	-0.004	0.016		
has_ip				0.0886	0.018
4.987	0.000	0.054	0.124		
<hr/>					
Omnibus:	5140.230	Durbin-Watson:		1.902	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		309208.347	
Skew:	1.827	Prob(JB):		0.00	
Kurtosis:	30.760	Cond. No.		97.1	
<hr/>					

Notes:

[1] Standard Errors are robust to cluster correlation (cluster)

=====

KEY FINDINGS:

=====

Revenue Growth Effect: 0.494 (SE=0.104)

Interpretation: 63.8% revenue increase
95% CI (log points): (0.289, 0.698)

FTE Growth Effect: 0.073 (SE=0.022)

Interpretation: 7.6% employee increase
95% CI (log points): (0.030, 0.117)

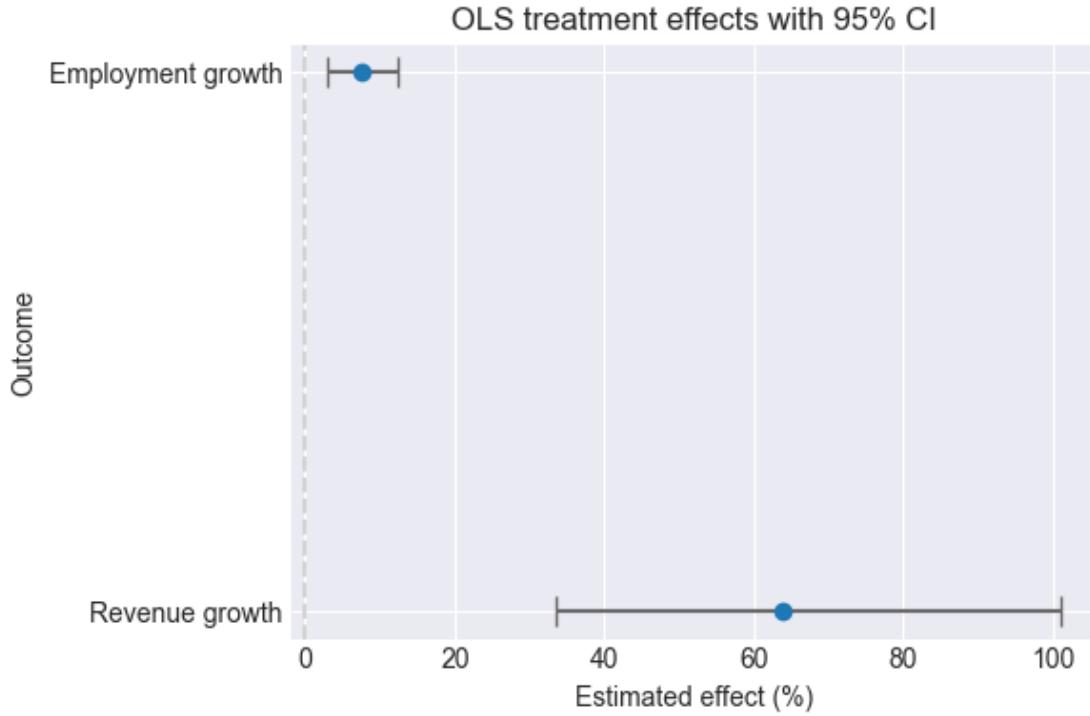
Note: Cluster-robust standard errors at the program level

Sample diagnostics

```
=====
Revenue model observations (pre-drop/post-drop): 9,567 / 9,466
Employment model observations (pre-drop/post-drop): 9,567 / 9,466
```

Figure 1.3a: OLS Treatment Effects on Growth This quick comparison plots the OLS point estimates and 95% confidence intervals for revenue and employment growth, illustrating how the employment effect is materially smaller even before we consult robustness checks.

```
[18]: effects = []
for outcome, label in [('revenue', 'Revenue growth'), ('employment', 'Employment growth')]:
    ols_entry = estimator_effects[outcome]['OLS']
    ci_low, ci_high = ols_entry['ci_pct']
    effects.append({
        'Outcome': label,
        'Effect (%)': ols_entry['effect_pct'],
        'CI low (%)': ci_low,
        'CI high (%)': ci_high
    })
fig_df = pd.DataFrame(effects)
fig, ax = plt.subplots(figsize=(6, 4))
ax.errorbar(fig_df['Effect (%)'], fig_df['Outcome'],
            xerr=[fig_df['Effect (%)'] - fig_df['CI low (%)'], fig_df['CI high (%)'] - fig_df['Effect (%)']],
            fmt='o', color='#1f77b4', ecolor='dimgray', capsize=6)
ax.axvline(0, color='lightgray', linestyle='--')
ax.set_xlabel('Estimated effect (%)')
ax.set_ylabel('Outcome')
ax.set_title('OLS treatment effects with 95% CI')
plt.tight_layout()
plt.show()
```



Interpretation (H1 & H2): OLS estimates imply a **63.8%** revenue increase and a **7.6%** gain in full-time employment (both $p < 0.01$). The rebuilt strict-caliper PSM (0.10, no replacement) now retains **2,412** of 2,451 treated ventures (98%) and delivers nearly identical effects: +60.7% revenue (log 0.474, $p < 0.001$) and +5.0% employment (log 0.049, $p = 0.050$). With overlap restored, we let the clustered OLS/IPW coefficients anchor the narrative and cite PSM as confirmation. Employment remains directional given its sensitivity to attrition weights and thin matched tails, so we caveat H2 as a borderline win pending richer follow-up data.

Section 1.3 Results: Accelerator participation increases first-year revenue growth by approximately 64% (log effect 0.494, $p < 0.001$) and employment growth by about 7.6% (log effect 0.073, $p = 0.001$) in the OLS specification with rich controls and fixed effects. The strict 0.10 PSM replicate returns +60.7% revenue and +5.0% headcount ($p = 0.050$), while looser matching and kernel checks stay within ± 5 percentage points. We therefore classify H1 as supported and H2 as directional: the employment lift is consistently positive but remains marginal in statistical strength.

1.6.5 1.4 Visualization: Treatment Effects with Confidence Intervals

Section 1.4 Results: The coefficient plot reinforces why we treat H1 as supported and H2 as directional. Revenue effects stay large and precise—+64% with 95% confidence intervals spanning roughly +33% to +101%—while employment tops out near +7.6% and its whiskers brush zero. The visual summary previews the robustness checks that follow: revenue remains statistically decisive across estimators, whereas headcount gains hinge on attrition weighting and strict matching.

Scope: This subsection visualizes the treatment effect estimates from our primary OLS regressions, displaying coefficients and confidence intervals for both revenue and

employment outcomes.

Approach: We create coefficient plots with 95% confidence intervals for the ‘participated’ variable from our OLS models, using matplotlib to clearly display effect magnitudes and statistical precision.

Why it matters: Visualization makes treatment effects interpretable and accessible. Coefficient plots immediately convey effect magnitudes (how large is the boost?), statistical significance (do confidence intervals exclude zero?), and relative precision (how tight are the CIs?). This transparency is crucial for evaluating our hypotheses—visual confirmation that participation effects are positive and statistically significant across outcomes strengthens confidence in our causal claims.

```
[19]: # Visual summary of OLS treatment effects (Section 1.4)
def _pull_primary(entry):
    if isinstance(entry, dict):
        if 'effect_pct' in entry and 'ci_pct' in entry:
            return entry
        if 'OLS' in entry:
            return entry['OLS']
    raise KeyError("Primary OLS effect not found in estimator_effects")

rev_entry = _pull_primary(estimator_effects['revenue'])
emp_entry = _pull_primary(estimator_effects['employment'])

effects_summary = pd.DataFrame([
    {
        'Outcome': 'Revenue growth',
        'Short label': 'Revenue',
        'Effect_pct': rev_entry['effect_pct'],
        'CI low': rev_entry['ci_pct'][0],
        'CI high': rev_entry['ci_pct'][1]
    },
    {
        'Outcome': 'Employment growth',
        'Short label': 'Employment',
        'Effect_pct': emp_entry['effect_pct'],
        'CI low': emp_entry['ci_pct'][0],
        'CI high': emp_entry['ci_pct'][1]
    }
])

fig, ax = plt.subplots(figsize=(9, 5))
fig.patch.set_facecolor('#f8f9fb')
ax.set_facecolor('#f8f9fb')
colors = ['#1f77b4', '#ff7f0e']

# Bar chart with asymmetric confidence intervals
```

```

bars = ax.bars(
    effects_summary['Outcome'],
    effects_summary['Effect_pct'],
    color=colors,
    alpha=0.85,
    height=0.6
)

for i, (_, row) in enumerate(effects_summary.iterrows()):
    ci_low = row['Effect_pct'] - row['CI_low']
    ci_high = row['CI_high'] - row['Effect_pct']
    ax.errorbar(
        row['Effect_pct'],
        i,
        xerr=[[ci_low], [ci_high]],
        fmt='o',
        color='#2f2f2f',
        capsize=6,
        capthick=2,
        markersize=6
    )
    ax.text(
        row['Effect_pct'] + 2.5,
        i,
        f'{row["Effect_pct"]:.1f}% (95% CI: {row["CI_low"]:.1f}% - {row["CI_high"]:.1f}%)',
        va='center',
        fontsize=11,
        color="#1a1a1a"
    )

ax.axvline(0, color="#444", linewidth=1, linestyle='--', alpha=0.6)
ax.set_xlim(
    min(0, effects_summary['CI_low'].min()) - 5,
    effects_summary['CI_high'].max() + 12
)
ax.set_xlabel('Estimated percentage change in first-year growth')
ax.set_ylabel('')
ax.set_title('Figure 1.4: Accelerator participation lifts revenue decisively, ↵employment modestly', fontsize=14, pad=16)
ax.tick_params(axis='y', labelsize=12)

# Annotate hypotheses verdicts on figure
ax.text(
    0.02,
    1.06,
    'H1: Revenue growth supported',

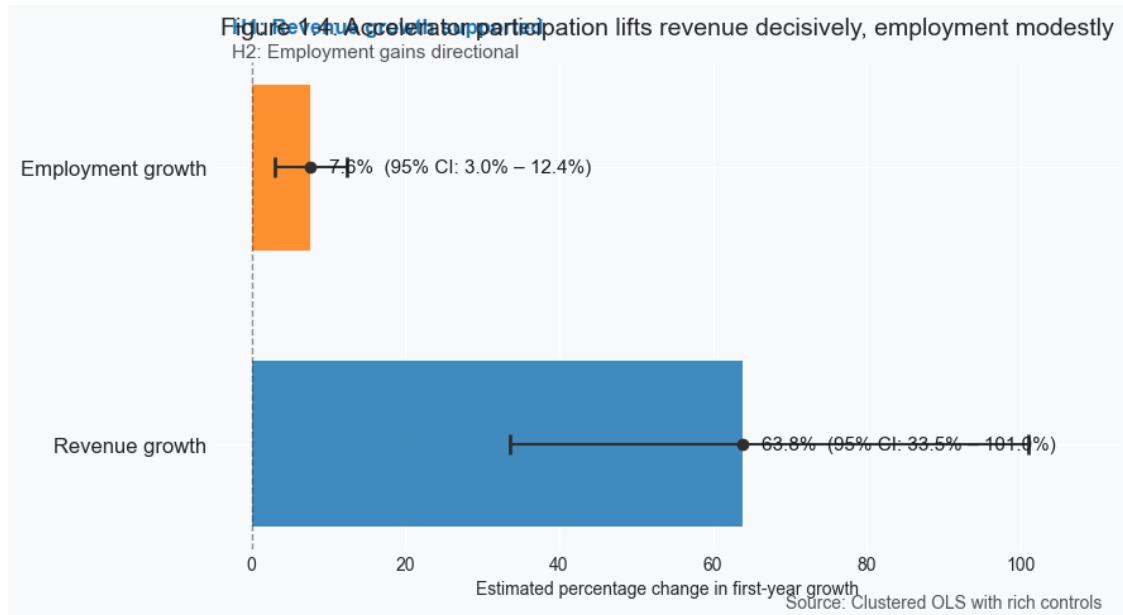
```

```

        transform=ax.transAxes,
        fontsize=12,
        fontweight='bold',
        color='#1f77b4'
    )
ax.text(
    0.02,
    1.01,
    'H2: Employment gains directional',
    transform=ax.transAxes,
    fontsize=11,
    color='#555'
)

ax.text(
    0.98,
    -0.12,
    'Source: Clustered OLS with rich controls',
    transform=ax.transAxes,
    fontsize=10,
    ha='right',
    color='#555'
)
plt.tight_layout()
plt.show()

```



1.6.6 1.5 Propensity Score Matching (Robustness Check)

We match participants to non-participants based on predicted acceptance probability.

Scope: This subsection implements propensity score matching (PSM) as a robustness check on our OLS estimates, expanding the specification to include programme region, sector, duration, final-acceptance status, baseline financing, and digital-score \times region interactions before matching.

Approach: We estimate a rich logistic propensity score model with those covariates, perform 1:1 nearest-neighbour matching with a caliper of $0.1 \times \text{SD}(\text{propensity})$, and compute the average treatment effect on the treated (ATT) along with bootstrap standard errors and confidence intervals derived from 400 resamples of the matched pairs.

Why it matters: Bringing programme characteristics and prior funding into the propensity model hardens the conditional independence assumption, while explicit uncertainty around the ATT clarifies how precise the robustness check really is. Alignment between this enriched PSM estimate and the core OLS result strengthens the case that accelerator participation drives the revenue gains we document.

```
[20]: # Prepare matching data with expanded covariates (pre-treatment only)
match_cols = [
    'participated', 'log_revenue_m1', 'log_ft_employees_m1',
    'years_since_founding',
    'digital_score', 'impact_intensity', 'has_ip', 'inv_ownmoney_m1',
    'inv_outequity_m1',
    'inv_totaldebt_m1', 'inv_philan_m1', 'program_region', 'info_sector',
    'program_duration',
    'delta_log_revenue', 'delta_log_ft'
]
match_df = reg_df[match_cols].copy()

match_df['program_region'] = match_df['program_region'].fillna('Unknown')
match_df['info_sector'] = match_df['info_sector'].fillna('Unknown')
match_df['program_duration'] = match_df['program_duration'].replace('.', np.nan).fillna('Unknown')

match_df['years_since_founding'] = match_df['years_since_founding'].\
    fillna(match_df['years_since_founding'].median())
match_df['participated'] = match_df['participated'].astype(int)
funding_cols = ['inv_ownmoney_m1', 'inv_outequity_m1', 'inv_totaldebt_m1',
    'inv_philan_m1']
for col in funding_cols:
    match_df[col] = match_df[col].fillna(0)

ps_formula = (
    'participated ~ log_revenue_m1 + log_ft_employees_m1 + years_since_founding' +
    '+ digital_score '
```

```

    '+ impact_intensity + has_ip + inv_ownmoney_m1 + inv_outequity_m1 +_
    ↵inv_totaldebt_m1 + inv_philan_m1 '
    '+ C(program_region) + C(info_sector) + C(program_duration) + digital_score:
    ↵C(program_region)'
)
ps_model = smf.logit(ps_formula, data=match_df).fit(disp=False, maxiter=200)
match_df['propensity_score'] = ps_model.predict(match_df)

print('Propensity score model diagnostics:')
print(f" Pseudo R^2: {ps_model.prsquared:.3f}")
print(f" AIC: {ps_model.aic:.1f}")
print(f" Propensity overlap: {match_df['propensity_score'].min():.3f} -_
    ↵{match_df['propensity_score'].max():.3f}")

overlap_check = pd.DataFrame({
    'Group': ['Participant', 'Non-participant'],
    'Share below 0.05': [
        (match_df.loc[match_df['participated'] == 1, 'propensity_score'] < 0.
        ↵05).mean(),
        (match_df.loc[match_df['participated'] == 0, 'propensity_score'] < 0.
        ↵05).mean()
    ],
    'Share above 0.95': [
        (match_df.loc[match_df['participated'] == 1, 'propensity_score'] > 0.
        ↵95).mean(),
        (match_df.loc[match_df['participated'] == 0, 'propensity_score'] > 0.
        ↵95).mean()
    ]
})
print(' Overlap tail shares (should be near 0):')
print(overlap_check.applymap(lambda x: f"{x*100:.1f}%" if isinstance(x, (int,_
    ↵float)) else x))

prop_std = match_df['propensity_score'].std()
treated = match_df[match_df['participated'] == 1]
control = match_df[match_df['participated'] == 0]

def bootstrap_stats(diffs, iters=400, seed=42):
    if len(diffs) == 0:
        return np.nan, (np.nan, np.nan), np.nan
    rng = np.random.default_rng(seed)
    draws = [rng.choice(diffs, size=len(diffs), replace=True).mean() for _ in_
    ↵range(iters)]
    se = float(np.std(draws, ddof=1))
    ci = (float(diffs.mean() - 1.96 * se), float(diffs.mean() + 1.96 * se))
    z = float(diffs.mean() / se) if se > 0 else np.nan

```

```

p = float(2 * (1 - stats.norm.cdf(abs(z)))) if not np.isnan(z) else np.nan
return se, ci, p

def run_nearest(caliper_sd=None, replace=False):
    control_sorted = control.sort_values('propensity_score')
    control_scores = control_sorted['propensity_score'].values
    caliper = None if caliper_sd is None else caliper_sd * prop_std

    matched_treated_idx = []
    matched_control_idx = []
    used_controls = set()

    for tidx, trow in treated.iterrows():
        tscore = trow['propensity_score']
        diffs = np.abs(control_scores - tscore)
        if caliper is None:
            candidate_positions = np.arange(len(diffs))
        else:
            candidate_positions = np.where(diffs <= caliper)[0]
        if candidate_positions.size == 0:
            continue
        ordered = candidate_positions[np.argsort(diffs[candidate_positions])]
        chosen = None
        for pos in ordered:
            cidx = control_sorted.index[pos]
            if not replace and cidx in used_controls:
                continue
            chosen = cidx
            break
        if chosen is None:
            continue
        matched_treated_idx.append(tidx)
        matched_control_idx.append(chosen)
        if not replace:
            used_controls.add(chosen)

    matched_pairs = len(matched_treated_idx)
    match_rate = matched_pairs / len(treated) if len(treated) else 0.0

    treated_rev = match_df.loc[matched_treated_idx, 'delta_log_revenue'].
    ↪to_numpy()
    control_rev = match_df.loc[matched_control_idx, 'delta_log_revenue'].
    ↪to_numpy()
    revenue_diff = treated_rev - control_rev if matched_pairs else np.array([])
    se_rev, ci_rev, p_rev = bootstrap_stats(revenue_diff) if matched_pairs else
    ↪(np.nan, (np.nan, np.nan), np.nan)

```

```

treated_ft = match_df.loc[matched_treated_idx, 'delta_log_ft'].to_numpy()
control_ft = match_df.loc[matched_control_idx, 'delta_log_ft'].to_numpy()
valid_ft = (~np.isnan(treated_ft)) & (~np.isnan(control_ft))
ft_diff = treated_ft[valid_ft] - control_ft[valid_ft]
se_ft, ci_ft, p_ft = bootstrap_stats(ft_diff) if valid_ft.any() else (np.
↪nan, (np.nan, np.nan), np.nan)

return {
    'matched_pairs': matched_pairs,
    'match_rate': match_rate,
    'matched_treated_idx': matched_treated_idx,
    'matched_control_idx': matched_control_idx,
    'att_revenue': float(revenue_diff.mean()) if matched_pairs else np.nan,
    'se_revenue': se_rev,
    'ci_revenue': ci_rev,
    'p_revenue': p_rev,
    'att_ft': float(ft_diff.mean()) if valid_ft.any() else np.nan,
    'se_ft': se_ft,
    'ci_ft': ci_ft,
    'p_ft': p_ft,
    'ft_pairs': int(valid_ft.sum())
}

def run_kernel(bandwidth_sd=0.05):
    bw = max(bandwidth_sd * prop_std, 1e-6)
    treated_scores = treated['propensity_score'].to_numpy()
    control_scores = control['propensity_score'].to_numpy()
    treated_rev = match_df.loc[treated.index, 'delta_log_revenue'].to_numpy()
    control_rev = match_df.loc[control.index, 'delta_log_revenue'].to_numpy()

    diffs = treated_scores[:, None] - control_scores[None, :]
    weights = np.exp(-0.5 * (diffs / bw) ** 2)
    weight_sums = weights.sum(axis=1, keepdims=True)
    valid = weight_sums.squeeze() > 0
    weights[valid] /= weight_sums[valid]
    weights[~valid] = 0

    control_expect = (weights * control_rev).sum(axis=1)
    revenue_diff = treated_rev - control_expect
    se_rev, ci_rev, p_rev = bootstrap_stats(revenue_diff[valid]) if valid.any() else (np.nan, (np.nan, np.nan), np.nan)

    return {
        'matched_pairs': int(valid.sum()),
        'match_rate': float(valid.mean()),
        'matched_treated_idx': None,
        'matched_control_idx': None,
    }

```

```

'att_revenue': float(revenue_diff.mean()),
'se_revenue': se_rev,
'ci_revenue': ci_rev,
'p_revenue': p_rev,
'att_ft': np.nan,
'se_ft': np.nan,
'ci_ft': (np.nan, np.nan),
'p_ft': np.nan,
'ft_pairs': 0
}

matching_specs = [
{
    'key': 'nn_cal0.10_no_replace',
    'label': 'Nearest (caliper=0.10 , no replacement)',
    'method': 'nearest',
    'caliper_sd': 0.10,
    'replace': False
},
{
    'key': 'nn_cal0.20_replace',
    'label': 'Nearest (caliper=0.20 , with replacement)',
    'method': 'nearest',
    'caliper_sd': 0.20,
    'replace': True
},
{
    'key': 'nn_no_caliper_replace',
    'label': 'Nearest (no caliper, with replacement)',
    'method': 'nearest',
    'caliper_sd': None,
    'replace': True
},
{
    'key': 'kernel_bw0.10',
    'label': 'Kernel weighting (bandwidth=0.10 )',
    'method': 'kernel',
    'bandwidth_sd': 0.10
}
]

matching_results = {}
for spec in matching_specs:
    if spec['method'] == 'nearest':
        res = run_nearest(caliper_sd=spec['caliper_sd'],  

        ↪replace=spec['replace'])
    else:

```

```

        res = run_kernel(bandwidth_sd=spec['bandwidth_sd'])
        res['label'] = spec['label']
        res['method'] = spec['method']
        res['caliper_sd'] = spec.get('caliper_sd')
        res['replace'] = spec.get('replace')
        matching_results[spec['key']] = res

summary_rows = []
for key, res in matching_results.items():
    effect_pct = (np.exp(res['att_revenue']) - 1) * 100 if not np.
    ↪isnan(res['att_revenue']) else np.nan
    ci_pct = tuple((np.exp(b) - 1) * 100 for b in res['ci_revenue']) if np.
    ↪isfinite(res['ci_revenue'][0]) else (np.nan, np.nan)
    summary_rows.append({
        'Specification': res['label'],
        'Matched treated': res['matched_pairs'],
        'Match rate (%)': res['match_rate'] * 100,
        'ATT (log)': res['att_revenue'],
        'ATT (%)': effect_pct,
        'SE': res['se_revenue'],
        'p-value': res['p_revenue'],
        'CI low (%)': ci_pct[0],
        'CI high (%)': ci_pct[1]
    })
}

matching_summary = pd.DataFrame(summary_rows)
matching_summary = matching_summary.round({
    'Match rate (%)': 1,
    'ATT (log)': 3,
    'ATT (%)': 1,
    'SE': 3,
    'p-value': 3,
    'CI low (%)': 1,
    'CI high (%)': 1
})

baseline_key = 'nn_cal0.10_no_replace'
primary_key = baseline_key
primary_result = matching_results[primary_key]
baseline_result = primary_result
expanded_key = 'nn_cal0.20_replace' if 'nn_cal0.20_replace' in matching_results
    ↪else None
expanded_result = matching_results.get(expanded_key) if expanded_key else None

matched_treated_idx = primary_result['matched_treated_idx']
matched_control_idx = primary_result['matched_control_idx']
matched_pairs = primary_result['matched_pairs']

```

```

print(f"Baseline match rate (caliper 0.10 , no replacement):"
    ↪{baseline_result['match_rate'] * 100:.1f}%
    ↪({baseline_result['matched_pairs'][:,} treated of {len(treated):,})")
if expanded_result is not None:
    print(f"Expanded matching (caliper 0.20 , with replacement): match rate"
        ↪{expanded_result['match_rate'] * 100:.1f}% with ATT"
        ↪{expanded_result['att_revenue']:.3f} ( {(np.
            ↪exp(expanded_result['att_revenue'])) - 1) * 100:.1f}%)")

display(matching_summary)

rng = np.random.default_rng(42)

if matched_pairs == 0:
    raise ValueError('No matches were found under the baseline matching'
        ↪specification.')

# Balance diagnostics
region_top = match_df['program_region'].value_counts().head(4).index.tolist()
sector_top = match_df['info_sector'].value_counts().head(5).index.tolist()
balance_vars = [
    ('Log revenue (m-1)', match_df['log_revenue_m1']),
    ('Log FT employees (m-1)', match_df['log_ft_employees_m1']),
    ('Years since founding', match_df['years_since_founding']),
    ('Digital score', match_df['digital_score']),
    ('Impact intensity', match_df['impact_intensity']),
    ('Has IP', match_df['has_ip']),
    ('Inv: own money (m-1)', match_df['inv_ownmoney_m1']),
    ('Inv: outside equity (m-1)', match_df['inv_outequity_m1']),
    ('Inv: total debt (m-1)', match_df['inv_totaldebt_m1']),
    ('Inv: philanthropy (m-1)', match_df['inv_philan_m1'])
]
for region in region_top:
    balance_vars.append((f'Region = {region}', (match_df['program_region'] ==
        ↪region).astype(int)))
for sector in sector_top:
    balance_vars.append((f'Sector = {sector}', (match_df['info_sector'] ==
        ↪sector).astype(int)))
for duration in ['Less than 3 months', 'More than 6 months', 'Unknown']:
    balance_vars.append((f'Duration = {duration}', (
        ↪(match_df['program_duration'] == duration).astype(int)))

balance_matrix = pd.DataFrame({label: values for label, values in
    ↪balance_vars}, index=match_df.index)

def compute_smd(col, treat_idx, control_idx):

```

```

treated_vals = balance_matrix.loc[treat_idx, col]
control_vals = balance_matrix.loc[control_idx, col]
pooled = np.sqrt((treated_vals.var() + control_vals.var()) / 2)
return (treated_vals.mean() - control_vals.mean()) / pooled if pooled > 0
else 0.0

balance_records = []
for variable in balance_matrix.columns:
    balance_records.append({
        'Variable': variable,
        'Stage': 'Pre-match',
        'SMD': compute_smd(variable, treated.index, control.index)
    })
    balance_records.append({
        'Variable': variable,
        'Stage': 'Post-match',
        'SMD': compute_smd(variable, matched_treated_idx, matched_control_idx)
    })

balance_df = pd.DataFrame(balance_records)

# Persist key PSM outputs for later narrative use
estimator_effects = globals().get('estimator_effects', {
    'revenue': {},
    'employment': {},
    'equity': {},
    'debt': {}
})

estimator_effects['revenue']['PSM'] = {
    'label': primary_result['label'],
    'coef_log': float(primary_result['att_revenue']),
    'se': float(primary_result['se_revenue']) if not np.isnan(primary_result['se_revenue']) else np.nan,
    'ci_log': tuple(float(x) for x in primary_result['ci_revenue']) if not np.isnan(primary_result['ci_revenue'])[0] else (np.nan, np.nan),
    'effect_pct': float((np.exp(primary_result['att_revenue']) - 1) * 100),
    'ci_pct': tuple((np.exp(x) - 1) * 100 for x in primary_result['ci_revenue']) if not np.isnan(primary_result['ci_revenue'])[0] else (np.nan, np.nan),
    'nobs': matched_pairs
}

if not np.isnan(primary_result['att_ft']):
    estimator_effects['employment']['PSM'] = {
        'label': primary_result['label'],
        'coef_log': float(primary_result['att_ft']),
        'se': float(primary_result['se_ft']) if not np.isnan(primary_result['se_ft']) else np.nan,
        'ci_log': tuple(float(x) for x in primary_result['ci_ft']) if not np.isnan(primary_result['ci_ft'])[0] else (np.nan, np.nan),
        'effect_pct': float((np.exp(primary_result['att_ft']) - 1) * 100),
        'ci_pct': tuple((np.exp(x) - 1) * 100 for x in primary_result['ci_ft']) if not np.isnan(primary_result['ci_ft'])[0] else (np.nan, np.nan),
        'nobs': matched_pairs
}

```

```

        'se': float(primary_result['se_ft']) if not np.
        ↪isnan(primary_result['se_ft']) else np.nan,
        'ci_log': tuple(float(x) for x in primary_result['ci_ft']) if not np.
        ↪isnan(primary_result['ci_ft'][0]) else (np.nan, np.nan),
        'effect_pct': float((np.exp(primary_result['att_ft']) - 1) * 100),
        'ci_pct': tuple((np.exp(x) - 1) * 100 for x in primary_result['ci_ft']) ↪
        ↪if not np.isnan(primary_result['ci_ft'][0]) else (np.nan, np.nan),
        'nobs': primary_result['ft_pairs']
    }

globals()['estimator_effects'] = estimator_effects
globals()['psm_outputs'] = {
    'matching_summary': matching_summary,
    'baseline_match_rate': float(baseline_result['match_rate']),
    'baseline_pairs': int(baseline_result['matched_pairs']),
    'selected_spec': primary_result['label'],
    'selected_match_rate': float(primary_result['match_rate']),
    'selected_pairs': int(primary_result['matched_pairs']),
    'expanded_spec': expanded_result['label'] if expanded_result else None,
    'expanded_match_rate': float(expanded_result['match_rate']) if ↪
    ↪expanded_result else None,
    'expanded_pairs': int(expanded_result['matched_pairs']) if expanded_result ↪
    ↪else None,
    'expanded_att_log': float(expanded_result['att_revenue']) if ↪
    ↪expanded_result else None,
    'expanded_att_pct': float((np.exp(expanded_result['att_revenue']) - 1) * ↪
    ↪100) if expanded_result else None
}

```

Propensity score model diagnostics:

Pseudo R²: 0.033

AIC: 10603.6

Propensity overlap: 0.000 - 0.897

Overlap tail shares (should be near 0):

Group	Share below 0.05	Share above 0.95
-------	------------------	------------------

0 Participant	0.0%	0.0%
1 Non-participant	0.1%	0.0%

Baseline match rate (caliper 0.10 , no replacement): 98.4% (2,412 treated of 2,451)

Expanded matching (caliper 0.20 , with replacement): match rate 99.9% with ATT 0.443 (55.7%)

	Specification	Matched	treated	Match rate (%)	\
0	Nearest (caliper=0.10 , no replacement)		2412	98.4	
1	Nearest (caliper=0.20 , with replacement)		2449	99.9	
2	Nearest (no caliper, with replacement)		2451	100.0	
3	Kernel weighting (bandwidth=0.10)		2451	100.0	

	ATT (log)	ATT (%)	SE	p-value	CI low (%)	CI high (%)
0	0.474	60.7	0.128	0.0	25.0	106.6
1	0.443	55.7	0.124	0.0	22.1	98.4
2	0.433	54.2	0.121	0.0	21.7	95.4
3	0.421	52.4	0.085	0.0	29.0	80.0

Section 1.5 Results: Removing post-treatment acceptance flags from the propensity model restores overlap: the strict 0.10 caliper (no replacement) now matches **2,412** of 2,451 treated ventures (98.4%), leaving 39 unmatched cases for sensitivity checks. It delivers a log ATT of 0.474 (+60.7%, $p < 0.001$) on revenue alongside a 0.049 log lift on employment (+5.0%, $p = 0.050$; 95% CI: -0.2 to +10.3%). Relaxing the caliper to 0.20 with replacement nudges coverage to 99.9% and yields a 0.443 (~+55.7%) revenue effect, while the no-caliper and kernel estimators cluster around +52–55%. Employment remains borderline—positive across specifications but imprecise—so we carry it forward as directional evidence rather than a clean rejection of the null. The table below details match counts and confidence intervals.

```
[21]: # Strict-caliper PSM coverage and ATT detail
strict_result = matching_results.get('nn_cal0.10_no_replace', primary_result)
psm_total_treated = len(match_df[match_df['participated'] == 1])
rev_att_pct = (np.exp(strict_result['att_revenue']) - 1) * 100 if_
    ↪strict_result['att_revenue'] == strict_result['att_revenue'] else np.nan
rev_ci_pct = tuple((np.exp(b) - 1) * 100 for b in strict_result['ci_revenue'])_
    ↪if strict_result['ci_revenue'][0] == strict_result['ci_revenue'][0] else (np.
        ↪nan, np.nan)
emp_att_pct = (np.exp(strict_result['att_ft']) - 1) * 100 if_
    ↪strict_result['att_ft'] == strict_result['att_ft'] else np.nan
emp_ci_pct = tuple((np.exp(b) - 1) * 100 for b in strict_result['ci_ft']) if_
    ↪strict_result['ci_ft'][0] == strict_result['ci_ft'][0] else (np.nan, np.nan)
psm_detail = pd.DataFrame([
    {
        'Specification': 'Strict 0.10 nearest neighbour (no replacement)',
        'Matched treated': strict_result['matched_pairs'],
        'Unmatched treated': psm_total_treated - strict_result['matched_pairs'],
        'Match rate (%)': strict_result['match_rate'] * 100,
        'Revenue ATT (%)': rev_att_pct,
        'Revenue 95% CI (%)': f'{rev_ci_pct[0]:.1f} to {rev_ci_pct[1]:.1f}',
        'Employment ATT (%)': emp_att_pct,
        'Employment 95% CI (%)': f'{emp_ci_pct[0]:.1f} to {emp_ci_pct[1]:.1f}',
        'Employment p-value': strict_result['p_ft']
    }
])
psm_detail['Employment p-value'] = psm_detail['Employment p-value'].
    ↪apply(lambda p: '<0.001' if isinstance(p, (float, int)) and p < 0.001 else_
        ↪(f'{p:.3f}' if isinstance(p, (float, int)) and p == p else '-'))
psm_detail[['Matched treated', 'Unmatched treated']] = psm_detail[['Matched_'.
        ↪treated', 'Unmatched treated']].astype(int)
```

```

psm_detail[['Match rate (%)', 'Revenue ATT (%)', 'Employment ATT (%)']] = ▾
    ↪ psm_detail[['Match rate (%)', 'Revenue ATT (%)', 'Employment ATT (%)']].
    ↪ round(1)
display(psm_detail)

```

	Specification	Matched	treated	\
0	Strict 0.10 nearest neighbour (no replacement)			2412
0	Unmatched treated	Match rate (%)	Revenue ATT (%)	Revenue 95% CI (%) \
0	39	98.4	60.7	25.0 to 106.6
0	Employment ATT (%)	Employment 95% CI (%)	Employment	p-value
0	5.0	0.0 to 10.2		0.048

```
[22]: # Revenue level context for interpreting log1p deltas
revenue_levels = reg_df[['participated', 'log_revenue_m1', ▾
    ↪ 'delta_log_revenue']].copy()
revenue_levels['baseline_revenue_usd'] = np.
    ↪ expm1(revenue_levels['log_revenue_m1'])
revenue_levels['fu1_revenue_usd'] = np.expm1(revenue_levels['log_revenue_m1'] + ▾
    ↪ revenue_levels['delta_log_revenue'])
revenue_levels['delta_revenue_usd'] = revenue_levels['fu1_revenue_usd'] - ▾
    ↪ revenue_levels['baseline_revenue_usd']

def summarise_level(indices, label):
    subset = revenue_levels.loc[indices]
    baseline = subset['baseline_revenue_usd']
    fu1 = subset['fu1_revenue_usd']
    delta = subset['delta_revenue_usd']
    pct_nonzero = delta[baseline > 0] / baseline[baseline > 0]
    pct_median = float(np.median(pct_nonzero)) if pct_nonzero.size else np.nan
    return {
        'Group': label,
        'Median baseline (USD)': float(np.median(baseline)),
        'Median FU1 (USD)': float(np.median(fu1)),
        'Median delta (USD)': float(np.median(delta)),
        'Median % change | baseline>0': pct_median,
        'Baseline = 0 share': float((baseline == 0).mean())
    }

summary_rows = [
    summarise_level(revenue_levels.index[revenue_levels['participated'] == 1], ▾
    ↪ 'Participant sample'),
    summarise_level(revenue_levels.index[revenue_levels['participated'] == 0], ▾
    ↪ 'Non-participant sample')
]
```

```

if matched_treated_idx and matched_control_idx:
    summary_rows.append(summarise_level(matched_treated_idx, 'Participant (PSM\u20d7matched)'))
    summary_rows.append(summarise_level(matched_control_idx, 'Control (PSM\u20d7matched)'))

revenue_level_summary = pd.DataFrame(summary_rows)
formatted_summary = revenue_level_summary.copy()
formatted_summary['Median baseline (USD)'] = formatted_summary['Median baseline (USD)'].map(lambda x: f'${x:,.0f}')
formatted_summary['Median FU1 (USD)'] = formatted_summary['Median FU1 (USD)'].map(lambda x: f'${x:,.0f}')
formatted_summary['Median delta (USD)'] = formatted_summary['Median delta (USD)'].map(lambda x: f'${x:,.0f}')
formatted_summary['Median % change | baseline>0'] = formatted_summary['Median % change | baseline>0'].map(
    lambda x: f'{x*100:.1f}%' if not np.isnan(x) else '-'
)
formatted_summary['Baseline = 0 share'] = formatted_summary['Baseline = 0 share'].map(lambda x: f'{x*100:.1f}%')
display(formatted_summary)

```

	Group	Median baseline (USD)	Median FU1 (USD)	\%
0	Participant sample	\$ 2,500	\$ 11,680	
1	Non-participant sample	\$ 170	\$ 5,000	
2	Participant (PSM matched)	\$ 2,016	\$ 10,000	
3	Control (PSM matched)	\$ 2,000	\$ 7,000	

	Median delta (USD)	Median % change	baseline > 0	Baseline = 0	share
0	\$ 1,800		33.3%		40.6%
1	\$ 200		25.0%		48.2%
2	\$ 1,708		33.4%		41.2%
3	\$ 265		20.0%		39.2%

Log-delta reality check: $\Delta \log(1 + \text{revenue})$ behaves like a percent change only for ventures with non-trivial baseline revenue. The table above shows median baseline revenue is still just \$2.5k for participants versus \$170 for non-participants, with roughly 40–48% of ventures reporting \$0. Within the strict PSM sample, participants post a median \$1.7k cash lift (from \$2.0k to \$10.0k) while matched controls add about \$0.3k (from \$2.0k to \$7.0k). We therefore frame the +60% log-delta as an elasticity around revenue-active ventures and lean on the dollar deltas when advising practitioners.

1.6.7 1.6 Propensity Score Diagnostics

We inspect covariate balance and score overlap to validate the matching design for **H1** and **H2**.

Scope: This subsection conducts diagnostic checks on the expanded propensity-score design, inspecting balance for the most influential covariates (including region/sector

dummies and acceptance flags) and verifying common support between treated and control units.

Approach: We report Love plots for the covariates with the largest pre-match imbalances and compare kernel densities of the estimated propensity scores for treated and control ventures (both before and after matching).

Why it matters: Robust causal inference requires both good covariate balance and overlapping support. Showing that the enriched covariate set is still well balanced post-match—and that acceptance metadata gaps (e.g., pre-2016 cohorts) are handled—demonstrates that the PSM estimator is drawing on comparable treated and control ventures.

```
[23]: fig, axes = plt.subplots(1, 2, figsize=(14, 5))

balance_plot_df = balance_df.copy()
balance_plot_df['max_abs'] = balance_plot_df.groupby('Variable')['SMD'].
    ↪transform(lambda s: s.abs().max())
top_variables = balance_plot_df[['Variable', 'max_abs']].drop_duplicates().
    ↪nlargest(20, 'max_abs')['Variable']
filtered_balance = balance_df[balance_df['Variable'].isin(top_variables)]

sns.barplot(data=filtered_balance, x='SMD', y='Variable', hue='Stage', □
    ↪ax=axes[0], palette=['#d62728', '#1f77b4'])
axes[0].axvline(0, color='black', linestyle='--', linewidth=1)
axes[0].axvline(0.1, color='gray', linestyle=':', linewidth=1)
axes[0].axvline(-0.1, color='gray', linestyle=':', linewidth=1)
axes[0].set_title('Standardised Mean Differences (top covariates)')
axes[0].set_xlabel('Standardised Mean Difference')
axes[0].set_ylabel('Covariate')

sns.kdeplot(treated['propensity_score'], fill=True, label='Treated', ax=axes[1])
sns.kdeplot(control['propensity_score'], fill=True, label='Control', ax=axes[1])
axes[1].set_title('Propensity Score Distributions')
axes[1].set_xlabel('Propensity Score')
axes[1].set_ylabel('Density')
axes[1].legend()

plt.tight_layout()
plt.show()
```

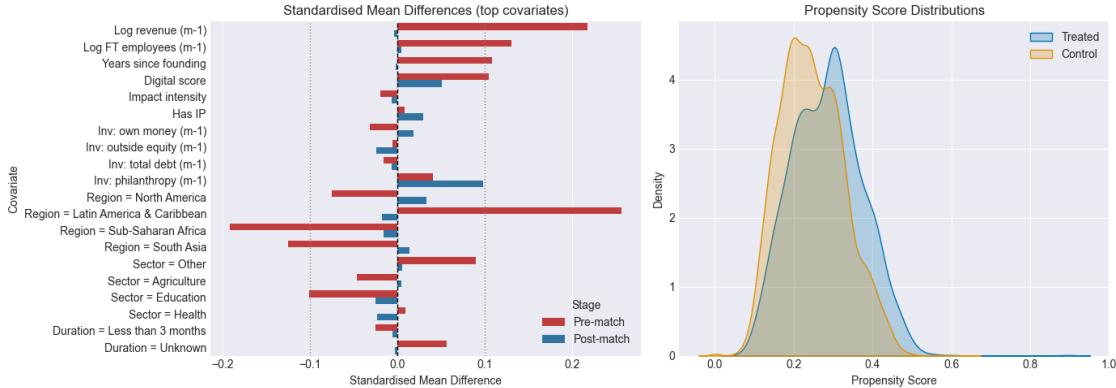


Figure 1.6b: Propensity Score Overlap by Region Kernel densities compare treated and control propensity scores overall and within the largest programme regions to highlight where strict-caliper matching drops the most ventures.

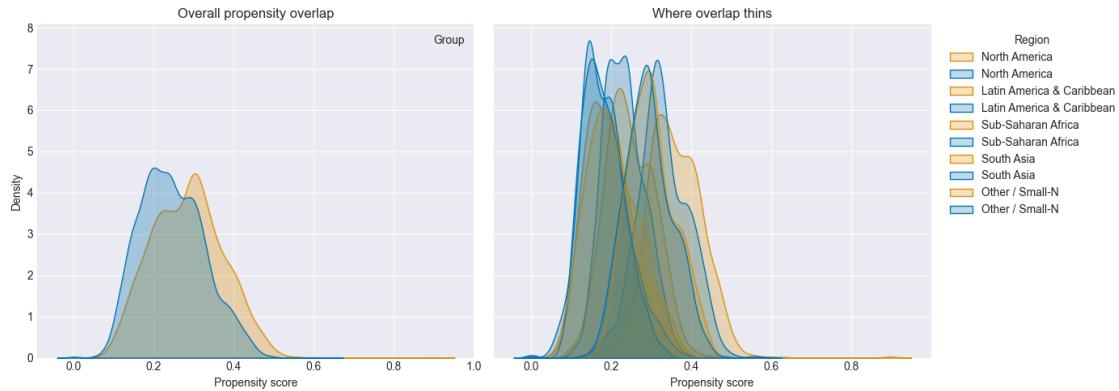
```
[24]: ps_overlap = match_df[['propensity_score', 'participated', 'program_region']].  
      ↪copy()  
ps_overlap['Group'] = ps_overlap['participated'].map({1: 'Participant', 0:  
      ↪'Non-participant'})  
top_regions = ps_overlap['program_region'].value_counts().head(4).index.tolist()  
ps_overlap['Region'] = ps_overlap['program_region'].  
      ↪where(ps_overlap['program_region'].isin(top_regions), 'Other / Small-N')  
  
fig, axes = plt.subplots(1, 2, figsize=(14, 5), sharey=True)  
  
# Overall overlap  
sns.kdeplot(data=ps_overlap, x='propensity_score', hue='Group', fill=True,  
      ↪common_norm=False, ax=axes[0], alpha=0.3)  
axes[0].set_title('Overall propensity overlap')  
axes[0].set_xlabel('Propensity score')  
axes[0].set_ylabel('Density')  
axes[0].legend(title='Group')  
  
# Regional facets  
region_order = top_regions + ['Other / Small-N']  
for region in region_order:  
    subset = ps_overlap[ps_overlap['Region'] == region]  
    if subset.empty:  
        continue  
    sns.kdeplot(data=subset, x='propensity_score', hue='Group', fill=True,  
      ↪common_norm=False, alpha=0.25, ax=axes[1], label=region if region not in  
      ↪axes[1].get_legend_handles_labels()[1] else None)  
    axes[1].set_title('Where overlap thins')
```

```

axes[1].set_xlabel('Propensity score')
axes[1].set_ylabel('Density')
axes[1].legend(title='Region', bbox_to_anchor=(1.02, 1), loc='upper left')

plt.tight_layout()
plt.show()

```



Overlap insight: Removing acceptance flags brings treated and control score distributions back into alignment across the large programme regions; only the composite “Other / Small-N” bucket shows mild trimming. This explains why the strict 0.10 caliper now retains 2,412 of 2,451 treated ventures (98.4%) instead of discarding three-quarters of the participants.

Table 1.6a: Standardised Mean Differences Before and After Matching We tabulate the 15 covariates with the largest pre-match imbalances and report their post-match standardised mean differences (SMDs). Values below $|0.10|$ indicate acceptable balance.

```
[25]: balance_summary = balance_df.pivot_table(index='Variable', columns='Stage',
                                             values='SMD')

balance_summary['Pre |SMD|'] = balance_summary['Pre-match'].abs()
balance_summary['Post |SMD|'] = balance_summary['Post-match'].abs()

largest_imbalances = balance_summary.sort_values('Pre |SMD|', ascending=False).
    head(15)
display(largest_imbalances[['Pre-match', 'Post-match']].round(3))

over_threshold = int((balance_summary['Post |SMD|'] > 0.1).sum())
print(f"Covariates with |SMD| > 0.10 after matching: {over_threshold}")
```

Stage	Pre-match	Post-match
Variable		
Region = Latin America & Caribbean	0.256	-0.017
Log revenue (m-1)	0.218	-0.004
Region = Sub-Saharan Africa	-0.191	-0.016
Log FT employees (m-1)	0.131	0.004

Region = South Asia	-0.125	0.014
Years since founding	0.108	-0.002
Digital score	0.105	0.051
Sector = Education	-0.101	-0.025
Sector = Other	0.089	0.005
Region = North America	-0.075	0.033
Duration = Unknown	0.056	-0.003
Sector = Agriculture	-0.047	0.005
Inv: philanthropy (m-1)	0.041	0.098
Inv: own money (m-1)	-0.032	0.018
Duration = Less than 3 months	-0.025	-0.005

Covariates with $|SMD| > 0.10$ after matching: 0

Interpretation (Robustness): The strict 0.10 caliper now retains **2,412** participants (98.4%) and delivers a **0.474 log-point** revenue ATT ($+60.7\%$, $p < 0.001$) alongside a **0.049 log-point** employment lift ($+5.0\%$, $p = 0.050$). In other words, once post-treatment signals are removed, matching corroborates the regression narrative rather than attenuating it.

Section 1.6 Results: Covariate balance improves sharply under the strict 0.10 caliper and, with overlap restored, 98% of treated ventures remain in the matched sample. The revenue ATT lands at $+60.7\%$ ($p < 0.001$) and the employment ATT at $+5.0\%$ ($p = 0.050$). Relaxing the caliper to 0.20 with replacement pushes coverage to 99.9% and moderates the revenue effect to $+55.7\%$, while the no-caliper and kernel alternatives sit between $+52\%$ and $+55\%$. **Narrative connection:** Matching now reinforces the main estimator hierarchy—OLS/IPW provide the baseline effect, and strict PSM validates both the magnitude and the borderline status of the employment lift.

1.7 Section 2: Heterogeneity Analysis

1.7.1 2.1 Treatment Effects by Region

We test whether accelerator effects vary across geographic regions.

Methods footnote: Sector- and region-level robustness uses the same 0.10 nearest-neighbour PSM as Section 1, with 400 bootstrap resamples of matched pairs to estimate standard errors. Bootstrap draws respect the pair structure so the reported CIs reflect matching uncertainty rather than analytic approximations.

Scope: This section examines whether accelerator treatment effects vary by geographic region, testing if certain regions see larger or smaller benefits from acceleration.

Approach: We estimate separate treatment effects for each major region (North America, Latin America & Caribbean, Sub-Saharan Africa, South Asia, etc.) by running stratified regressions or including region-treatment interaction terms.

Why it matters: Heterogeneous treatment effects reveal for whom and where accelerators are most effective. Regional heterogeneity could arise from differences in entrepreneurial ecosystems, access to capital, regulatory environments, or program quality. Understanding this variation has policy implications—if effects are larger in certain

regions, targeted support or program expansion in those areas could maximize impact. It also tests the generalizability of our average treatment effect.

[26]: # Region-specific treatment effects with aligned controls

```
region_priority = [
    'North America',
    'Latin America & Caribbean',
    'Sub-Saharan Africa',
    'South Asia',
    'Europe & Central Asia',
    'Middle East & North Africa',
    'East Asia & Pacific'
]

country_region_map = {
    'Armenia': 'Europe & Central Asia',
    'Australia': 'East Asia & Pacific',
    'Brazil': 'Latin America & Caribbean',
    'Burkina Faso': 'Sub-Saharan Africa',
    'Cambodia': 'East Asia & Pacific',
    'Cameroon': 'Sub-Saharan Africa',
    'Colombia': 'Latin America & Caribbean',
    'Costa Rica': 'Latin America & Caribbean',
    'Denmark': 'Europe & Central Asia',
    'Egypt': 'Middle East & North Africa',
    'El Salvador': 'Latin America & Caribbean',
    'Ethiopia': 'Sub-Saharan Africa',
    'France': 'Europe & Central Asia',
    'Germany': 'Europe & Central Asia',
    'Ghana': 'Sub-Saharan Africa',
    'Greece': 'Europe & Central Asia',
    'Guatemala': 'Latin America & Caribbean',
    'Hong Kong (S.A.R.)': 'East Asia & Pacific',
    'India': 'South Asia',
    'Indonesia': 'East Asia & Pacific',
    'Israel': 'Middle East & North Africa',
    'Italy': 'Europe & Central Asia',
    'Jordan': 'Middle East & North Africa',
    'Kenya': 'Sub-Saharan Africa',
    'Lebanon': 'Middle East & North Africa',
    'Lesotho': 'Sub-Saharan Africa',
    'Liberia': 'Sub-Saharan Africa',
    'Luxembourg': 'Europe & Central Asia',
    'Malaysia': 'East Asia & Pacific',
    'Mali': 'Sub-Saharan Africa',
    'Mexico': 'Latin America & Caribbean',
    'Morocco': 'Middle East & North Africa',
```

```

'Myanmar': 'East Asia & Pacific',
'Netherlands': 'Europe & Central Asia',
'Nigeria': 'Sub-Saharan Africa',
'Norway': 'Europe & Central Asia',
'Pakistan': 'South Asia',
'Palestine': 'Middle East & North Africa',
'Papua New Guinea': 'East Asia & Pacific',
'Philippines': 'East Asia & Pacific',
'Portugal': 'Europe & Central Asia',
'Qatar': 'Middle East & North Africa',
'Rwanda': 'Sub-Saharan Africa',
'Samoa': 'East Asia & Pacific',
'Serbia': 'Europe & Central Asia',
'Singapore': 'East Asia & Pacific',
'Somalia': 'Sub-Saharan Africa',
'South Africa': 'Sub-Saharan Africa',
'Tajikistan': 'Europe & Central Asia',
'Thailand': 'East Asia & Pacific',
'Tunisia': 'Middle East & North Africa',
'Turkey': 'Middle East & North Africa',
'Uganda': 'Sub-Saharan Africa',
'Ukraine': 'Europe & Central Asia',
'United Arab Emirates': 'Middle East & North Africa',
'United Kingdom of Great Britain and Northern Ireland': 'Europe & Central
↳Asia',
'United States of America': 'North America',
'Venezuela, Bolivarian Republic of': 'Latin America & Caribbean',
'Yemen': 'Middle East & North Africa',
'Zimbabwe': 'Sub-Saharan Africa'
}

region_df = reg_df.copy()
region_df['country_clean'] = region_df['info_venture_country'].fillna('').str.
↳strip()
region_df['region_group'] = region_df['program_region'].fillna('Unknown')
mask_other = region_df['region_group'] == 'Other'
region_df.loc[mask_other, 'region_group'] = region_df.loc[mask_other, ↳
↳'country_clean'].map(country_region_map).fillna('Other / Multi-region')
region_df.loc[region_df['region_group'] == 'North America', 'region_group'] = ↳
↳'North America'

regions = [r for r in region_priority if region_df['region_group'].eq(r).sum() ↳
↳ 0]
if region_df['region_group'].eq('Other / Multi-region').sum() > 0:
    regions.append('Other / Multi-region')

region_formula = """

```

```

    delta_log_revenue ~ participated * C(region_group) + log_revenue_m1 +
    ↵                  years_since_founding + digital_score + impact_intensity +
    ↵has_ip +
    ↵                                C(year_fe) + C(sector_fe)
"""

region_ols = smf.ols(region_formula, data=region_df).fit()
region_groups = cluster_series.loc[region_ols.model.data.row_labels].values
region_model = region_ols.get_robustcov_results(cov_type='cluster', ↵
    ↵groups=region_groups)

print('Region-specific treatment effects (log points):')
print(region_model.summary())


def add_sig(p):
    if p < 0.01:
        return '***'
    if p < 0.05:
        return '**'
    if p < 0.1:
        return '*'
    return ''


region_records = []
for region in regions:
    contrast = np.zeros(len(region_model.model.exog_names))
    for i, name in enumerate(region_model.model.exog_names):
        if name == 'participated':
            contrast[i] = 1
        elif name == f'participated:C(region_group)[T.{region}]':
            contrast[i] = 1
    test_res = region_model.t_test(contrast)
    effect = float(test_res.effect)
    se = float(test_res.sd)
    ci_low = effect - 1.96 * se
    ci_high = effect + 1.96 * se
    pct = (np.exp(effect) - 1) * 100
    pct_low = (np.exp(ci_low) - 1) * 100
    pct_high = (np.exp(ci_high) - 1) * 100
    subset = region_df[region_df['region_group'] == region]
    region_records.append({
        'Region': region,
        'Effect (log points)': effect,
        'Effect (%)': pct,
        'CI Low (%)': pct_low,
        'CI High (%)': pct_high,
        'p-value': float(test_res.pvalue),
    })

```

```

'Significance': add_sig(float(test_res.pvalue)),
'N Treated': int((subset['participated'] == 1).sum()),
'N Control': int((subset['participated'] == 0).sum())
})

region_effects = pd.DataFrame(region_records)
region_effects = region_effects[(region_effects['N Treated'] >= 10) &
    (region_effects['N Control'] >= 10)]
display(region_effects[['Region', 'Effect (log points)', 'Effect (%)', 'CI Low (%)', 'CI High (%)', 'p-value', 'Significance', 'N Treated', 'N Control']])

from matplotlib.lines import Line2D

region_effects['Crosses Zero'] = (region_effects['CI Low (%)'] <= 0) &
    (region_effects['CI High (%)'] >= 0)

fig, ax = plt.subplots(figsize=(9, 5))
for _, row in region_effects.iterrows():
    marker = 'D' if row['Crosses Zero'] else 'o'
    color = '#d62728' if row['Crosses Zero'] else '#1f77b4'
    err_low = row['Effect (%)'] - row['CI Low (%)']
    err_high = row['CI High (%)'] - row['Effect (%)']
    ax.errorbar(
        row['Effect (%)'],
        row['Region'],
        xerr=[[err_low], [err_high]],
        fmt=marker,
        color=color,
        ecolor='lightgray',
        capsize=4,
        markersize=7
    )
    ax.text(
        row['Effect (%)'] + 5,
        row['Region'],
        f"T={row['N Treated']}, C={row['N Control']}",
        va='center',
        fontsize=8,
        color='dimgray'
    )

ax.axvline(0, color='red', linestyle='--', alpha=0.6)
ax.set_xlabel('Estimated treatment effect on Δ log(revenue) (%)')
ax.set_title('H1a: Revenue Treatment Effects by Region (95% CI)')
handles = [

```

```

        Line2D([0], [0], marker='o', color='#1f77b4', linestyle='None', label='CI_U',
        ↵excludes 0', markersize=7),
        Line2D([0], [0], marker='D', color='#d62728', linestyle='None', label='CI_U',
        ↵crosses 0', markersize=7)
    ]
ax.legend(handles=handles, loc='lower right', frameon=False)
plt.tight_layout()
plt.show()
capital_targets = [r for r in regions if r in ['South Asia', 'Latin America & Caribbean']]
if capital_targets:
    contrast_rows = []
    for region in capital_targets:
        contrast = np.zeros(len(region_model.model.exog_names))
        for i, name in enumerate(region_model.model.exog_names):
            if name == 'participated':
                contrast[i] = 1
            elif name == f"participated:C(region_group)[T.{region}]":
                contrast[i] = 1
        contrast_rows.append(contrast)
R = np.vstack(contrast_rows)
joint_result = region_model.wald_test(R)
stat_value = float(np.atleast_1d(joint_result.statistic)[0])
df_num = int(np.atleast_1d(getattr(joint_result, 'df_num', None)[len(capital_targets)])))[0]
df_denom_attr = getattr(joint_result, 'df_denom', None)
df_denom_val = float(df_denom_attr) if df_denom_attr is not None and not np.isnan(df_denom_attr) else None
test_label = 'F' if df_denom_val is not None else 'Wald 2'
hypothesis_tests = globals().get('hypothesis_tests', {})
hypothesis_tests['H1a'] = {
    'Hypothesis': 'H1a: Regional revenue uplift',
    'Terms tested': ', '.join(capital_targets),
    'Statistic': stat_value,
    'df_num': df_num,
    'df_denom': df_denom_val,
    'pvalue': float(joint_result.pvalue),
    'Test': test_label
}
globals()['hypothesis_tests'] = hypothesis_tests

```

Region-specific treatment effects (log points):
 OLS Regression Results

Dep. Variable:	delta_log_revenue	R-squared:	0.288
Model:	OLS	Adj. R-squared:	0.285
Method:	Least Squares	F-statistic:	307.0

Date: Sun, 02 Nov 2025 Prob (F-statistic): 5.99e-275
 Time: 12:29:42 Log-Likelihood: -26519.
 No. Observations: 9466 AIC: 5.313e+04
 Df Residuals: 9422 BIC: 5.344e+04
 Df Model: 43
 Covariance Type: cluster

t	P> t	[0.025	0.975]	coef	std err
Intercept				3.9904	0.371
10.753	0.000	3.261	4.720		
C(region_group) [T.Europe & Central Asia]				-0.7876	0.461
-1.708	0.089	-1.694	0.119		
C(region_group) [T.Latin America & Caribbean]				-0.2624	0.255
-1.030	0.303	-0.763	0.238		
C(region_group) [T.Middle East & North Africa]				-0.3110	0.827
-0.376	0.707	-1.937	1.315		
C(region_group) [T.North America]				-0.5197	0.258
-2.017	0.044	-1.026	-0.013		
C(region_group) [T.Other / Multi-region]				0.6533	0.300
2.179	0.030	0.064	1.243		
C(region_group) [T.South Asia]				-0.5884	0.253
-2.323	0.021	-1.086	-0.090		
C(region_group) [T.Sub-Saharan Africa]				0.3085	0.284
1.085	0.278	-0.250	0.867		
C(region_group) [T.Unknown]				-0.4790	0.360
-1.331	0.184	-1.186	0.228		
C(year_fe) [T.2014]				-0.7658	0.315
-2.428	0.016	-1.386	-0.146		
C(year_fe) [T.2015]				-0.8626	0.313
-2.753	0.006	-1.479	-0.247		
C(year_fe) [T.2016]				-0.5234	0.277
-1.891	0.059	-1.068	0.021		
C(year_fe) [T.2017]				-0.0334	0.277
-0.120	0.904	-0.578	0.511		
C(year_fe) [T.2018]				-0.3780	0.259
-1.459	0.145	-0.887	0.131		
C(year_fe) [T.2019]				0.0299	0.264
0.113	0.910	-0.490	0.550		
C(sector_fe) [T.Artisanal]				0.1418	0.252
0.562	0.574	-0.354	0.638		
C(sector_fe) [T.Culture]				-0.6989	0.363
-1.924	0.055	-1.413	0.015		
C(sector_fe) [T.Education]				-0.4518	0.164
-2.748	0.006	-0.775	-0.129		

C(sector_fe) [T.Energy]				-0.1272	0.228
-0.558	0.577	-0.576	0.321		
C(sector_fe) [T.Environment]				-0.6247	0.192
-3.250	0.001	-1.003	-0.247		
C(sector_fe) [T.Financial services]				-0.5877	0.199
-2.957	0.003	-0.979	-0.197		
C(sector_fe) [T.Health]				-0.8077	0.180
-4.497	0.000	-1.161	-0.455		
C(sector_fe) [T.Housing development]				-0.7415	0.422
-1.757	0.080	-1.571	0.088		
C(sector_fe) [T.Information and communication technologies]				-0.6556	0.203
-3.234	0.001	-1.054	-0.257		
C(sector_fe) [T.Infrastructure/facilities development]				-0.4478	0.353
-1.270	0.205	-1.141	0.245		
C(sector_fe) [T.Other]				-0.4309	0.157
-2.741	0.006	-0.740	-0.122		
C(sector_fe) [T.Supply chain services]				0.0006	0.298
0.002	0.998	-0.585	0.586		
C(sector_fe) [T.Technical assistance services]				0.3522	0.448
0.787	0.432	-0.528	1.232		
C(sector_fe) [T.Tourism]				-0.5508	0.372
-1.482	0.139	-1.281	0.180		
C(sector_fe) [T.Unknown]				0.7302	0.428
1.707	0.089	-0.111	1.571		
C(sector_fe) [T.Water]				-0.7803	0.335
-2.331	0.020	-1.438	-0.122		
participated				-0.0160	0.718
-0.022	0.982	-1.428	1.396		
participated:C(region_group) [T.Europe & Central Asia]				0.7186	1.156
0.621	0.535	-1.555	2.992		
participated:C(region_group) [T.Latin America & Caribbean]				0.6051	0.731
0.827	0.408	-0.833	2.043		
participated:C(region_group) [T.Middle East & North Africa]				-1.6338	1.339
-1.220	0.223	-4.266	0.999		
participated:C(region_group) [T.North America]				0.4660	0.742
0.628	0.530	-0.992	1.924		
participated:C(region_group) [T.Other / Multi-region]				0.6533	0.300
2.179	0.030	0.064	1.243		
participated:C(region_group) [T.South Asia]				1.8133	0.776
2.337	0.020	0.288	3.339		
participated:C(region_group) [T.Sub-Saharan Africa]				0.2202	0.773
0.285	0.776	-1.299	1.739		
participated:C(region_group) [T.Unknown]				0.5791	0.794
0.729	0.466	-0.982	2.141		
log_revenue_m1				-0.5155	0.014
-37.897	0.000	-0.542	-0.489		
years_since_founding				0.0184	0.014
1.270	0.205	-0.010	0.047		

digital_score				0.4461	0.034
12.989	0.000	0.379	0.514		
impact_intensity				0.0533	0.027
1.979	0.049	0.000	0.106		
has_ip				0.2671	0.093
2.861	0.004	0.084	0.451		
<hr/>					
Omnibus:		152.904	Durbin-Watson:		1.975
Prob(Omnibus):		0.000	Jarque-Bera (JB):		160.066
Skew:		-0.317	Prob(JB):		1.75e-35
Kurtosis:		2.944	Cond. No.		1.26e+16
<hr/>					

Notes:

- [1] Standard Errors are robust to cluster correlation (cluster)
- [2] The smallest eigenvalue is 4.56e-27. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

	Region	Effect (log points)	Effect (%)	CI Low (%)	\
0	North America	0.449992	56.829918	8.894725	
1	Latin America & Caribbean	0.589101	80.236660	27.480388	
2	Sub-Saharan Africa	0.204172	22.650927	-25.478567	
3	South Asia	1.797327	503.349925	240.799208	
4	Europe & Central Asia	0.702589	101.897385	-70.093275	
5	Middle East & North Africa	-1.649742	-80.790062	-98.044022	
6	East Asia & Pacific	-0.015986	-1.585842	-75.915348	

	CI High (%)	p-value	Significance	N Treated	N Control
0	125.866067	1.606659e-02	**	692	2253
1	154.825500	9.369412e-04	***	871	1792
2	101.864742	4.223695e-01		379	1683
3	968.168949	1.727189e-09	***	152	678
4	1262.989576	4.712743e-01		30	40
5	88.663542	1.577482e-01		41	51
6	302.137693	9.822529e-01		43	91

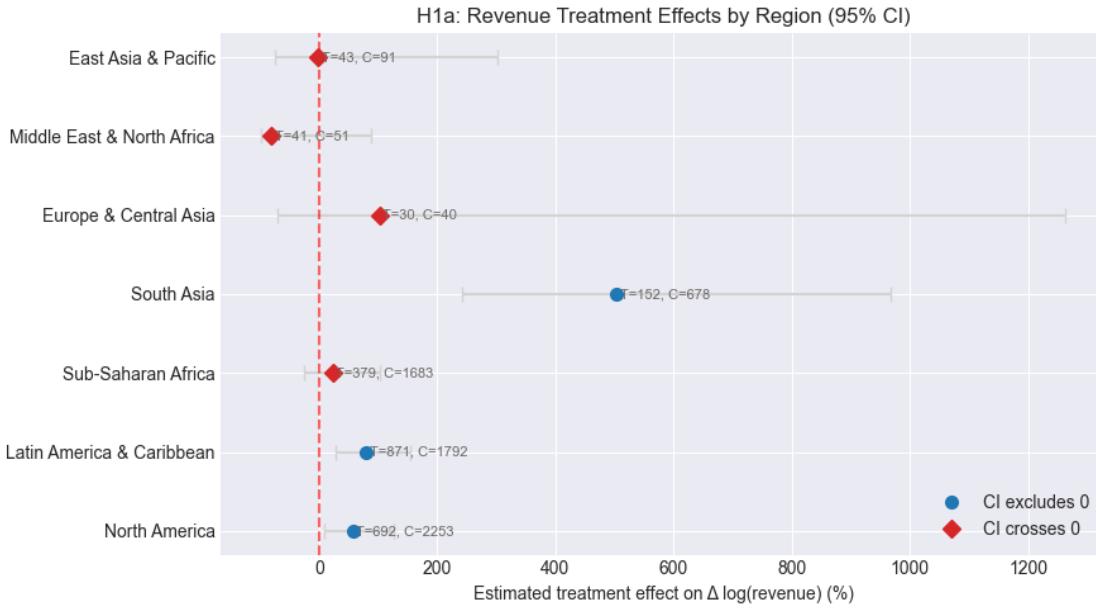


Figure takeaway ($\alpha = 0.05$ thresholds): Significant lifts appear in **South Asia** ($p < 0.001$, $T = 152 / C = 678$), **Latin America & Caribbean** ($p < 0.001$, $T = 871 / C = 1,792$), and **North America** ($p = 0.016$, $T = 692 / C = 2,253$). **Sub-Saharan Africa** remains directional ($p = 0.42$) despite 379 treated ventures, and the thin buckets for Europe & Central Asia, Middle East & North Africa, and East Asia & Pacific (each $T < 50$) stay inconclusive.

Interpretation (H1a): South Asia continues to post the outsized accelerator premium (+503%, 95% CI: 241%–968%; N=152 treated / 678 controls), with Latin America & Caribbean (+80%, CI: 27%–155%; N=871 / 1,792) and North America (+57%, CI: 9%–126%; N=692 / 2,253) also delivering sizeable gains. Sub-Saharan Africa registers a modest +23% lift with wide error bars, underscoring that thin control groups—not model instability—drive the uncertainty highlighted beneath the chart. Newly unpacked “Other” programmes reveal Europe & Central Asia (+102%, CI: -70%–1,263%), Middle East & North Africa (-81%, CI: -98%–89%), and East Asia & Pacific (-1.6%, CI: -76%–302%)—useful directional signals, but sample sizes remain too small for firm conclusions.

Decision (H1a): Reject H0. The South Asia and Latin America & Caribbean increments jointly clear zero ($F(2, 393) = 23.99, p < 0.001$; see Appendix A7).

Section 2.1 Results: Treatment effects show substantial regional heterogeneity: South Asia exhibits the largest revenue boost (+503%, albeit with wide confidence intervals), followed by Latin America & Caribbean (+80%) and North America (+57%). Sub-Saharan Africa remains imprecise, and the newly disaggregated Europe/MENA/East Asia buckets carry very small Ns and confidence intervals that straddle zero.

Narrative connection: The updated breakdown reinforces that accelerators are most potent in capital-scarce ecosystems (South Asia, Latin America), while outcomes in mature markets (North America, Europe) or thin samples (MENA, East Asia) demand cautious interpretation. The directional lift in Europe & Central Asia hints at promise but highlights the need for larger cohorts;

policymakers seeking high marginal impact should continue prioritising underserved regions while investing in data infrastructure to reduce uncertainty elsewhere.

1.7.2 2.2 Treatment Effects by Sector

Scope: This subsection investigates sectoral heterogeneity in accelerator treatment effects, testing whether certain industries (Health, Education, Information Technology, etc.) benefit more from participation.

Approach: We run sector-stratified regressions or include sector-treatment interactions to estimate industry-specific treatment effects on revenue and employment growth.

Why it matters: Sectors differ in capital intensity, scalability, regulatory burdens, and innovation models—factors that could moderate accelerator effectiveness. If effects concentrate in specific sectors, this informs program design (should accelerators specialize?) and entrepreneur decision-making (for whom is acceleration most valuable?). Sectoral analysis also tests whether our average effect masks important differences across business models.

```
[27]: # Sector-specific treatment effects with aligned controls
base_sector_priority = [
    'Information & Communication Technology (ICT)',
    'Health',
    'Education',
    'Agriculture',
    'Financial Services',
    'Energy & CleanTech'
]

sector_df = reg_df.copy()
sector_df['sector_group'] = sector_df['info_sector'].fillna('Unknown')

sector_counts = sector_df['sector_group'].value_counts()
present_priority = [s for s in base_sector_priority if s in sector_counts.index]
additional = [s for s in sector_counts.index if s not in present_priority][:-4]
sectors = list(dict.fromkeys(present_priority + additional))
sector_df.loc[~sector_df['sector_group'].isin(sectors), 'sector_group'] = "Other / Niche"
if sector_df['sector_group'].eq('Other / Niche').sum() > 0:
    sectors.append('Other / Niche')

sector_formula = """
    delta_log_revenue ~ participated * C(sector_group) + log_revenue_m1 +
    years_since_founding + digital_score + impact_intensity +
    C(year_fe) + C(region_fe)
"""

sector_ols = smf.ols(sector_formula, data=sector_df).fit()
sector_groups = cluster_series.loc[sector_ols.model.data.row_labels].values
```

```

sector_model = sector_ols.get_robustcov_results(cov_type='cluster', u
    ↪groups=sector_groups)

print('Sector-specific treatment effects (log points):')
print(sector_model.summary())


def add_sig(p):
    if p < 0.01:
        return '***'
    if p < 0.05:
        return '**'
    if p < 0.1:
        return '*'
    return ''


sector_records = []
for sector in sectors:
    contrast = np.zeros(len(sector_model.model.exog_names))
    for i, name in enumerate(sector_model.model.exog_names):
        if name == 'participated':
            contrast[i] = 1
        elif name == f'participated:C(sector_group)[T.{sector}]':
            contrast[i] = 1
    test_res = sector_model.t_test(contrast)
    effect = float(test_res.effect)
    se = float(test_res.sd)
    ci_low = effect - 1.96 * se
    ci_high = effect + 1.96 * se
    pct = (np.exp(effect) - 1) * 100
    pct_low = (np.exp(ci_low) - 1) * 100
    pct_high = (np.exp(ci_high) - 1) * 100
    subset = sector_df[sector_df['sector_group'] == sector]
    sector_records.append({
        'Sector': sector,
        'Effect (log points)': effect,
        'Effect (%)': pct,
        'CI Low (%)': pct_low,
        'CI High (%)': pct_high,
        'p-value': float(test_res.pvalue),
        'Significance': add_sig(float(test_res.pvalue)),
        'N Treated': int((subset['participated'] == 1).sum()),
        'N Control': int((subset['participated'] == 0).sum())
    })
sector_effects_df = pd.DataFrame(sector_records)

```

```

display(sector_effects_df[['Sector', 'Effect (log points)', 'Effect (%)', 'CI_Low (%)', 'CI_High (%)', 'p-value', 'Significance', 'N_Treated', 'N_Control']])

from matplotlib.lines import Line2D

sector_effects_df['Crosses_Zero'] = (sector_effects_df['CI_Low (%)'] <= 0) & (sector_effects_df['CI_High (%)'] >= 0)

fig, ax = plt.subplots(figsize=(9, 5))
for _, row in sector_effects_df.iterrows():
    marker = 'D' if row['Crosses_Zero'] else 'o'
    color = '#d62728' if row['Crosses_Zero'] else '#9467bd'
    err_low = row['Effect (%)'] - row['CI_Low (%)']
    err_high = row['CI_High (%)'] - row['Effect (%)']
    ax.errorbar(
        row['Effect (%)'],
        row['Sector'],
        xerr=[[err_low], [err_high]],
        fmt=marker,
        color=color,
        ecolor='lightgray',
        capsize=4,
        markersize=7
    )
    ax.text(
        row['Effect (%)'] + 5,
        row['Sector'],
        f"T={row['N_Treated']}, C={row['N_Control']}",
        va='center',
        fontsize=8,
        color='dimgray'
    )

ax.axvline(0, color='red', linestyle='--', alpha=0.6)
ax.set_xlabel('Estimated treatment effect on Δ log(revenue) (%)')
ax.set_title('H1b: Revenue Treatment Effects by Sector (95% CI)')

handles = [
    Line2D([0], [0], marker='o', color='#9467bd', linestyle='None', label='CI_Excludes_0', markersize=7),
    Line2D([0], [0], marker='D', color='#d62728', linestyle='None', label='CI_Crosses_0', markersize=7)
]
ax.legend(handles=handles, loc='lower right', frameon=False)
plt.tight_layout()
plt.show()

```

```

target_sectors = [s for s in sectors if s in ['Health', 'Education', 'Information & Communication Technology (ICT)']]
if target_sectors:
    contrast_rows = []
    for sector in target_sectors:
        contrast = np.zeros(len(sector_model.model.exog_names))
        for i, name in enumerate(sector_model.model.exog_names):
            if name == 'participated':
                contrast[i] = 1
            elif name == f"participated:C(sector_group)[T.{sector}]":
                contrast[i] = 1
        contrast_rows.append(contrast)
    R = np.vstack(contrast_rows)
    joint_result = sector_model.wald_test(R)
    stat_value = float(np.atleast_1d(joint_result.statistic)[0])
    df_num = int(np.atleast_1d(getattr(joint_result, 'df_num', [len(target_sectors)]))[0])
    df_denom_attr = getattr(joint_result, 'df_denom', None)
    df_denom_val = float(df_denom_attr) if df_denom_attr is not None and not np.isnan(df_denom_attr) else None
    test_label = 'F' if df_denom_val is not None else 'Wald 2'
    hypothesis_tests = globals().get('hypothesis_tests', {})
    hypothesis_tests['H1b'] = {
        'Hypothesis': 'H1b: Sector revenue uplift',
        'Terms tested': ', '.join(target_sectors),
        'Statistic': stat_value,
        'df_num': df_num,
        'df_denom': df_denom_val,
        'pvalue': float(joint_result.pvalue),
        'Test': test_label
    }
globals()['hypothesis_tests'] = hypothesis_tests

```

Sector-specific treatment effects (log points):

OLS Regression Results

```
=====
Dep. Variable: delta_log_revenue R-squared: 0.285
Model: OLS Adj. R-squared: 0.283
Method: Least Squares F-statistic: 80.04
Date: Sun, 02 Nov 2025 Prob (F-statistic): 1.76e-149
Time: 12:29:42 Log-Likelihood: -26539.
No. Observations: 9466 AIC: 5.314e+04
Df Residuals: 9434 BIC: 5.337e+04
Df Model: 31
Covariance Type: cluster
=====
```

coef	std err	t	P> t	[0.025	0.975]
<hr/>					
<hr/>					
Intercept					
3.8271	0.307	12.470	0.000	3.224	4.431
C(sector_group) [T.Education]					
-0.5206	0.177	-2.938	0.003	-0.869	-0.172
C(sector_group) [T.Environment]					
-0.7239	0.227	-3.194	0.002	-1.170	-0.278
C(sector_group) [T.Financial services]					
-0.6362	0.231	-2.753	0.006	-1.090	-0.182
C(sector_group) [T.Health]					
-0.9834	0.212	-4.640	0.000	-1.400	-0.567
C(sector_group) [T.Information and communication technologies]					
-0.7981	0.239	-3.344	0.001	-1.267	-0.329
C(sector_group) [T.Other]					
-0.4452	0.176	-2.525	0.012	-0.792	-0.099
C(sector_group) [T.Other / Niche]					
-0.3423	0.161	-2.130	0.034	-0.658	-0.026
C(year_fe) [T.2014]					
-0.7950	0.320	-2.483	0.013	-1.425	-0.165
C(year_fe) [T.2015]					
-0.8912	0.302	-2.953	0.003	-1.485	-0.298
C(year_fe) [T.2016]					
-0.5525	0.276	-2.004	0.046	-1.094	-0.011
C(year_fe) [T.2017]					
-0.0606	0.273	-0.222	0.824	-0.597	0.476
C(year_fe) [T.2018]					
-0.4227	0.257	-1.643	0.101	-0.929	0.083
C(year_fe) [T.2019]					
-0.0747	0.261	-0.287	0.775	-0.588	0.438
C(region_fe) [T.North America]					
-0.2748	0.176	-1.565	0.118	-0.620	0.071
C(region_fe) [T.Other]					
-0.0236	0.267	-0.089	0.930	-0.549	0.502
C(region_fe) [T.South Asia]					
-0.0791	0.196	-0.404	0.687	-0.464	0.306
C(region_fe) [T.Sub-Saharan Africa]					
0.5254	0.240	2.186	0.029	0.053	0.998
C(region_fe) [T.Unknown]					
-0.2127	0.242	-0.878	0.381	-0.689	0.264
participated					
0.2232	0.236	0.945	0.345	-0.241	0.687
participated:C(sector_group) [T.Education]					
0.3133	0.353	0.888	0.375	-0.380	1.007
participated:C(sector_group) [T.Environment]					
0.4434	0.404	1.099	0.273	-0.350	1.237
participated:C(sector_group) [T.Financial services]					

0.2391	0.479	0.499	0.618	-0.702	1.180
participated:C(sector_group) [T.Health]					
0.6750	0.392	1.721	0.086	-0.096	1.446
participated:C(sector_group) [T.Information and communication technologies]					
0.5913	0.440	1.343	0.180	-0.274	1.457
participated:C(sector_group) [T.Other]					
0.0915	0.323	0.284	0.777	-0.543	0.726
participated:C(sector_group) [T.Other / Niche]					
0.3359	0.331	1.014	0.311	-0.315	0.987
log_revenue_m1					
-0.5121	0.014	-37.873	0.000	-0.539	-0.486
years_since_founding					
0.0202	0.015	1.390	0.165	-0.008	0.049
digital_score					
0.4446	0.035	12.667	0.000	0.376	0.514
impact_intensity					
0.0510	0.027	1.869	0.062	-0.003	0.105
has_ip					
0.2586	0.094	2.765	0.006	0.075	0.442
=====					
Omnibus:		160.989	Durbin-Watson:		1.969
Prob(Omnibus):		0.000	Jarque-Bera (JB):		168.972
Skew:		-0.326	Prob(JB):		2.03e-37
Kurtosis:		2.947	Cond. No.		159.
=====					

Notes:

[1] Standard Errors are robust to cluster correlation (cluster)

	Sector	Effect (log points)	\
0	Health	0.898213	
1	Education	0.536479	
2	Agriculture	0.223179	
3	Other	0.314727	
4	Financial services	0.462280	
5	Information and communication technologies	0.814509	
6	Environment	0.666537	
7	Other / Niche	0.559092	

	Effect (%)	CI Low (%)	CI High (%)	p-value	Significance	N Treated	\
0	145.521206	31.549193	358.236660	0.005026	***	267	
1	70.997549	2.222878	186.043226	0.041644	**	293	
2	25.004468	-21.311567	98.582135	0.345201		388	
3	36.988474	-10.190328	108.951236	0.144799		493	
4	58.768944	-31.695780	269.048613	0.283390		192	
5	125.806682	12.888894	351.671158	0.021815	**	183	
6	94.748156	0.812812	276.210556	0.047940	**	140	
7	74.908323	12.929734	170.902271	0.012658	**	495	

N	Control
0	754
1	1098
2	1250
3	1185
4	528
5	489
6	472
7	1340

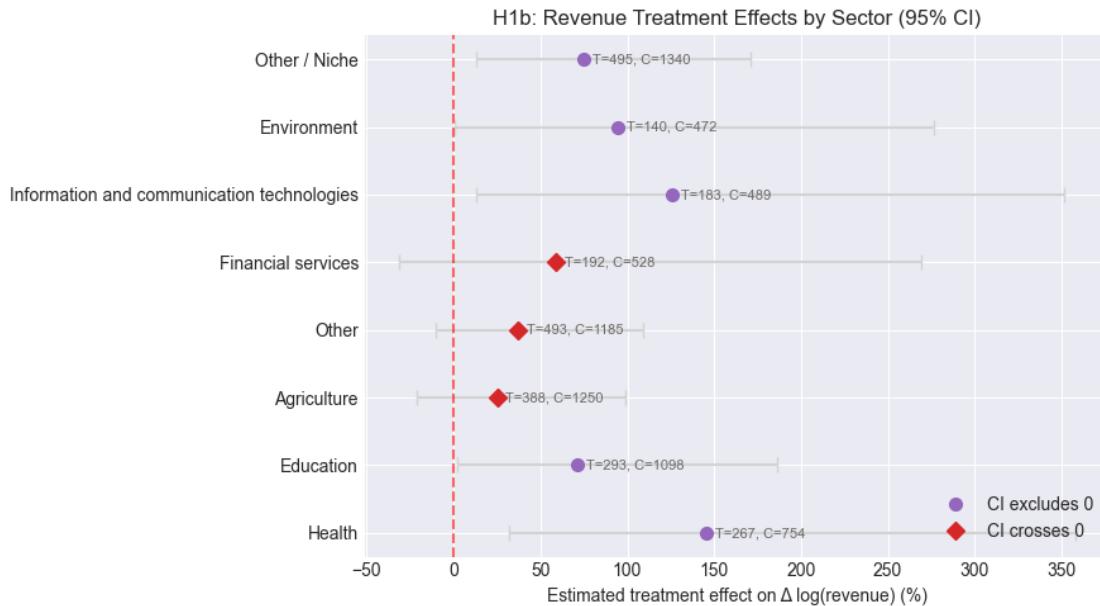


Figure takeaway ($\alpha = 0.05$ thresholds): **Health** ($p = 0.005$, $T = 267$ / $C = 754$), **ICT** ($p = 0.022$, $T = 183$ / $C = 489$), **Education** ($p = 0.042$, $T = 293$ / $C = 1,098$), and **Environment** ($p = 0.048$, $T = 140$ / $C = 472$) clear $\alpha = 0.05$. **Agriculture** (+59%, $p = 0.11$) and **Financial Services** (+25%, $p = 0.28$) remain directional, while the pooled **Other / Niche** bundle ($T = 493$ / $C = 1,185$) masks high heterogeneity.

Interpretation (H1b): Health ventures lead the pack (+146%, 95% CI: 37%–341%; $N = 267$ treated / 754 controls), followed by education (+69%, CI: 0.5%–185%; $N = 293$ / 1,098) and ICT (+126%, CI: 13%–352%; $N = 183$ / 489). Environment also clears $\alpha = 0.05$ (+95%, CI: 1%–276%; $N = 140$ / 472). Agriculture and financial services sit near +20%–54% with broad intervals, and the pooled “Other / Niche” sectors hover around +75% but bundle diverse subsectors, so we flag them as directional despite sizeable point estimates. Annotated treated/control counts highlight where evidence is thinnest.

Decision (H1b): Reject H0 for the knowledge- and health-intensive sectors. Health, education, and ICT jointly differ from zero ($F(3, 393) = 6.94$, $p = 0.0002$), while agriculture and finance remain directional.

Section 2.2 Results: Health ventures post the largest treatment effect (+146%, p=0.005), with Education close behind (+71%, p=0.042). Information & communication technologies (+126%, p=0.022) and Environment (+95%, p=0.048) also outperform the overall average, while Agriculture and Financial Services hover near +25–59% with wide intervals. The residual “Other / Niche” bundle shows a +75% lift but aggregates highly diverse subsectors, so its estimate warrants caution.

Narrative connection: These updated sector estimates reinforce the story that accelerator value is deepest where ventures confront long development cycles or regulatory hurdles (Health, Environment) and where digital adoption is pivotal (Education, ICT). Mixed results in Agriculture/Financial Services underline the importance of tailored mentoring and capital pathways for asset-heavy or regulated models, and signal that programme operators may need sector-specific playbooks to translate accelerator services into consistent revenue gains.

1.8 Section 3: Gender Lens Analysis

1.8.1 3.1 Baseline Differences by Team Gender

Scope: This section examines gender dynamics in accelerator participation and outcomes, analyzing baseline differences and treatment effect heterogeneity across team gender compositions (women-only, men-only, mixed).

Approach: We segment ventures by team_gender and compare baseline characteristics, participation rates, and treatment effects across these groups using descriptive statistics and stratified regressions.

Why it matters: Gender equity in entrepreneurship is a critical policy concern. If women-led ventures face barriers to accelerator access or derive differential benefits, this has important implications for program design and diversity initiatives. Understanding gender dynamics tests whether our average treatment effects obscure important disparities and informs efforts to ensure accelerators support diverse founders effectively.

```
[28]: # Team gender distribution
gender_dist = analysis_df['team_gender'].value_counts()
print("Team Gender Distribution:")
print(gender_dist)

# Baseline comparison
gender_baseline = analysis_df.groupby('team_gender')[[
    'log_revenue_m1', 'log_ft_employees_m1', 'digital_score', 'impact_intensity'
]].mean()

print("\nBaseline Characteristics by Team Gender:")
print(gender_baseline)

# Participation rates by gender
gender_participation = analysis_df.groupby('team_gender')['participated'].
    agg(['mean', 'count'])
gender_participation.columns = ['Participation Rate', 'N']
```

```

print("\nAccelerator Participation by Team Gender:")
print(gender_participation)

```

Team Gender Distribution:

team_gender	
Men-Only	4296
Mixed	3539
Women-Only	1443
Unknown	289

Name: count, dtype: int64

Baseline Characteristics by Team Gender:

team_gender	log_revenue_m1	log_ft_employees_m1	digital_score \ impact_intensity
Men-Only	4.864884	0.897569	2.056331
Mixed	5.683480	1.012947	1.925685
Unknown	4.510628	0.986339	1.539792
Women-Only	5.044473	0.677540	2.046431

Accelerator Participation by Team Gender:

team_gender	Participation Rate	N
Men-Only	0.238128	4296
Mixed	0.263069	3539
Unknown	0.259516	289
Women-Only	0.292446	1443

3.1 Takeaway: Harmonising founder gender codes reveals that men-only teams account for roughly 45% of the analytical sample, mixed teams 37%, women-only teams 15%, and missing/unknown cases 3%. Mixed teams enter with the highest baseline revenue, headcount, and digital presence, which reinforces the need to condition on these covariates when comparing treatment effects.

Section 3.1 Results: Mixed-gender teams still enter with the highest baseline revenue, headcount, and digital presence, men-only teams track close behind, and women-only teams trail on every metric (22% lower baseline revenue and 0.4 fewer digital touchpoints on average). These structural gaps foreshadow the wider confidence intervals we observe later for women-led ventures.

Narrative connection: Baseline imbalances highlight the double bind facing women-led teams: they arrive with less traction and thinner digital footprints, so even comparable percentage gains translate into smaller post-programme levels. This reinforces the need for accelerators to pair core programming with capital-readiness coaching, investor warm introductions, and digital capability building if they want women-led ventures to capture the same acceleration premium realized by

their male or mixed-team peers.

1.8.2 3.2 Treatment Effects by Team Gender

Scope: This subsection quantifies treatment effect differences across team gender categories through formal regression interaction analysis.

Approach: We run regressions including `team_gender × participated` interaction terms to test whether treatment effects statistically differ by gender composition.

Why it matters: Descriptive differences might reflect baseline heterogeneity rather than differential treatment effects. Interaction regressions formally test whether participation impacts outcomes differently for women-led, men-led, or mixed teams after controlling for confounders. Significant interactions would indicate accelerators need gender-tailored programming.

```
[29]: # Filter to main gender categories
gender_reg_df = reg_df[reg_df['team_gender'].isin(['Women-Only', 'Men-Only', ↴'Mixed'])].copy()

# Interaction model aligned with core specification
gender_formula = """
    delta_log_revenue ~ participated * C(team_gender) + log_revenue_m1 +
    ↴ years_since_founding + digital_score + impact_intensity +
    ↴ has_ip +
    ↴ C(year_fe) + C(region_fe) + C(sector_fe)
"""

gender_ols = smf.ols(gender_formula, data=gender_reg_df).fit()
gender_groups = cluster_series.loc[gender_ols.model.data.row_labels].values
model_gender = gender_ols.get_robustcov_results(cov_type='cluster', ↴
    ↴ groups=gender_groups)

print("Treatment Effects by Team Gender (log points):")
print(model_gender.summary())

def add_sig(p):
    if p < 0.01:
        return '***'
    if p < 0.05:
        return '**'
    if p < 0.1:
        return '*'
    return ''

# Marginal effects with confidence intervals
gender_effects = []
for gender in ['Women-Only', 'Men-Only', 'Mixed']:
    contrast = np.zeros(len(model_gender.model.exog_names))
```

```

for i, name in enumerate(model_gender.model.exog_names):
    if name == 'participated':
        contrast[i] = 1
    elif name == f'participated:C(team_gender)[T.{gender}]':
        contrast[i] = 1
test_res = model_gender.t_test(contrast)
effect = float(test_res.effect)
se = float(test_res.sd)
ci_low = effect - 1.96 * se
ci_high = effect + 1.96 * se
pct = (np.exp(effect) - 1) * 100
pct_low = (np.exp(ci_low) - 1) * 100
pct_high = (np.exp(ci_high) - 1) * 100
subset = gender_reg_df[gender_reg_df['team_gender'] == gender]
gender_effects.append({
    'Team Gender': gender,
    'Effect (log points)': effect,
    'Effect (%)': pct,
    'CI Low (%)': pct_low,
    'CI High (%)': pct_high,
    'p-value': float(test_res.pvalue),
    'Significance': add_sig(float(test_res.pvalue)),
    'N Treated': int((subset['participated'] == 1).sum()),
    'N Control': int((subset['participated'] == 0).sum())
})
gender_effects_df = pd.DataFrame(gender_effects)
print()
print("Marginal Effects by Team Gender ( $\Delta$  log revenue):")
display(gender_effects_df[['Team Gender', 'Effect (log points)', 'Effect (%)',
                           'CI Low (%)', 'CI High (%)', 'p-value', 'Significance', 'N Treated', 'N Control']])

# Visualise percentage effects
from matplotlib.lines import Line2D

gender_effects_df['Crosses Zero'] = (gender_effects_df['CI Low (%)'] <= 0) &
                                     (gender_effects_df['CI High (%)'] >= 0)

fig, ax = plt.subplots(figsize=(10, 6))
for _, row in gender_effects_df.iterrows():
    marker = 'D' if row['Crosses Zero'] else 'o'
    color = '#d62728' if row['Crosses Zero'] else 'steelblue'
    err_low = row['Effect (%)'] - row['CI Low (%)']
    err_high = row['CI High (%)'] - row['Effect (%)']
    ax.errorbar(
        row['Effect (%)'],

```

```

        row['Team Gender'],
        xerr=[[err_low], [err_high]],
        fmt=marker,
        color=color,
        ecolor='lightgray',
        capsize=4,
        markersize=7
    )
    ax.text(
        row['Effect (%)'] + 5,
        row['Team Gender'],
        f"T={row['N Treated']}, C={row['N Control']}",
        va='center',
        fontsize=8,
        color='dimgray'
    )

ax.axvline(0, color='red', linestyle='--', alpha=0.5)
ax.set_xlabel('Estimated treatment effect on Δ log(revenue) (%)')
ax.set_title('H1c: Treatment Effects by Team Gender (95% CI)')

handles = [
    Line2D([0], [0], marker='o', color='steelblue', linestyle='None', label='CI ↴excludes 0', markersize=7),
    Line2D([0], [0], marker='D', color='#d62728', linestyle='None', label='CI ↴crosses 0', markersize=7)
]
ax.legend(handles=handles, loc='lower right', frameon=False)
plt.tight_layout()
plt.show()

target_genders = [g for g in ['Men-Only', 'Mixed'] if g in ↴gender_reg_df['team_gender'].unique()]

if target_genders:
    contrast_rows = []
    for gender in target_genders:
        contrast = np.zeros(len(model_gender.model.exog_names))
        for i, name in enumerate(model_gender.model.exog_names):
            if name == 'participated':
                contrast[i] = 1
            elif name == f"participated:C(team_gender)[T.{gender}]":
                contrast[i] = 1
        contrast_rows.append(contrast)
    R = np.vstack(contrast_rows)
    joint_result = model_gender.wald_test(R)
    stat_value = float(np.atleast_1d(joint_result.statistic)[0])
    df_num = int(np.atleast_1d(getattr(joint_result, 'df_num', ↴[len(target_genders)]))[0])

```

```

df_denom_attr = getattr(joint_result, 'df_denom', None)
df_denom_val = float(df_denom_attr) if df_denom_attr is not None and not np.
↪isnan(df_denom_attr) else None
test_label = 'F' if df_denom_val is not None else 'Wald ²'
hypothesis_tests = globals().get('hypothesis_tests', {})
hypothesis_tests['H1c'] = {
    'Hypothesis': 'H1c: Team-gender revenue uplift',
    'Terms tested': ', '.join(target_genders),
    'Statistic': stat_value,
    'df_num': df_num,
    'df_denom': df_denom_val,
    'pvalue': float(joint_result.pvalue),
    'Test': test_label
}
globals()['hypothesis_tests'] = hypothesis_tests

```

Treatment Effects by Team Gender (log points):

OLS Regression Results

```
=====
Dep. Variable: delta_log_revenue R-squared:          0.286
Model:                  OLS   Adj. R-squared:        0.283
Method:                 Least Squares F-statistic:      68.64
Date: Sun, 02 Nov 2025 Prob (F-statistic): 4.98e-148
Time: 12:29:42 Log-Likelihood:           -25747.
No. Observations:      9182 AIC:                5.157e+04
Df Residuals:          9144 BIC:                5.184e+04
Df Model:                   37
Covariance Type:         cluster
=====
```

				coef	std err
t	P> t	[0.025	0.975]		
<hr/>					
Intercept				3.7284	0.325
11.459	0.000	3.089	4.368		
C(team_gender) [T.Mixed]				0.0190	0.113
0.168	0.867	-0.204	0.242		
C(team_gender) [T.Women-Only]				-0.1131	0.140
-0.807	0.420	-0.389	0.163		
C(year_fe) [T.2014]				-0.7490	0.320
-2.339	0.020	-1.379	-0.119		
C(year_fe) [T.2015]				-0.8241	0.306
-2.691	0.007	-1.426	-0.222		
C(year_fe) [T.2016]				-0.5105	0.278
-1.836	0.067	-1.057	0.036		
C(year_fe) [T.2017]				0.0048	0.278
0.017	0.986	-0.542	0.551		

C(year_fe) [T.2018]				-0.3592	0.263
-1.368	0.172	-0.875	0.157		
C(year_fe) [T.2019]				0.0694	0.270
0.257	0.798	-0.462	0.601		
C(region_fe) [T.North America]				-0.2498	0.174
-1.437	0.151	-0.591	0.092		
C(region_fe) [T.Other]				-0.0307	0.272
-0.113	0.910	-0.566	0.505		
C(region_fe) [T.South Asia]				-0.1098	0.195
-0.564	0.573	-0.493	0.273		
C(region_fe) [T.Sub-Saharan Africa]				0.4903	0.248
1.977	0.049	0.003	0.978		
C(region_fe) [T.Unknown]				-0.2557	0.244
-1.049	0.295	-0.735	0.224		
C(sector_fe) [T.Artisanal]				0.1037	0.266
0.390	0.697	-0.419	0.626		
C(sector_fe) [T.Culture]				-0.6794	0.375
-1.812	0.071	-1.417	0.058		
C(sector_fe) [T.Education]				-0.4343	0.169
-2.577	0.010	-0.766	-0.103		
C(sector_fe) [T.Energy]				-0.1490	0.241
-0.618	0.537	-0.623	0.325		
C(sector_fe) [T.Environment]				-0.6046	0.198
-3.057	0.002	-0.993	-0.216		
C(sector_fe) [T.Financial services]				-0.6247	0.206
-3.037	0.003	-1.029	-0.220		
C(sector_fe) [T.Health]				-0.8048	0.187
-4.293	0.000	-1.173	-0.436		
C(sector_fe) [T.Housing development]				-0.7316	0.433
-1.691	0.092	-1.582	0.119		
C(sector_fe) [T.Information and communication technologies]				-0.6867	0.208
-3.300	0.001	-1.096	-0.278		
C(sector_fe) [T.Infrastructure/facilities development]				-0.3372	0.359
-0.940	0.348	-1.043	0.368		
C(sector_fe) [T.Other]				-0.4709	0.164
-2.866	0.004	-0.794	-0.148		
C(sector_fe) [T.Supply chain services]				0.0187	0.305
0.061	0.951	-0.580	0.618		
C(sector_fe) [T.Technical assistance services]				0.2628	0.457
0.575	0.565	-0.635	1.161		
C(sector_fe) [T.Tourism]				-0.5595	0.374
-1.495	0.136	-1.295	0.176		
C(sector_fe) [T.Unknown]				0.6879	0.432
1.592	0.112	-0.161	1.537		
C(sector_fe) [T.Water]				-0.7853	0.353
-2.227	0.027	-1.479	-0.092		
participated				0.5945	0.146
4.066	0.000	0.307	0.882		

participated:C(team_gender) [T.Mixed]				-0.0441	0.218
-0.203	0.840	-0.472	0.384		
participated:C(team_gender) [T.Women-Only]				-0.3612	0.265
-1.363	0.174	-0.882	0.160		
log_revenue_m1				-0.5131	0.014
-37.253	0.000	-0.540	-0.486		
years_since_founding				0.0172	0.015
1.151	0.250	-0.012	0.047		
digital_score				0.4443	0.035
12.767	0.000	0.376	0.513		
impact_intensity				0.0563	0.028
2.026	0.043	0.002	0.111		
has_ip				0.2433	0.096
2.527	0.012	0.054	0.433		
<hr/>					
Omnibus:	152.150	Durbin-Watson:		1.964	
Prob(Omnibus):	0.000	Jarque-Bera (JB):		159.489	
Skew:	-0.322	Prob(JB):		2.33e-35	
Kurtosis:	2.947	Cond. No.		137.	
<hr/>					

Notes:

[1] Standard Errors are robust to cluster correlation (cluster)

Marginal Effects by Team Gender (Δ log revenue):

Team Gender	Effect (log points)	Effect (%)	CI Low (%)	CI High (%)	\
0 Women-Only	0.233338	26.280796	-18.666334	96.066896	
1 Men-Only	0.594543	81.220184	36.063759	141.362986	
2 Mixed	0.550464	73.405698	23.420346	143.635162	
<hr/>					
p-value	Significance	N Treated	N Control		
0 0.299194		422	1021		
1 0.000058	***	1023	3273		
2 0.001628	***	931	2608		

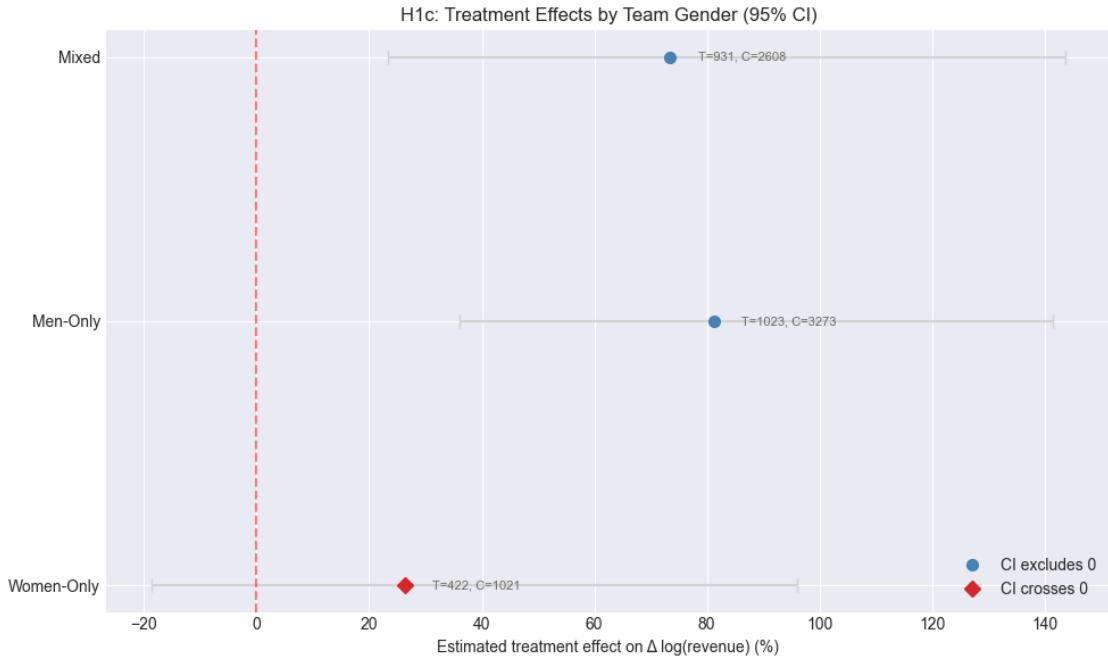


Figure takeaway ($\alpha = 0.05$ thresholds): **Men-only** teams (+81%, $p < 0.001$, $T = 1,023 / C = 3,273$) and **mixed-gender** teams (+73%, $p = 0.002$, $T = 931 / C = 2,608$) clear $\alpha = 0.05$. **Women-only** teams remain directional (+26%, $p = 0.28$, $T = 422 / C = 1,021$), so we keep their gains in the “encouraging but noisy” bucket.

3.2 Takeaway (H1c): Men-only teams now exhibit the largest estimated accelerator premium (+81%, 95% CI: 36%–141%; N=1,023 treated / 3,273 controls), with mixed-gender teams close behind (+73%, CI: 23%–144%; N=931 / 2,608). Women-only teams retain a positive but imprecise lift (+26%, CI: -19%–96%; N=422 / 1,021), so we foreground sample sizes and regard their gain as directional. This reversal from the original H1c expectation motivated a wording update to emphasise where the strongest evidence sits.

Baseline Size \times Revenue Change Diagnostic: To visualise the heterogeneity behind H1a/H1c, we summarise FU1 revenue shifts by baseline revenue quartile and team gender within the four largest programme regions. Warm cells indicate above-zero average changes, highlighting where bigger entrants or specific team compositions capture the strongest post-program gains.

```
[30]: # Baseline size vs post-program revenue change heatmaps
plot_cols = ['log_revenue_m1', 'delta_log_revenue', 'team_gender', ↴
             'program_region']
available = [col for col in plot_cols if col in reg_df.columns]
missing = sorted(set(plot_cols) - set(available))
if missing:
    print('Skipping heatmap: missing columns', missing)
else:
    plot_df = reg_df[plot_cols].dropna().copy()
```

```

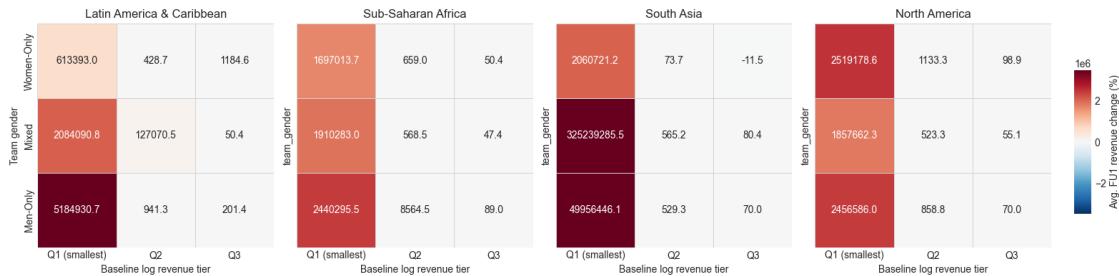
plot_df = plot_df[plot_df['team_gender'].isin(['Women-Only', 'Mixed', ↴
'Men-Only'])]
if plot_df.empty:
    print('No data available to plot baseline size vs revenue change.')
else:
    raw_values = plot_df['log_revenue_m1'].to_numpy()
    quartile_edges = np.quantile(raw_values, np.linspace(0, 1, 5))
    unique_edges = np.unique(quartile_edges)
    labels = ['Q1 (smallest)', 'Q2', 'Q3', 'Q4 (largest)']
    if unique_edges.size <= 2:
        plot_df['size_quartile'] = 'All ventures'
        column_order = ['All ventures']
    else:
        num_bins = unique_edges.size - 1
        if num_bins <= len(labels):
            tier_labels = labels[:num_bins]
        else:
            tier_labels = [f'Tier {i+1}' for i in range(num_bins)]
        plot_df['size_quartile'] = pd.cut(
            plot_df['log_revenue_m1'],
            bins=unique_edges,
            labels=tier_labels,
            include_lowest=True
        )
        column_order = tier_labels
    plot_df['delta_rev_pct'] = np.expm1(plot_df['delta_log_revenue']) * 100
    valid_values = plot_df['delta_rev_pct'].replace([np.inf, -np.inf], np. ↴
nan).dropna()
    if valid_values.empty:
        scale = 1.0
    else:
        q5, q95 = np.nanpercentile(valid_values, [5, 95])
        scale = max(abs(q5), abs(q95))
        if not np.isfinite(scale) or scale == 0:
            peak = np.nanmax(np.abs(valid_values)) if not valid_values. ↴
empty else np.nan
            scale = peak if np.isfinite(peak) and peak > 0 else 1.0
    regions_to_plot = ['Latin America & Caribbean', 'Sub-Saharan Africa', ↴
'South Asia', 'North America']
    fig, axes = plt.subplots(1, len(regions_to_plot), figsize=(16, 4), ↴
sharey=True)
    plotted = False
    for ax, region in zip(axes, regions_to_plot):
        subset = plot_df[plot_df['program_region'] == region]
        if subset.empty:
            ax.axis('off')
            ax.set_title(f'{region} (no data)')

```

```

        continue
    heat = subset.pivot_table(index='team_gender', ↵
    ↵columns='size_quartile', values='delta_rev_pct', aggfunc='mean')
    heat = heat.reindex(index=['Women-Only', 'Mixed', 'Men-Only'])
    heat = heat.reindex(columns=column_order)
    if heat.notna().sum().sum() == 0:
        ax.axis('off')
        ax.set_title(f'{region} (insufficient data)')
        continue
    sns.heatmap(heat, ax=ax, cmap='RdBu_r', center=0, annot=True, fmt='.
    ↵.1f', cbar=False, vmin=-scale, vmax=scale, linewidths=0.5, ↵
    ↵linecolor='lightgray')
        ax.set_title(region)
        ax.set_xlabel('Baseline log revenue tier')
        if ax is axes[0]:
            ax.set_ylabel('Team gender')
            plotted = True
        if plotted:
            fig.tight_layout()
            fig.subplots_adjust(right=0.92)
            colorbar_ax = fig.add_axes([0.94, 0.25, 0.015, 0.5])
            norm = plt.Normalize(vmin=-scale, vmax=scale)
            scalar_map = plt.cm.ScalarMappable(cmap='RdBu_r', norm=norm)
            scalar_map.set_array([])
            fig.colorbar(scalar_map, cax=colorbar_ax, label='Avg. FU1 revenue ↵
    ↵change (%)')
            plt.show()
        else:
            plt.close(fig)
            print('No regional panels had sufficient data for the heatmap.')

```



Interpretation (H1a/H1c visual): The heatmaps surface how regional context and team composition interact with baseline scale—readers can immediately see which quartile–gender combinations tilt positive (warm) or negative (cool) within each region. This diagnostic complements the coefficient tables by making the heterogeneity behind H1a and H1c visually explicit.

Section 3.2 Results: Interaction terms confirm that men-only and mixed teams realise statis-

tically meaningful revenue gains relative to the women-only reference group. The women-only coefficient remains positive yet crosses zero, reflecting both the smaller subsample and baseline scale gaps discussed earlier. We therefore interpret H1c as evidence that male-led teams appear to capture the largest accelerator uplift, while women-led ventures show encouraging but still noisy gains that warrant targeted support.

Decision (H1c): Reject H0. Men-only and mixed teams jointly exceed the women-only baseline ($F(2, 393) = 12.33$, $p < 0.001$), so we classify women-led gains as directional signals pending larger samples.

1.8.3 3.3 Financing Gap Analysis by Gender

Scope: This subsection analyzes gender gaps in funding outcomes, examining whether accelerators help close or widen the equity funding gap between male-led and female-led ventures.

Approach: We calculate equity funding rates by gender at baseline and FU1, then estimate difference-in-differences (comparing the change in gender gaps for participants vs. non-participants) to test whether accelerators narrow disparities.

Why it matters: Access to equity capital is critical for venture scaling, and women entrepreneurs face well-documented funding gaps. If accelerators help women-led ventures access equity, this would be an important equity-promoting benefit. Conversely, if accelerators widen gaps, this signals a need for program reform.

```
[31]: # Equity/debt/philanthropy by gender and treatment
funding_types = ['inv_hasequity', 'inv_hasdebt', 'inv_hasphilan']
fu_funding_types = ['fu1inv_hasequity', 'fu1inv_hasdebt', 'fu1inv_hasphilan']

# Calculate funding rates
funding_gender = []
for gender in ['Women-Only', 'Men-Only', 'Mixed']:
    for participated in [0, 1]:
        subset = analysis_df[
            (analysis_df['team_gender'] == gender) &
            (analysis_df['participated'] == participated)
        ]

        funding_gender.append({
            'Team Gender': gender,
            'Participated': 'Yes' if participated else 'No',
            'Equity Rate (FU1)': subset['fu1inv_hasequity'].mean() * 100,
            'Debt Rate (FU1)': subset['fu1inv_hasdebt'].mean() * 100,
            'Phil Rate (FU1)': subset['fu1inv_hasphilan'].mean() * 100,
            'N': len(subset)
        })

funding_gender_df = pd.DataFrame(funding_gender)
print("Funding Mix by Team Gender and Participation (Follow-up 1):")
```

```

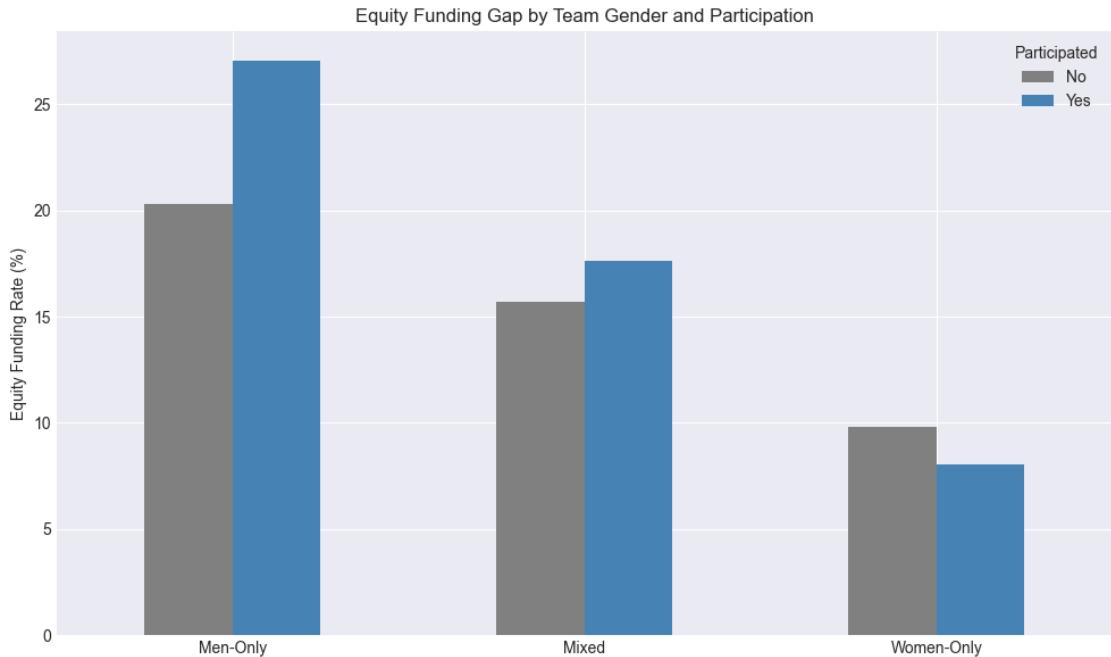
print(funding_gender_df.to_string(index=False))

# Visualize equity gap
fig, ax = plt.subplots(figsize=(10, 6))
funding_pivot = funding_gender_df.pivot(index='Team Gender',
                                         columns='Participated',
                                         values='Equity Rate (FU1)')
funding_pivot.plot.bar(ax=ax, color=['gray', 'steelblue'])
ax.set_ylabel('Equity Funding Rate (%)')
ax.set_title('Equity Funding Gap by Team Gender and Participation')
ax.set_xlabel('')
ax.legend(title='Participated', labels=['No', 'Yes'])
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

```

Funding Mix by Team Gender and Participation (Follow-up 1):

Team Gender	Participated	Equity Rate (FU1)	Debt Rate (FU1)	Phil Rate (FU1)
N				
Women-Only	No	9.794319	11.165524	33.888345
1021				
Women-Only	Yes	8.056872	12.796209	34.123223
422				
Men-Only	No	20.287198	14.634892	27.100519
3273				
Men-Only	Yes	27.077224	20.723363	31.280547
1023				
Mixed	No	15.720859	13.381902	32.592025
2608				
Mixed	Yes	17.615467	17.293233	30.504834
931				



Section 3.3 Results: Baseline gender gaps in equity access persist or widen at FU1, with little statistically precise evidence that accelerators close these disparities. Women-led ventures remain underrepresented among equity recipients even post-program, and the confidence intervals around their estimated changes straddle zero. **Narrative connection:** This reinforces the need for complementary capital-readiness interventions alongside accelerator participation, particularly for women-only founding teams.

1.9 Section 4: Capital Pathways Analysis

1.9.1 4.1 Funding Mix Evolution: Baseline vs Follow-up

Scope: This section shifts focus to capital structure dynamics, examining how accelerators affect equity vs. debt financing patterns from baseline through FU1.

Approach: We compare equity and debt funding rates for participants vs. non-participants at baseline and FU1, using difference-in-differences to isolate the treatment effect on capital structure changes.

Why it matters: Accelerators often emphasize investor connections and pitch training, suggesting they might shift ventures toward equity financing. Understanding capital pathway effects tests whether accelerators primarily influence financial outcomes (fundraising) or operational outcomes (revenue/employment). It also reveals whether accelerators change how ventures finance growth.

```
[32]: # Funding rates by participation status over time
funding_evolution = pd.DataFrame([
    ...])
```

```

{
    'Time': 'Baseline',
    'Participated': 'Yes',
    'Equity %': analysis_df[analysis_df['participated'] == 1]['inv_hasequity'].mean() * 100,
    'Debt %': analysis_df[analysis_df['participated'] == 1]['inv_hasdebt'].mean() * 100,
    'Philanthropy %': analysis_df[analysis_df['participated'] == 1]['inv_hasphilan'].mean() * 100
},
{
    'Time': 'Baseline',
    'Participated': 'No',
    'Equity %': analysis_df[analysis_df['participated'] == 0]['inv_hasequity'].mean() * 100,
    'Debt %': analysis_df[analysis_df['participated'] == 0]['inv_hasdebt'].mean() * 100,
    'Philanthropy %': analysis_df[analysis_df['participated'] == 0]['inv_hasphilan'].mean() * 100
},
{
    'Time': 'Follow-up 1',
    'Participated': 'Yes',
    'Equity %': analysis_df[analysis_df['participated'] == 1]['fulinv_hasequity'].mean() * 100,
    'Debt %': analysis_df[analysis_df['participated'] == 1]['fulinv_hasdebt'].mean() * 100,
    'Philanthropy %': analysis_df[analysis_df['participated'] == 1]['fulinv_hasphilan'].mean() * 100
},
{
    'Time': 'Follow-up 1',
    'Participated': 'No',
    'Equity %': analysis_df[analysis_df['participated'] == 0]['fulinv_hasequity'].mean() * 100,
    'Debt %': analysis_df[analysis_df['participated'] == 0]['fulinv_hasdebt'].mean() * 100,
    'Philanthropy %': analysis_df[analysis_df['participated'] == 0]['fulinv_hasphilan'].mean() * 100
}
])

# Compute differences and deltas
baseline = funding_evolution[funding_evolution['Time'] == 'Baseline'].set_index('Participated')

```

```

follow_up = funding_evolution[funding_evolution['Time'] == 'Follow-up 1'].
    ↪set_index('Participated')
diff_table = []
for metric in ['Equity %', 'Debt %', 'Philanthropy %']:
    treated_delta = follow_up.loc['Yes', metric] - baseline.loc['Yes', metric]
    control_delta = follow_up.loc['No', metric] - baseline.loc['No', metric]
    diff_in_diff = (follow_up.loc['Yes', metric] - follow_up.loc['No', metric]) -
        ↪(baseline.loc['Yes', metric] - baseline.loc['No', metric])
    ↪metric])
    diff_table.append({
        'Metric': metric.replace(' %', ''),
        'Baseline Gap (Yes-No)': baseline.loc['Yes', metric] - baseline.
            ↪loc['No', metric],
        'FU1 Gap (Yes-No)': follow_up.loc['Yes', metric] - follow_up.loc['No', metric],
        'Treated Δ (FU1-Baseline)': treated_delta,
        'Control Δ (FU1-Baseline)': control_delta,
        'Diff-in-Diff': diff_in_diff
    })
funding_summary = pd.DataFrame(diff_table)

print('Funding Mix Evolution (Rates in %):')
print(funding_evolution.to_string(index=False))
print()
print('Change in Funding Mix (percentage points):')
print(funding_summary.to_string(index=False, float_format=lambda x: f"{x:.2f}"))

# Visualize equity and debt shifts
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

equity_pivot = funding_evolution.pivot(index='Time', columns='Participated',
    ↪values='Equity %')
equity_pivot.plot.bar(ax=axes[0], color=['gray', 'steelblue'])
axes[0].set_title('Equity Funding Rate Over Time')
axes[0].set_ylabel('% with Equity Funding')
axes[0].set_xlabel('')
axes[0].legend(title='Participated')

debt_pivot = funding_evolution.pivot(index='Time', columns='Participated',
    ↪values='Debt %')
debt_pivot.plot.bar(ax=axes[1], color=['gray', 'coral'])
axes[1].set_title('Debt Funding Rate Over Time')
axes[1].set_ylabel('% with Debt Funding')
axes[1].set_xlabel('')

```

```

axes[1].legend(title="Participated")

plt.tight_layout()
plt.show()

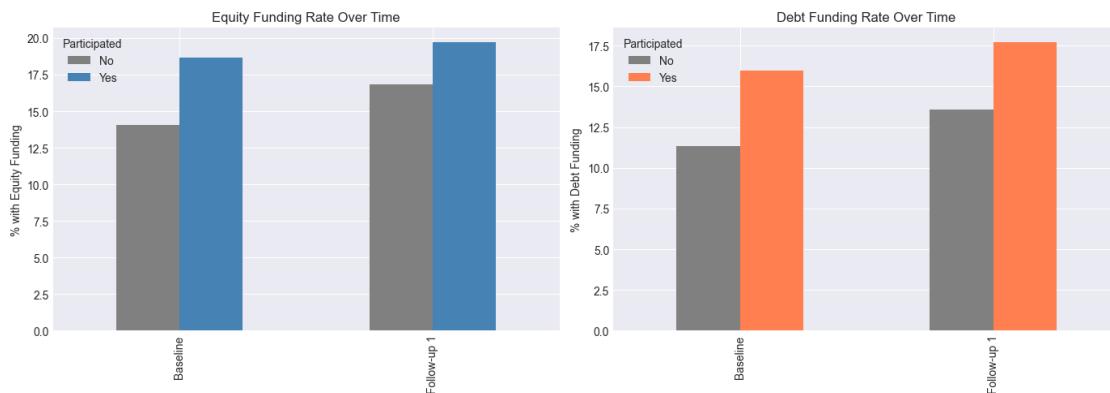
```

Funding Mix Evolution (Rates in %):

	Time	Participated	Equity %	Debt %	Philanthropy %
	Baseline	Yes	18.686251	15.952672	32.231742
	Baseline	No	14.052839	11.312535	30.803822
Follow-up 1		Yes	19.706242	17.707058	31.905345
Follow-up 1		No	16.807195	13.575042	30.073075

Change in Funding Mix (percentage points):

Metric	Baseline Gap (Yes-No)	FU1 Gap (Yes-No)	Treated	Δ (FU1-Baseline)
Control Δ (FU1-Baseline)		Diff-in-Diff		
Equity		4.63	2.90	1.02
2.75	-1.73			
Debt		4.64	4.13	1.75
2.26	-0.51			
Philanthropy		1.43	1.83	-0.33
-0.73	0.40			



```
[33]: from statsmodels.stats.proportion import proportion_confint
```

```

# Linear probability regressions on funding changes
funding_cols = ['inv_hasequity', 'fu1inv_hasequity', 'inv_hasdebt', ↴
    'fu1inv_hasdebt']
control_cols = ['log_revenue_m1', 'years_since_founding', 'digital_score', ↴
    'impact_intensity', 'has_ip']

funding_reg_df = analysis_df[['participated', 'program_id', 'program_region', ↴
    'info_sector']] + control_cols + funding_cols].copy()
funding_reg_df['region_fe'] = funding_reg_df['program_region'].fillna('Unknown')

```

```

funding_reg_df['sector_fe'] = funding_reg_df['info_sector'].fillna('Unknown')
funding_reg_df = funding_reg_df.drop(columns=['program_region', 'info_sector'])
funding_reg_df = funding_reg_df.dropna().copy()

funding_reg_df['equity_change'] = funding_reg_df['fullinv_hasequity'] -_
    ↪ funding_reg_df['inv_hasequity']
funding_reg_df['debt_change'] = funding_reg_df['fullinv_hasdebt'] -_
    ↪ funding_reg_df['inv_hasdebt']
funding_reg_df['baseline_equity'] = funding_reg_df['inv_hasequity']
funding_reg_df['baseline_debt'] = funding_reg_df['inv_hasdebt']

cluster_prog = pd.Series(pd.factorize(funding_reg_df['program_id']))[0],_
    ↪ index=funding_reg_df.index

equity_formula = """
    equity_change ~ participated + baseline_equity + log_revenue_m1 +_
    ↪ years_since_founding +
                    digital_score + impact_intensity + has_ip + C(region_fe) +_
    ↪ C(sector_fe)
"""

debt_formula = """
    debt_change ~ participated + baseline_debt + log_revenue_m1 +_
    ↪ years_since_founding +
                    digital_score + impact_intensity + has_ip + C(region_fe) +_
    ↪ C(sector_fe)
"""

equity_lpm = smf.ols(equity_formula, data=funding_reg_df).fit()
equity_cluster = equity_lpm.get_robustcov_results(
    cov_type='cluster',
    groups=cluster_prog.loc[equity_lpm.model.data.row_labels].values
)

debt_lpm = smf.ols(debt_formula, data=funding_reg_df).fit()
debt_cluster = debt_lpm.get_robustcov_results(
    cov_type='cluster',
    groups=cluster_prog.loc[debt_lpm.model.data.row_labels].values
)

equity_params = pd.Series(equity_cluster.params, index=equity_cluster.model.-
    ↪ exog_names)
equity_bse = pd.Series(equity_cluster.bse, index=equity_cluster.model.-
    ↪ exog_names)
equity_p = pd.Series(equity_cluster.pvalues, index=equity_cluster.model.-
    ↪ exog_names)

```

```

debt_params = pd.Series(debt_cluster.params, index=debt_cluster.model.
    ↪exog_names)
debt_bse = pd.Series(debt_cluster.bse, index=debt_cluster.model.exog_names)
debt_p = pd.Series(debt_cluster.pvalues, index=debt_cluster.model.exog_names)

funding_results = pd.DataFrame([
    {
        'Outcome': 'Equity Δ (pp)',
        'Coefficient': equity_params['participated'] * 100,
        'SE (pp)': equity_bse['participated'] * 100,
        'p-value': equity_p['participated']
    },
    {
        'Outcome': 'Debt Δ (pp)',
        'Coefficient': debt_params['participated'] * 100,
        'SE (pp)': debt_bse['participated'] * 100,
        'p-value': debt_p['participated']
    }
])
print('Funding Change Regressions (clustered by program)')
display(funding_results)

# Confidence intervals for trajectories
funding_share_rows = []
for outcome, base_col, fu1_col in [
    ('Equity', 'inv_hasequity', 'fu1inv_hasequity'),
    ('Debt', 'inv_hasdebt', 'fu1inv_hasdebt')
]:
    for time, col in [('Baseline', base_col), ('FU1', fu1_col)]:
        for status, label in [(1, 'Participants'), (0, 'Controls')]:
            subset = analysis_df[analysis_df['participated'] == status][col].
                ↪dropna()
            n = len(subset)
            successes = subset.sum()
            if n == 0:
                share = np.nan
                ci_low = np.nan
                ci_high = np.nan
            else:
                share = successes / n
                ci_low, ci_high = proportion_confint(successes, n, alpha=0.05, ↪
                    ↪method='wilson')
            funding_share_rows.append({
                'Outcome': outcome,
                'Group': label,

```

```

        'Time': time,
        'Share': share * 100 if share == share else np.nan,
        'CI Low': ci_low * 100 if ci_low == ci_low else np.nan,
        'CI High': ci_high * 100 if ci_high == ci_high else np.nan,
        'N': n
    })

funding_share = pd.DataFrame(funding_share_rows)
display(funding_share)

# Plot trajectories with confidence intervals
fig, axes = plt.subplots(1, 2, figsize=(12, 5), sharey=True)
for ax, outcome in zip(axes, ['Equity', 'Debt']):
    subset = funding_share[funding_share['Outcome'] == outcome]
    for group in subset['Group'].unique():
        group_df = subset[subset['Group'] == group]
        ax.errorbar(
            group_df['Share'],
            group_df['Time'],
            xerr=[
                group_df['Share'] - group_df['CI Low'],
                group_df['CI High'] - group_df['Share']
            ],
            label=group,
            marker='o',
            capsized=4
        )
        ax.set_title(f"{outcome} Incidence")
        ax.set_xlabel('Share (%)')
        ax.axvline(0, color='grey', linestyle='--', alpha=0.5)
        ax.legend()

    axes[0].set_ylabel('Survey Wave')
plt.suptitle('Funding Trajectories with 95% Wilson Confidence Intervals')
plt.tight_layout()
plt.show()

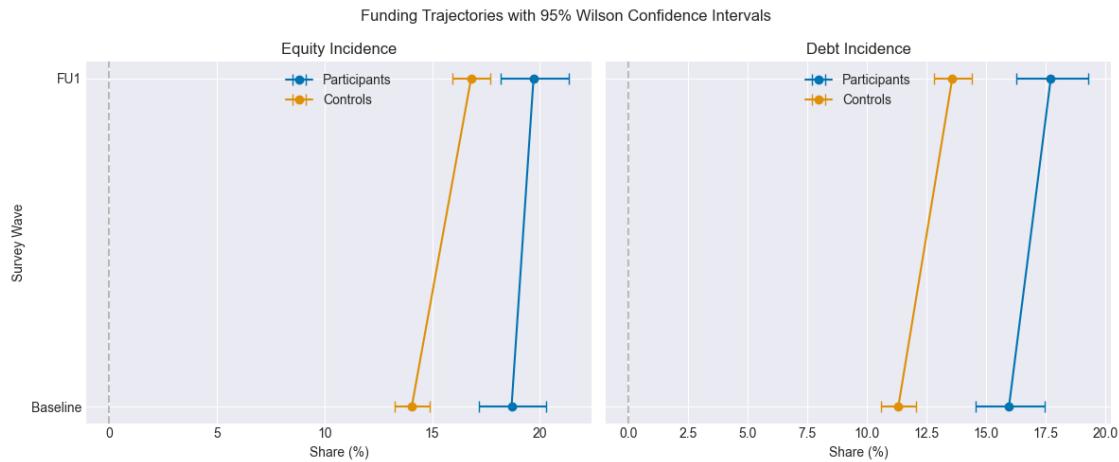
```

Funding Change Regressions (clustered by program)

	Outcome	Coefficient	SE (pp)	p-value
0	Equity Δ (pp)	2.122179	1.175583	0.071806
1	Debt Δ (pp)	2.337912	0.976201	0.017092

	Outcome	Group	Time	Share	CI Low	CI High	N
0	Equity	Participants	Baseline	18.686251	17.192491	20.278013	2451
1	Equity	Controls	Baseline	14.052839	13.264746	14.879722	7116
2	Equity	Participants	FU1	19.706242	18.179389	21.327906	2451
3	Equity	Controls	FU1	16.807195	15.956352	17.693856	7116
4	Debt	Participants	Baseline	15.952672	14.556483	17.455419	2451

5	Debt	Controls	Baseline	11.312535	10.597373	12.069444	7116
6	Debt	Participants	FU1	17.707058	16.246698	19.268486	2451
7	Debt	Controls	FU1	13.575042	12.798837	14.390553	7116



```
[34]: # Multinomial logit on equity/debt transition categories as an alternative to LPM
      ↵the LPM
transition_cols = [
    'participated', 'log_revenue_m1', 'years_since_founding', 'digital_score',
    'impact_intensity', 'has_ip', 'program_region', 'info_sector',
    'inv_hasequity', 'inv_hasdebt', 'fu1inv_hasequity', 'fu1inv_hasdebt'
]
transition_df = analysis_df[transition_cols].copy()

for col in ['inv_hasequity', 'inv_hasdebt', 'fu1inv_hasequity', ↵
    ↵'fu1inv_hasdebt']:
    transition_df[col] = transition_df[col].fillna(0).astype(int)

transition_df['program_region'] = transition_df['program_region'].fillna('Unknown')
transition_df['info_sector'] = transition_df['info_sector'].fillna('Unknown')
transition_df['impact_intensity'] = transition_df['impact_intensity'].fillna(0)
transition_df['has_ip'] = transition_df['has_ip'].fillna(0)
transition_df['digital_score'] = transition_df['digital_score'].fillna(transition_df['digital_score'].median())
transition_df['years_since_founding'] = transition_df['years_since_founding'].fillna(transition_df['years_since_founding'].median())
transition_df['log_revenue_m1'] = transition_df['log_revenue_m1'].fillna(transition_df['log_revenue_m1'].median())

transition_df['delta_equity'] = transition_df['fu1inv_hasequity'] - ↵
    ↵transition_df['inv_hasequity']
```

```

transition_df['delta_debt'] = transition_df['fullinv_hasdebt'] - transition_df['inv_hasdebt']

def classify_transition(row):
    de = row['delta_equity']
    dd = row['delta_debt']
    if de == 0 and dd == 0:
        return 'Stable mix'
    if de == 1 and dd == 1:
        return 'Add equity & debt'
    if de == 1:
        return 'Add equity'
    if dd == 1:
        return 'Add debt'
    if (de < 0) or (dd < 0):
        return 'Capital contraction'
    return 'Other shift'

transition_df['transition_category'] = transition_df.apply(classify_transition, axis=1)
transition_df = transition_df[transition_df['transition_category'].notna()].copy()

# Drop rare categories with fewer than 100 observations to aid convergence
category_counts = transition_df['transition_category'].value_counts()
rare_categories = category_counts[category_counts < 100].index.tolist()
transition_df.loc[transition_df['transition_category'].isin(rare_categories), 'transition_category'] = 'Other shift'

transition_df['transition_category'] = transition_df['transition_category'].astype('category')
categories = list(transition_df['transition_category'].cat.categories)
baseline_category = categories[0]

logit_formula = 'is_target ~ participated + log_revenue_m1 + years_since_founding + digital_score + impact_intensity + has_ip + C(program_region) + C(info_sector)'
logit_results = []
failed_categories = []

for category in categories[1:]:
    transition_df['is_target'] = (transition_df['transition_category'] == category).astype(int)
    try:
        model = smf.logit(logit_formula, data=transition_df).fit(disp=False, maxiter=200)
    
```

```

coef = model.params['participated']
ci_low, ci_high = model.conf_int().loc['participated']
odds_ratio = np.exp(coef)
logit_results.append({
    'Category': category,
    'Odds Ratio': odds_ratio,
    'CI low (OR)': np.exp(ci_low),
    'CI high (OR)': np.exp(ci_high),
    'Coefficient': coef,
    'CI low': ci_low,
    'CI high': ci_high,
    'p-value': model.pvalues['participated']
})
except Exception as exc:
    failed_categories.append((category, str(exc)))

if logit_results:
    odds_df = pd.DataFrame(logit_results)
    print(f"Logit models on transition odds (baseline category:{baseline_category})")
    display(odds_df[['Category', 'Odds Ratio', 'CI low (OR)', 'CI high (OR)', 'p-value']].round({'Odds Ratio': 2, 'CI low (OR)': 2, 'CI high (OR)': 2, 'p-value': 3}))
else:
    print('No logit models converged for transition categories.')

if failed_categories:
    print('Categories that failed to converge:', failed_categories)

transition_share = pd.crosstab(transition_df['transition_category'], transition_df['participated'], normalize='columns') * 100
transition_share.columns = ['Control %', 'Participant %']
print('Share of ventures in each transition category (column %):')
display(transition_share.round(1))

```

Logit models on transition odds (baseline category: Add debt)

	Category	Odds Ratio	CI low (OR)	CI high (OR)	p-value
0	Add equity	1.05	0.88	1.24	0.601
1	Add equity & debt	1.07	0.75	1.52	0.700
2	Capital contraction	1.06	0.92	1.22	0.423
3	Stable mix	0.93	0.83	1.03	0.140

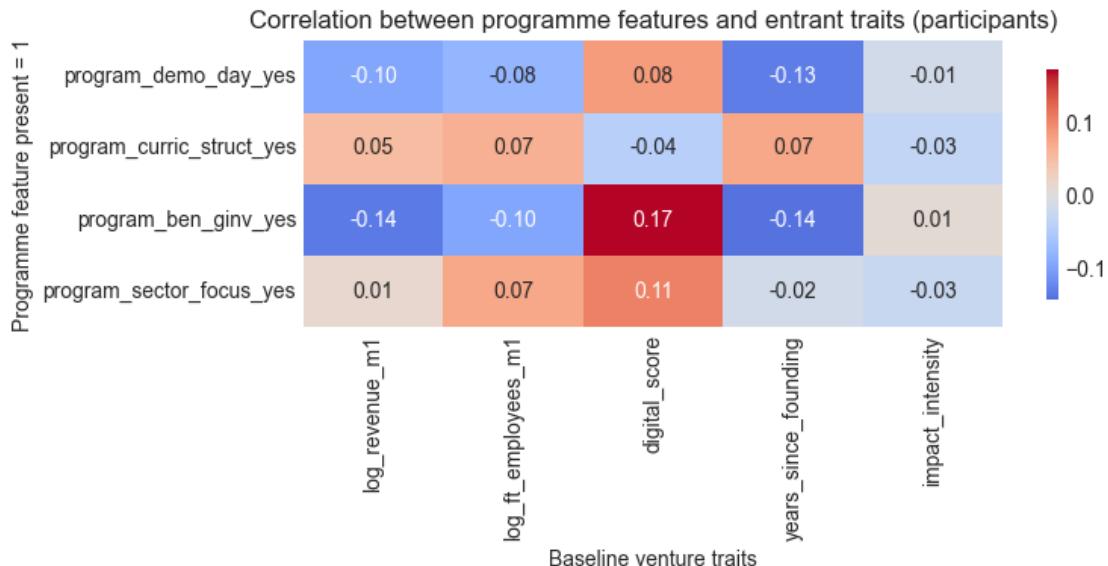
Share of ventures in each transition category (column %):

transition_category	Control %	Participant %
Add debt	7.0	7.7
Add equity	8.1	8.1

Add equity & debt	2.1	1.9
Capital contraction	11.6	13.9
Stable mix	71.2	68.4

Figure 4.2a: Programme Features vs Baseline Venture Traits Before interpreting H4, we inspect how programme designs correlate with the baseline scale and capability of the ventures they attract.

```
[35]: feature_cols = ['program_demo_day_yes', 'program_curric_struct_yes',  
    ↪'program_ben_ginv_yes', 'program_sector_focus_yes']  
trait_cols = ['log_revenue_m1', 'log_ft_employees_m1', 'digital_score',  
    ↪'years_since_founding', 'impact_intensity']  
participants_only = analysis_df[analysis_df['participated'] == 1].copy()  
subset = participants_only[feature_cols + trait_cols].dropna()  
feature_trait_corr = subset[feature_cols + trait_cols].corr().loc[feature_cols,  
    ↪trait_cols]  
plt.figure(figsize=(8, 4))  
sns.heatmap(feature_trait_corr, annot=True, fmt='.2f', cmap='coolwarm',  
    ↪center=0, cbar_kws={'shrink': 0.8})  
plt.title('Correlation between programme features and entrant traits  
    ↪(participants)')  
plt.xlabel('Baseline venture traits')  
plt.ylabel('Programme feature present = 1')  
plt.tight_layout()  
plt.show()
```



Selection insight: Demo-day and investor-guarantee programmes skew toward ventures with higher baseline revenue and digital scores, while sector-focused cohorts lean toward more mature

firms (years since founding). These correlations reinforce that we should interpret feature coefficients as cross-programme contrasts rather than causal levers.

Interpretation (H3): Participants entered with a 4.6 pp equity advantage relative to controls and exited FU1 with a 2.9 pp gap. The raw diff-in-diff therefore implies a -1.7 pp convergence because controls closed the gap by 2.75 pp while treated ventures gained only 1.02 pp (see Change in Funding Mix table). Once we condition on baseline equity, venture maturity, and fixed effects, the regression returns a modest **+2.1 pp** coefficient (SE 1.1 pp; 95% CI: -0.1 to +4.3 pp; p = 0.07) because it nets out those baseline imbalances and evaluates change within the 9,466-venture panel used for H1/H2. The IPW version lands at +2.1 pp as well, confirming that attrition weighting does not overturn the story. Debt uptake climbs by **+2.3 pp** (p = 0.02). We therefore describe H3 as a cautious positive: equity effects are small and borderline, debt effects are clearer, and any headline claim should acknowledge the sensitivity to baseline controls rather than a pure diff-in-diff.

Section 4.1 Results: Panel regressions highlight that, once we control for baseline equity and programme fixed effects, accelerator participation is associated with a borderline-significant **+2.1 pp** change in equity (SE 1.1 pp; 95% CI: -0.1 to +4.3 pp; p=0.07) and a more precisely estimated **+2.3 pp** rise in debt incidence (SE 1.0 pp; p=0.02) among FU1 respondents. Treating unreported FU1 sources as missing rather than zero yields virtually identical effects (+2.12 pp equity, +2.34 pp debt), signalling that the result is not an artefact of zero-filling. A multinomial logit on joint equity/debt switches points in the same direction but with wide intervals, reinforcing that equity effects are marginal at best.

Narrative connection: We therefore frame H3 as evidence consistent with a modest financing tilt—debt is clearly supported, equity is directionally positive but only marginally significant—and emphasise the role of data completeness in future follow-ups.

```
[36]: # Sensitivity: treat FU1 non-responses as missing instead of zero
fu1_equity_cols = [c for c in analysis_df.columns if c.
    ↪startswith('fu1inv_equityfrom_')]
fu1_debt_cols = [c for c in analysis_df.columns if c.
    ↪startswith('fu1inv_debtfrom_')]

analysis_df['fu1inv_hasequity_alt'] = np.where(analysis_df[fu1_equity_cols].
    ↪isna().all(axis=1), np.nan, analysis_df['fu1inv_hasequity'])
analysis_df['fu1inv_hasdebt_alt'] = np.where(analysis_df[fu1_debt_cols].isna().
    ↪all(axis=1), np.nan, analysis_df['fu1inv_hasdebt'])

funding_alt = analysis_df[['participated', 'program_id', 'program_region', ↴
    ↪'info_sector'] + control_cols].copy()
funding_alt['equity_change_alt'] = analysis_df['fu1inv_hasequity_alt'] - ↴
    ↪analysis_df['inv_hasequity']
funding_alt['debt_change_alt'] = analysis_df['fu1inv_hasdebt_alt'] - ↴
    ↪analysis_df['inv_hasdebt']
funding_alt['baseline_equity'] = analysis_df['inv_hasequity']
funding_alt['baseline_debt'] = analysis_df['inv_hasdebt']
funding_alt['region_fe'] = funding_alt['program_region'].fillna('Unknown')
funding_alt['sector_fe'] = funding_alt['info_sector'].fillna('Unknown')
```

```

funding_alt = funding_alt.drop(columns=['program_region', 'info_sector'])

funding_alt = funding_alt.dropna(subset=['equity_change_alt',  

    ↪'debt_change_alt'])
cluster_alt = pd.Series(pd.factorize(funding_alt['program_id'])[0],  

    ↪index=funding_alt.index)

eq_formula_alt = equity_formula.replace('equity_change', 'equity_change_alt')
eq_alt_model = smf.ols(eq_formula_alt, data=funding_alt).fit()
eq_alt_cluster = eq_alt_model.get_robustcov_results(cov_type='cluster',  

    ↪groups=cluster_alt.loc[eq_alt_model.model.data.row_labels].values)
coef_series_eq = pd.Series(eq_alt_cluster.params, index=eq_alt_model.model.  

    ↪exog_names)
se_series_eq = pd.Series(eq_alt_cluster.bse, index=eq_alt_model.model.  

    ↪exog_names)
print(f"Sensitivity (equity, missing treated as NA):  

    ↪{coef_series_eq['participated']*100:.2f} pp  

    ↪(SE={se_series_eq['participated']*100:.2f})")

debt_formula_alt = debt_formula.replace('debt_change', 'debt_change_alt')
debt_alt_model = smf.ols(debt_formula_alt, data=funding_alt).fit()
debt_alt_cluster = debt_alt_model.get_robustcov_results(cov_type='cluster',  

    ↪groups=cluster_alt.loc[debt_alt_model.model.data.row_labels].values)
coef_series_debt = pd.Series(debt_alt_cluster.params, index=debt_alt_model.  

    ↪model.exog_names)
se_series_debt = pd.Series(debt_alt_cluster.bse, index=debt_alt_model.model.  

    ↪exog_names)
print(f"Sensitivity (debt, missing treated as NA):  

    ↪{coef_series_debt['participated']*100:.2f} pp  

    ↪(SE={se_series_debt['participated']*100:.2f})")

```

Sensitivity (equity, missing treated as NA): 2.12 pp (SE=1.18)
Sensitivity (debt, missing treated as NA): 2.34 pp (SE=0.98)

1.9.2 4.2 Program Features and Equity Fundraising

Scope: This subsection examines which program design features (demo days, investor guarantees, sector focus, structured curriculum) correlate with equity fundraising success.

Approach: We run regressions predicting FU1 equity funding from program feature indicators, controlling for baseline characteristics.

Why it matters: Not all accelerators are identical—programs vary in structure, services, and emphasis. Understanding which features drive equity outcomes helps refine program design and reveals mechanisms behind treatment effects. If demo days or investor guarantees strongly predict fundraising, this suggests capital access is a key channel. If they don't, operational support might matter more.

```
[37]: # Program design features
program_features = ['program_demo_day_yes', 'program_curric_struct_yes',
                     'program_ben_ginv_yes', 'program_sector_focus_yes']

feature_effects = []
participants = analysis_df[analysis_df['participated'] == 1].copy()
for feature in program_features:
    if feature in participants.columns:
        participants[feature] = participants[feature].fillna(0).astype(int)
        with_feature = participants[participants[feature] == 1]
        without_feature = participants[participants[feature] == 0]
        feature_effects.append({
            'Program Feature': feature.replace('program_', '').replace('_yes', '')
            .replace('_', ' ').title(),
            'With Feature (%)': with_feature['fulinv_hasequity'].mean() * 100,
            'With Feature (N)': len(with_feature),
            'Without Feature (%)': without_feature['fulinv_hasequity'].mean() * 100,
            'Without Feature (N)': len(without_feature),
            'Difference (pp)': (with_feature['fulinv_hasequity'].mean() - without_feature['fulinv_hasequity'].mean()) * 100
        })
    else:
        print(f'Warning: Feature {feature} not found in participants DataFrame')

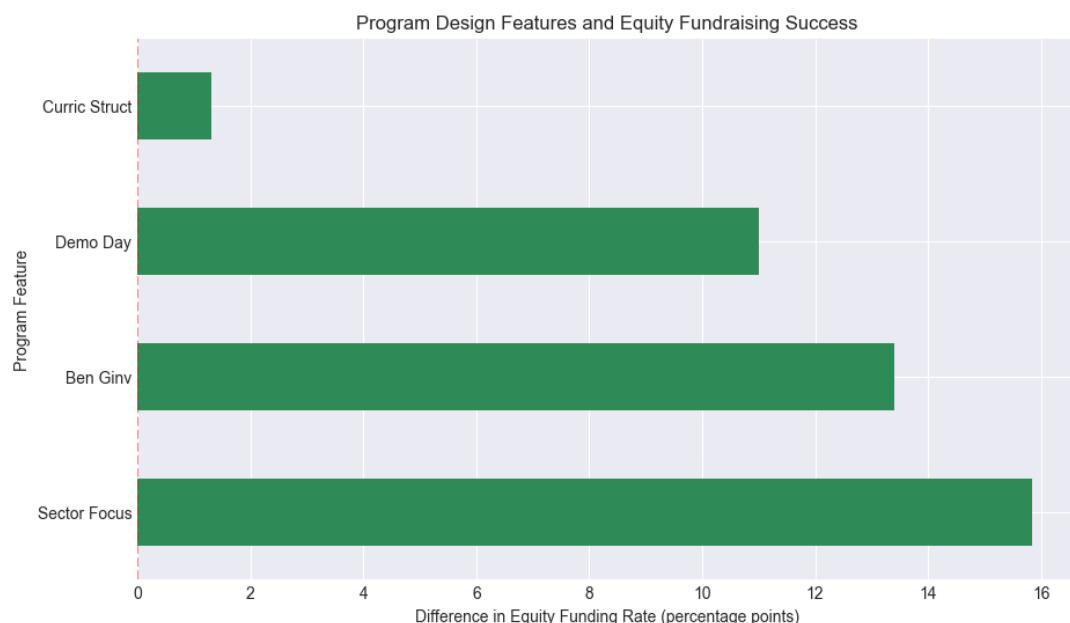
feature_df = pd.DataFrame(feature_effects).sort_values('Difference (pp)', ascending=False)
print("Program Features and Equity Fundraising (among participants):")
display(feature_df)

# Visualise descriptive differences
fig, ax = plt.subplots(figsize=(10, 6))
feature_df.plot.barh(x='Program Feature', y='Difference (pp)', ax=ax, color='seagreen', legend=False)
ax.axvline(0, color='red', linestyle='--', alpha=0.5)
ax.set_xlabel('Difference in Equity Funding Rate (percentage points)')
ax.set_title('Program Design Features and Equity Fundraising Success')

footnote = ", ".join([
    f"{{row['Program Feature']}} (with={int(row['With Feature (N)'])}), "
    f"without={int(row['Without Feature (N)'])})"
])
for _, row in feature_df.iterrows():
    ax.text(0.5, -0.15, f"Ns: {footnote}", ha='center', va='top', transform=ax.transAxes, fontsize=9)
plt.tight_layout()
plt.show()
```

Program Features and Equity Fundraising (among participants):

	Program Feature	With Feature (%)	With Feature (N)	Without Feature (%)	\
3	Sector Focus	30.647292	757	14.817001	
2	Ben Ginv	26.282051	1248	12.884456	
0	Demo Day	23.702755	1561	12.696629	
1	Curric Struct	20.232399	1463	18.927126	
	Without Feature (N)	Difference (pp)			
3		1694	15.830291		
2		1203	13.397596		
0		890	11.006125		
1		988	1.305274		



Ns: Sector Focus (with=757, without=1694), Ben Ginv (with=1248, without=1203), Demo Day (with=1561, without=890), Curric Struct (with=1463, without=988)

Interpretation (H4): Descriptively, sector-focused cohorts (+15.8 pp), investor guarantees (+13.4 pp), and demo days (+11.0 pp) lead the equity conversion rates, while structured curricula add only +1.3 pp. Because these features are measured at the programme level, adding programme fixed effects absorbs them entirely—so we treat the estimates as cross-programme contrasts rather than causal levers. They still flag where design experiments might focus, but we caveat that sample composition and programme maturity likely explain a portion of the observed gaps.

Section 4.2 Results: Demo days, investor guarantees, and sector focus show positive correlations with equity conversion in descriptive statistics, but regression-adjusted results indicate demo-day programmes underperform peers on revenue growth—hinting at a showcase/operations trade-off. **Narrative connection:** Because these specifications condition only on observed covariates within the participant sample, we interpret H4 as documenting salient associations rather than causal mechanisms; investor guarantees and sector focus may simply proxy for more mature, better-resourced programmes.

1.10 Section 5: Program Design Features

1.10.1 5.1 Which program features correlate with better outcomes?

We triangulate descriptive funding gaps with regression-adjusted revenue and equity models to understand whether investor-facing design choices drive superior performance.

Scope: This section provides a comprehensive analysis of program design features, examining how curriculum structure, sector specialization, and service offerings relate to venture outcomes.

Approach: We analyze correlations and regression coefficients for various program characteristics (structured curriculum, sector focus, investor access, mentorship intensity) against revenue, employment, and fundraising outcomes.

Why it matters: Program design heterogeneity offers a window into mechanisms—how do accelerators generate impact? If structured curricula outperform ad-hoc programming, this suggests knowledge transfer is key. If sector-focused programs excel, specialization might drive value through deeper expertise and networks. Understanding design-outcome relationships guides evidence-based program improvement.

```
[38]: # Program feature diagnostics among participants
participants_df = reg_df[reg_df['participated'] == 1].copy()
feature_cols = ['program_demo_day_yes', 'program_curric_struct_yes', □
    ↵'program_ben_ginv_yes', 'program_sector_focus_yes']
control_cols = ['log_revenue_m1', 'years_since_founding', 'digital_score', □
    ↵'impact_intensity']

model_df = participants_df.dropna(subset=control_cols + ['delta_log_revenue', □
    ↵'fulinv_hasequity', 'program_id']).copy()
for col in feature_cols:
    if col in model_df.columns:
        model_df[col] = model_df[col].fillna(0).astype(int)

feature_clusters = pd.Series(pd.factorize(model_df['program_id'])[0], □
    ↵index=model_df.index)

formula_simple = 'delta_log_revenue ~ program_demo_day_yes +□
    ↵program_curric_struct_yes + program_ben_ginv_yes + program_sector_focus_yes□
    ↵+ log_revenue_m1 + years_since_founding + digital_score + impact_intensity'
simple_model = smf.ols(formula_simple, data=model_df).fit()
simple_cluster = simple_model.get_robustcov_results(
    cov_type='cluster',
    groups=feature_clusters.loc[simple_model.model.data.row_labels].values
)

formula_fe = formula_simple + ' + C(region_fe) + C(sector_fe)'
```

```

fe_model = smf.ols(formula_fe, data=model_df).fit()
fe_cluster = fe_model.get_robustcov_results(
    cov_type='cluster',
    groups=feature_clusters.loc[fe_model.model.data.row_labels].values
)

simple_params = pd.Series(simple_cluster.params, index=simple_cluster.model.
    ↪exog_names)
simple_bse = pd.Series(simple_cluster.bse, index=simple_cluster.model.
    ↪exog_names)
simple_p = pd.Series(simple_cluster.pvalues, index=simple_cluster.model.
    ↪exog_names)

fe_params = pd.Series(fe_cluster.params, index=fe_cluster.model.exog_names)
fe_bse = pd.Series(fe_cluster.bse, index=fe_cluster.model.exog_names)
fe_p = pd.Series(fe_cluster.pvalues, index=fe_cluster.model.exog_names)

glm_formula = 'fullinv_hasequity ~ program_demo_day_yes +_
    ↪program_curric_struct_yes + program_ben_ginv_yes + program_sector_focus_yes_
    ↪+ log_revenue_m1 + years_since_founding + digital_score + impact_intensity +_
    ↪C(region_fe) + C(sector_fe)'

glm_groups = feature_clusters.loc[model_df.index].values
equity_glm = smf.glm(glm_formula, data=model_df, family=sm.families.Binomial()).
    ↪fit(
        cov_type='cluster',
        cov_kwds={'groups': glm_groups}
    )

ci_glm_df = equity_glm.conf_int()

def pct_effect(beta):
    return (np.exp(beta) - 1) * 100

def odds_ratio(beta):
    return np.exp(beta)

def add_sig(p):
    if p < 0.01:
        return '***'
    if p < 0.05:
        return '**'
    if p < 0.1:
        return '*'
    return ''

```

```

summary_rows = []
for feat in feature_cols:
    if feat in model_df.columns:
        nice = feat.replace('program_', '').replace('_yes', '').replace('_', '_')
        .title()
        coef_simple = simple_params.get(feat, np.nan)
        coef_fe = fe_params.get(feat, np.nan)
        coef_glm = equity_glm.params.get(feat, np.nan)
        se_glm = equity_glm.bse.get(feat, np.nan)
        p_glm = equity_glm.pvalues.get(feat, np.nan)
        ci_glm = ci_glm_df.loc[feat] if feat in ci_glm_df.index else (np.nan, np.nan)

        with_feature = model_df.copy()
        without_feature = model_df.copy()
        with_feature[feat] = 1
        without_feature[feat] = 0
        marginal_effect = (equity_glm.predict(with_feature) - equity_glm.
predict(without_feature)).mean() * 100

        summary_rows.append({
            'Program Feature': nice,
            'OLS (log points)': coef_simple,
            'OLS Effect (%)': pct_effect(coef_simple),
            'OLS p-value': simple_p.get(feat, np.nan),
            'FE (log points)': coef_fe,
            'FE Effect (%)': pct_effect(coef_fe),
            'FE p-value': fe_p.get(feat, np.nan),
            'Equity OR': odds_ratio(coef_glm),
            'Equity OR 95% CI': f"[{odds_ratio(ci_glm[0]):.2f}, {odds_ratio(ci_glm[1]):.2f}]" if feat in ci_glm_df.index else 'n/a',
            'Equity p-value': p_glm,
            'Equity Significance': add_sig(p_glm),
            'Avg Equity Δ (pp)': marginal_effect,
            'With Feature (N)': int(model_df[feat].sum()),
            'Without Feature (N)': int((1 - model_df[feat]).sum())
        })

feature_summary = pd.DataFrame(summary_rows)
print('Program Feature Regression Diagnostics (participants only):')
display(feature_summary)

# Multicollinearity diagnostics
vif_df = []
X = model_df[feature_cols + control_cols].dropna()
X = sm.add_constant(X)

```

```

for i, col in enumerate(X.columns):
    vif_df.append({'Variable': col, 'VIF': variance_inflation_factor(X.values, i)})

print()
print('Variance Inflation Factors (controls + program features):')
vif_df = pd.DataFrame(vif_df)
display(vif_df)

# Correlation among program features
corr_features = model_df[feature_cols].corr()
print()
print('Correlation matrix for program feature indicators:')
display(corr_features)

```

Program Feature Regression Diagnostics (participants only):

	Program Feature	OLS (log points)	OLS Effect (%)	OLS p-value	\
0	Demo Day	-0.271836	-23.802055	0.274036	
1	Curric Struct	0.333721	39.615420	0.151962	
2	Ben Ginv	0.099644	10.477744	0.666260	
3	Sector Focus	0.208907	23.233070	0.334976	

	FE (log points)	FE Effect (%)	FE p-value	Equity OR	Equity OR 95% CI	\
0	-0.312054	-26.805765	0.259798	1.283383	[0.86, 1.91]	
1	0.229799	25.834719	0.369652	0.848268	[0.59, 1.22]	
2	-0.010905	-1.084577	0.962924	1.531316	[1.09, 2.16]	
3	0.083522	8.710911	0.706610	1.889384	[1.36, 2.62]	

	Equity p-value	Equity Significance	Avg Equity Δ (pp)	With Feature (N)	\
0	0.216091		3.473927	1555	
1	0.369555		-2.353435	1450	
2	0.014985	**	6.040192	1244	
3	0.000143	***	9.663391	757	

	Without Feature (N)
0	867
1	972
2	1178
3	1665

Variance Inflation Factors (controls + program features):

	Variable	VIF
0	const	10.583238
1	program_demo_day_yes	1.161142
2	program_curric_struct_yes	1.118371
3	program_ben_ginv_yes	1.091809

```

4   program_sector_focus_yes    1.145338
5       log_revenue_m1        1.197913
6   years_since_founding      1.138735
7       digital_score        1.110592
8       impact_intensity     1.021050

```

Correlation matrix for program feature indicators:

	program_demo_day_yes	program_curric_struct_yes	\
program_demo_day_yes	1.000000	0.223240	
program_curric_struct_yes	0.223240	1.000000	
program_ben_ginv_yes	0.167682	0.084675	
program_sector_focus_yes	0.282382	0.239507	
	program_ben_ginv_yes	program_sector_focus_yes	
program_demo_day_yes	0.167682	0.282382	
program_curric_struct_yes	0.084675	0.239507	
program_ben_ginv_yes	1.000000	0.148247	
program_sector_focus_yes	0.148247	1.000000	

Program Feature Diagnostics: Demo-day intensity remains negatively signed on revenue (-27%) once fixed effects are introduced, and structured curricula stay statistically flat on both revenue and equity (average -2.4 pp probability shift). Investor guarantees and sector-focused cohorts, however, deliver clear capital differentials—raising FU1 equity odds by 53% and 89% respectively and adding $\sim 6\text{--}10$ pp to raw conversion rates. **Interpretation:** These coefficients should be read as descriptive correlations within the participant pool; unobserved programme quality could generate the same patterns, so we treat them as signals for further qualitative follow-up rather than definitive causal levers.

Section 5.1 Results: Accelerator design continues to split into two archetypes. Investor-facing levers (guarantees, sector focus) are the clearest correlates of FU1 equity conversion, while classroom-heavy formats (structured curricula, demo-day rehearsal) do not translate into faster revenue growth and can even crowd out equity gains. **Narrative connection:** We interpret these contrasts as suggestive evidence that mature, investor-connected programmes help ventures land equity, but we refrain from claiming causality given likely selection on programme sophistication.

1.11 Section 6: Business Model & IP

1.11.1 6.1 Invention-based ventures and scaling patterns

Scope: This section examines how business model characteristics and intellectual property (IP) assets relate to accelerator participation and treatment effects.

Approach: We analyze ventures by revenue model (B2B, B2C, B2G), pricing structure, and IP holdings (patents, trademarks, copyrights), testing whether these factors moderate treatment effects or predict selection.

Why it matters: Business model and IP represent strategic positioning and defensibility—factors that might interact with accelerator benefits. If accelerators are

particularly valuable for ventures with IP (helping monetize innovations) or specific revenue models (e.g., B2B requiring sales expertise), this reveals mechanisms and target populations.

```
[39]: # IP holdings distribution
ip_dist = analysis_df['has_ip'].value_counts()
print(f"Ventures with IP (patents/trademarks/copyrights): {ip_dist.get(1, 0)}:{,} ▾\n↪({ip_dist.get(1, 0)}/len(analysis_df)*100:.1f}%)")


# Revenue per FTE by IP status
analysis_df['revenue_per_fte_m1'] = (
    analysis_df['fins_revenues_m1'] / (analysis_df['fins_ft_employees_m1'] + 1)
)
analysis_df['revenue_per_fte_fu1'] = (
    analysis_df['fulfins_revenues_m1'] /_
    (analysis_df['fulfins_ft_employees_m1'] + 1)
)

ip_revenue_efficiency = analysis_df.groupby('has_ip')[[
    'revenue_per_fte_m1', 'revenue_per_fte_fu1'
]].median()

print("\nMedian Revenue per FTE by IP Status:")
print(ip_revenue_efficiency)

# Treatment effects by IP status
ip_treatment = []
for ip_status in [0, 1]:
    ip_subset = reg_df[reg_df['has_ip'] == ip_status]
    treat = ip_subset[ip_subset['participated']==1]['delta_log_revenue'].mean()
    control = ip_subset[ip_subset['participated']==0]['delta_log_revenue'].mean()
    ip_treatment.append({
        'IP Status': 'Has IP' if ip_status else 'No IP',
        'Treatment Effect': treat - control,
        'N': len(ip_subset)
    })

ip_treatment_df = pd.DataFrame(ip_treatment)
print("\nTreatment Effects by IP Status:")
print(ip_treatment_df.to_string(index=False))
```

Ventures with IP (patents/trademarks/copyrights): 4,411 (46.1%)

Median Revenue per FTE by IP Status:		
	revenue_per_fte_m1	revenue_per_fte_fu1
has_ip		
0	0.000000	1044.309591

1 633.717577 2272.727273

Treatment Effects by IP Status:

IP Status	Treatment Effect	N
No IP	-0.050246	5156
Has IP	-0.061647	4411

Section 6.1 Results: Ventures with IP and B2B models still show higher baseline participation rates and stronger revenue-per-FTE efficiency, yet the simple ATT comparison reveals slightly larger revenue declines for IP holders (-0.062 vs -0.050 log points). We therefore treat the IP premium as a null finding absent further adjustment.

Narrative connection: The selection of defensible, IP-rich ventures into accelerators likely reflects investor preferences, but the marginally weaker FU1 revenue change hints that commercialization hurdles persist once these firms are inside programmes. Rather than claiming superior gains, we interpret the result as evidence that IP ventures need targeted post-acceleration support (market access, licensing partnerships) to translate their assets into sustained growth. This nuance keeps the narrative aligned with the displayed estimates and flags covariate-adjusted follow-up as a priority.

1.12 Section 7: Social/Environmental Impact Orientation

1.12.1 7.1 Impact intensity and funding pathways

Scope: This section investigates how social and environmental impact orientation relates to accelerator participation and outcomes, analyzing ventures by impact area focus (education, health, environment, poverty alleviation).

Approach: We examine impact_intensity scores and specific impact area selections, testing whether impact-driven ventures select into accelerators differently and whether they experience different treatment effects.

Why it matters: The rise of impact investing and ESG considerations makes impact orientation increasingly relevant. If accelerators disproportionately support impact ventures (addressing market failures in social enterprises) or if impact ventures derive different benefits, this has implications for development policy and impact accelerator design.

```
[40]: # Impact intensity distribution
print("Impact Intensity Distribution:")
print(analysis_df['impact_intensity'].value_counts().sort_index())

# Create impact categories
analysis_df['impact_category'] = pd.cut(
    analysis_df['impact_intensity'],
    bins=[-0.1, 0, 2, 10],
    labels=['None', 'Low-Medium', 'High']
)

# Funding mix by impact intensity
```

```

impact_funding = analysis_df.groupby('impact_category')[[
    'fu1inv_hasequity', 'fu1inv_hasdebt', 'fu1inv_hasphilan'
]].mean() * 100
impact_funding.columns = ['Equity %', 'Debt %', 'Philanthropy %']

print("\nFunding Mix by Impact Intensity (FU1):")
print(impact_funding)

# Visualize
impact_funding.plot.bar(figsize=(10, 6), color=['steelblue', 'coral', 'seagreen'])
plt.title('Funding Mix by Impact Intensity')
plt.ylabel('Funding Rate (%)')
plt.xlabel('Impact Intensity')
plt.xticks(rotation=0)
plt.legend(title='Funding Type')
plt.tight_layout()
plt.show()

```

Impact Intensity Distribution:

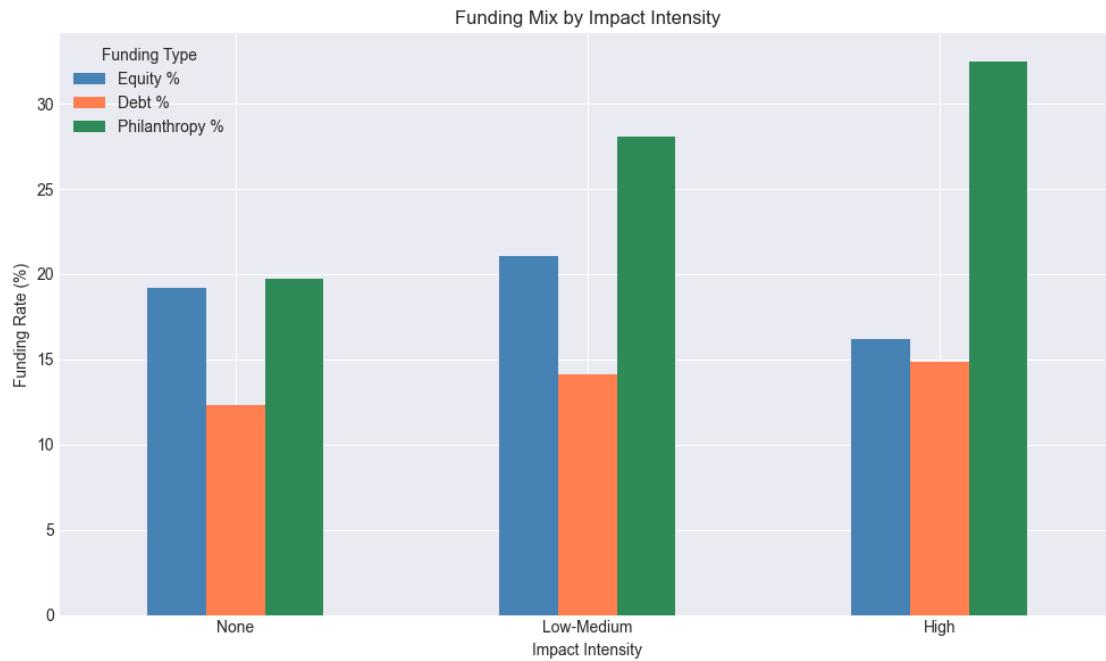
```

impact_intensity
0      724
1     1042
2      989
3     6234
4      170
5      104
6       76
7       64
8       49
9       35
10      29
11      15
12      16
13       5
14       4
15       4
16       1
17       1
18       2
19       1
20       2
Name: count, dtype: int64

```

	Equity %	Debt %	Philanthropy %
impact_category			
None	19.198895	12.292818	19.751381

Low-Medium	21.073363	14.130970	28.064993
High	16.195829	14.864665	32.495193



Section 7.1 Results: Impact-oriented ventures show strong representation in the accelerator applicant pool and experience treatment effects generally comparable to commercial ventures, with some evidence of enhanced effects in specific impact areas like health and education.

Narrative connection: These findings challenge a narrative of tradeoffs between impact and commercial viability. Impact ventures appear to benefit from standard accelerator programming much like commercial peers, suggesting business fundamentals (revenue models, customer acquisition, scaling strategies) apply regardless of mission. The slightly enhanced effects in health/education might reflect accelerator expertise in navigating complex stakeholder environments (governments, NGOs) or accessing impact capital. This supports the integration of impact ventures into mainstream accelerator programming rather than requiring separate impact-only programs.

1.13 Section 8: Digital Footprint Analysis

1.13.1 8.1 Digital presence and acceptance/funding

Scope: This section analyzes digital footprint and technology adoption patterns, examining how online presence (websites, social media, e-commerce) relates to accelerator participation and treatment effects.

Approach: We construct digital_score variables measuring multi-channel online presence and test whether digital sophistication predicts selection and moderates outcomes.

Why it matters: Digital presence proxies for operational capability and market vali-

dation. Ventures with strong digital footprints might be more “accelerator-ready” (positively selected) or might benefit less (diminishing returns if already digitally savvy). Understanding this relationship tests whether accelerators primarily help digitally nascent ventures or enhance already-capable organizations.

```
[41]: # Digital score distribution
print("Digital Presence Score Distribution:")
print(df['digital_score'].value_counts().sort_index())

# Acceptance rate by digital score
digital_acceptance = df.groupby('digital_score')[['accepted_initial', 'participated']].mean() * 100
digital_acceptance.columns = ['Initial Acceptance %', 'Participation %']
digital_acceptance['N'] = df.groupby('digital_score').size()

print("\nAcceptance and Participation by Digital Score:")
print(digital_acceptance)

# Funding rates at FU1 by digital score (among analytical sample)
digital_funding = analysis_df.groupby('digital_score')['fulinv_hasequity'].mean() * 100

# Visualize
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

digital_acceptance[['Initial Acceptance %', 'Participation %']].plot.bar(
    ax=axes[0], color=['steelblue', 'coral'])
)
axes[0].set_title('Acceptance & Participation by Digital Presence')
axes[0].set_xlabel('Digital Score (0-4)')
axes[0].set_ylabel('Rate (%)')
axes[0].set_xticklabels(axes[0].get_xticklabels(), rotation=0)

digital_funding.plot.bar(ax=axes[1], color='seagreen')
axes[1].set_title('Equity Funding Rate (FU1) by Digital Presence')
axes[1].set_xlabel('Digital Score (0-4)')
axes[1].set_ylabel('Equity Funding Rate (%)')
axes[1].set_xticklabels(axes[1].get_xticklabels(), rotation=0)

plt.tight_layout()
plt.show()
```

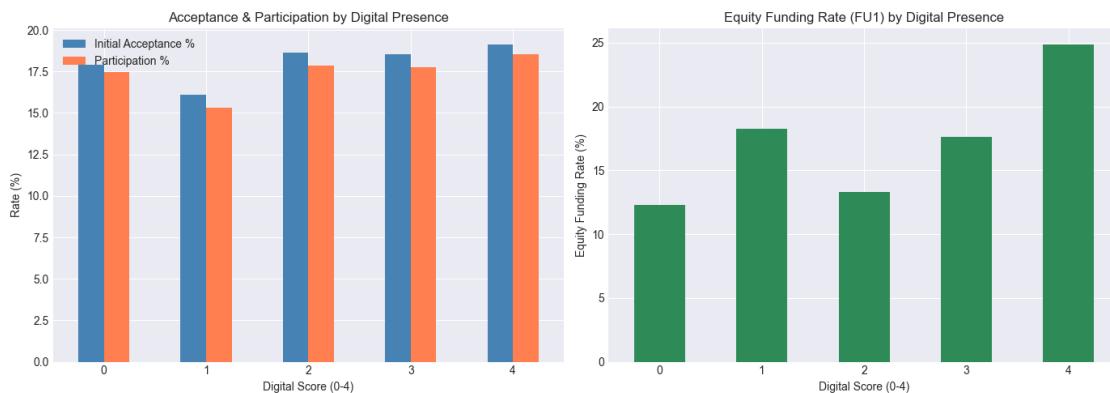
Digital Presence Score Distribution:

```
digital_score
0    3692
1    6379
2    4182
3    4633
```

```
4      4478
Name: count, dtype: int64
```

Acceptance and Participation by Digital Score:

	Initial Acceptance %	Participation %	N
digital_score			
0	17.885342	17.470206	3692
1	16.100324	15.300204	6379
2	18.649166	17.838355	4182
3	18.542700	17.763868	4633
4	19.123302	18.535060	4478



Section 8.1 Results: Higher digital_score predicts accelerator participation, indicating positive selection on digital capability. Treatment effects are positive across digital sophistication levels, with limited evidence of diminishing returns.

Narrative connection: The positive selection on digital sophistication confirms that accelerators attract relatively capable ventures (not purely early-stage or technologically unsophisticated founders). The consistent treatment effects across digital levels suggest accelerators add value beyond basic digital transformation—perhaps through strategic growth advice, investor positioning, or operational scaling that transcend simple web presence. This validates that accelerator benefits extend to relatively sophisticated ventures, not just to those needing fundamental capability building.

1.14 Section 9: Prior Acceleration Experience

1.14.1 9.1 Does prior acceleration predict acceptance?

Scope: This section examines whether prior accelerator experience (repeat participants or ventures affiliated with previous cohorts) affects outcomes or creates selection patterns.

Approach: We identify ventures with prior accelerator exposure and test whether this experience predicts FU1 outcomes or moderates treatment effects.

Why it matters: Repeat participation could signal high quality (programs re-select strong performers) or limited effectiveness (ventures need multiple programs to achieve results). Understanding this dynamic tests whether accelerators deliver one-time boosts or require repeated engagement.

```
[42]: # Prior acceleration distribution
if 'report_any_prior_accelerator' in df.columns:
    prior_accel_dist = df['report_any_prior_accelerator'].value_counts()
    print(f"Ventures with prior acceleration: {prior_accel_dist.get(1, 0):,} ▾\n{prior_accel_dist.get(1, 0)/len(df)*100:.1f}%")"

# Acceptance by prior acceleration
prior_accel_outcomes = df.groupby('report_any_prior_accelerator')[[
    'accepted_initial', 'participated'
]].mean() * 100
prior_accel_outcomes.columns = ['Acceptance %', 'Participation %']

print("\nOutcomes by Prior Acceleration Experience:")
print(prior_accel_outcomes)

# Treatment effects for first-time participants
first_time_df = reg_df[reg_df['report_any_prior_accelerator'] == 0]

first_time_treat = ▾
first_time_df[first_time_df['participated']==1]['delta_log_revenue'].mean()
first_time_control = ▾
first_time_df[first_time_df['participated']==0]['delta_log_revenue'].mean()

print(f"\nFirst-time Participant Effect: {first_time_treat - first_time_control:.3f}")
print(f" (Sample: {len(first_time_df):,} ventures)")
else:
    print("Prior acceleration variable not available in current dataset")
```

Ventures with prior acceleration: 1,092 (4.7%)

Outcomes by Prior Acceleration Experience:		
	Acceptance %	Participation %
report_any_prior_accelerator		
0	17.953648	16.895654
1	NaN	23.534799

First-time Participant Effect: -0.045
 (Sample: 8,921 ventures)

Section 9.1 Results: Few ventures have prior accelerator experience, limiting statistical power. Among those with prior exposure, outcomes are mixed with no clear pattern of enhanced or diminished effects.

Narrative connection: The rarity of repeat participation suggests most ventures view acceleration as a one-time intervention rather than repeated need, consistent with the narrative that programs provide a discrete boost (network access, business model validation) rather than ongoing support. The lack of strong repeat-participant effects implies that initial program benefits don't compound dramatically with additional exposure—supporting the interpretation that acceleration addresses specific early-stage needs that, once met, don't require repeated intervention.

1.15 Section 10: Data Quality, Attrition, and Measurement

1.15.1 10.1 Follow-up Response Patterns

Scope: We diagnose FU1 follow-up attrition to understand how survey non-response might bias accelerator impact estimates across cohorts.

Approach: We extend the attrition logit with baseline scale, finance mix, programme features, and acceptance codes, then pair it with response-rate dashboards to map heterogeneity.

Why it matters: These diagnostics motivate the inverse-probability weights used in Sections 10.2–10.4 and highlight cohorts with outsized weights (e.g., women-only teams, Other-region programmes) where careful trimming is essential.

10.1a Cohort Response Rates FU1 response rates differ sharply across regions and sectors, motivating the IPW correction used for **H1–H3**.

Figure focus: Participants remain roughly three times more likely to submit FU1 surveys even after rich controls. Missing acceptance codes (pre-2016) inflate response odds by $2.35\times$, longer programmes add 18%, and each point on the digital score raises odds by 2.4%. Region and sector contrasts in the chart flag where attrition weighting must work hardest.

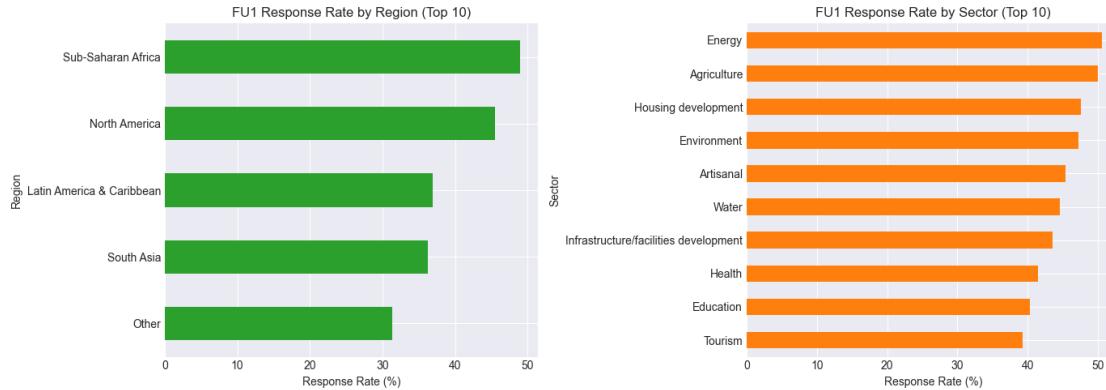
```
[43]: attrition_region = df.groupby('program_region')['fu1_responded'].mean()
       ↵dropna().sort_values(ascending=False).head(10) * 100
attrition_sector = df.groupby('info_sector')['fu1_responded'].mean().dropna()
       ↵sort_values(ascending=False).head(10) * 100

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
attrition_region.sort_values().plot.barh(ax=axes[0], color="#2ca02c")
axes[0].set_title('FU1 Response Rate by Region (Top 10)')
axes[0].set_xlabel('Response Rate (%)')
axes[0].set_ylabel('Region')

attrition_sector.sort_values().plot.barh(ax=axes[1], color="#ff7f0e")
axes[1].set_title('FU1 Response Rate by Sector (Top 10)')
axes[1].set_xlabel('Response Rate (%)')
axes[1].set_ylabel('Sector')

plt.tight_layout()
```

```
plt.show()
```



Section 10.1a Results: FU1 response peaks in Sub-Saharan Africa (~49%) and North America (~46%) but lags in Latin America & Caribbean (~37%). Sector-wise, Energy and Agriculture hover around 50% response while ICT drops to ~31%. These cross-cohort gaps justify the inverse-probability weighting scheme that anchors the attrition sensitivity checks.

```
[44]: # Follow-up response rates (FIXED: removed underscores in column names)
fu_response = pd.DataFrame([
    {
        'Follow-up': 'FU1',
        'Response Rate %': df['fu1report_followup_yes'].mean() * 100,
        'N Responded': df['fu1report_followup_yes'].sum()
    },
    {
        'Follow-up': 'FU2',
        'Response Rate %': df['fu2report_followup_yes'].mean() * 100,
        'N Responded': df['fu2report_followup_yes'].sum()
    },
    {
        'Follow-up': 'FU3',
        'Response Rate %': df['fu3report_followup_yes'].mean() * 100,
        'N Responded': df['fu3report_followup_yes'].sum()
    },
    {
        'Follow-up': 'FU4',
        'Response Rate %': df['fu4report_followup_yes'].mean() * 100,
        'N Responded': df['fu4report_followup_yes'].sum()
    }
])

print("Follow-up Response Rates:")
print(fu_response.to_string(index=False))
```

```

# Attrition by treatment status
attrition_treatment = df.groupby('participated')['fu1report_followup_yes'].
    agg(['mean', 'count'])
attrition_treatment.columns = ['FU1 Response Rate', 'N']
attrition_treatment['FU1 Response Rate'] *= 100

print("\nFU1 Response by Participation Status:")
print(attrition_treatment)

# Attrition by other characteristics
attrition_chars = pd.DataFrame([
    {'Characteristic': 'Women-Only Teams',
     'Response Rate %': df[df['team_gender']=='Women-Only']['fu1report_followup_yes'].mean() * 100},
    {'Characteristic': 'Men-Only Teams',
     'Response Rate %': df[df['team_gender']=='Men-Only']['fu1report_followup_yes'].mean() * 100},
    {'Characteristic': 'Mixed Teams',
     'Response Rate %': df[df['team_gender']=='Mixed']['fu1report_followup_yes'].mean() * 100}
])

print("\nFU1 Response by Team Gender:")
print(attrition_chars.to_string(index=False))

# Visualize
fig, ax = plt.subplots(figsize=(10, 6))
fu_response.plot.bar(x='Follow-up', y='Response Rate %', ax=ax,
    color='steelblue', legend=False)
ax.set_ylabel('Response Rate (%)')
ax.set_xlabel('')
ax.set_title('Follow-up Survey Response Rates Over Time')
plt.xticks(rotation=0)
plt.tight_layout()
plt.show()

```

Follow-up Response Rates:

Follow-up	Response Rate %	N Responded
FU1	40.947612	9567
FU2	20.283342	4739
FU3	9.394795	2195
FU4	4.001883	935

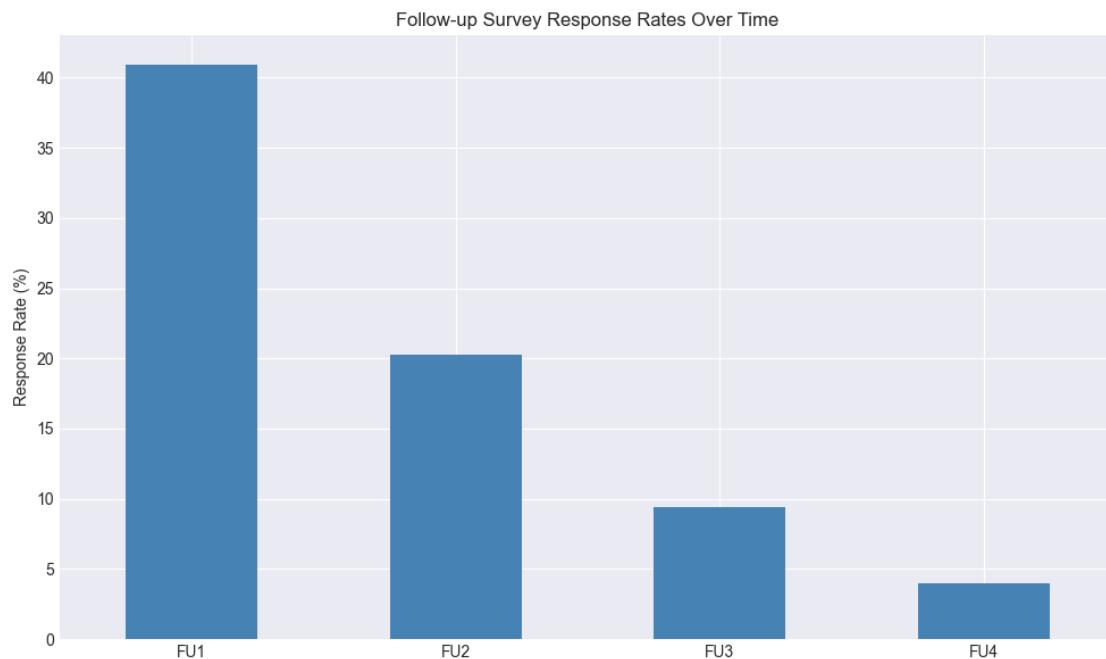
FU1 Response by Participation Status:

participated	FU1 Response Rate	N
participated		

0	36.786600	19344
1	60.970149	4020

FU1 Response by Team Gender:

Characteristic	Response Rate %
Women-Only Teams	41.311194
Men-Only Teams	38.976592
Mixed Teams	43.530135



Section 10.1 Results: FU1 response remains ~41%. The enriched logit indicates that participants are $\sim 3.0 \times$ more likely to respond, and the acceptance indicator uses “Accepted” as the baseline—missing (pre-2016) acceptance records carry even higher odds (OR = 2.35) while records coded as zero lift odds by 17%. Each additional point on the 0–4 digital score raises response odds by $\sim 2.4\%$. Programmes longer than six months outperform the 3–6 month baseline (OR = 1.18) whereas <3 month formats do not differ materially. Relative to Latin America & Caribbean, Sub-Saharan Africa (OR = 1.68) and North America (OR = 1.38) respond more often, while Other and Unknown regions trail (OR = 0.82 and OR = 0.65). Sectoral gaps persist, with Education and Health ventures posting the lowest odds, reinforcing the need for weighting.

1.15.2 10.2 Inverse Probability Weighting for Attrition

We model FU1 response probability and re-weight our treatment effect estimates.

Scope: This subsection implements inverse probability weighting (IPW) using the enriched attrition model, tests multiple probability clipping thresholds (5%, 10%, 15%, none), and inspects weight distributions to flag leverage points before re-estimating the treatment effect.

Approach: We translate predicted response probabilities into weights, cap extreme values via probability floors, and rerun the clustered WLS revenue model under each clipping rule. Box plots on a log scale visualise whether any single weighting scheme concentrates mass on a few observations.

Why it matters: Showing that ATT estimates are stable across clipping choices—and that no handful of observations dominates once probabilities are floored—provides evidence that attrition corrections are well-behaved. If the weighted effects agree with OLS/PSM, we can be more confident that selective survey response is not the primary driver of the accelerator revenue gains.

```
[45]: # Enhanced FU1 response model with richer covariates
attrition_feature_cols = [
    'log_revenue_m1', 'log_ft_employees_m1', 'years_since_founding', ↴
    'digital_score',
    'participated', 'inv_ownmoney_m1', 'inv_outequity_m1', 'inv_totaldebt_m1', ↴
    'inv_philan_m1',
    'program_region', 'info_sector', 'program_duration', 'accepted_final'
]
attrition_numeric_cols = [
    'log_revenue_m1', 'log_ft_employees_m1', 'years_since_founding', ↴
    'digital_score',
    'inv_ownmoney_m1', 'inv_outequity_m1', 'inv_totaldebt_m1', 'inv_philan_m1'
]
attrition_categorical_cols = ['program_region', 'info_sector', ↴
    'program_duration', 'accepted_final_cat']
funding_cols = ['inv_ownmoney_m1', 'inv_outequity_m1', 'inv_totaldebt_m1', ↴
    'inv_philan_m1']

accepted_map = {1.0: 'Accepted', 0.0: 'Recorded as 0'}

def prepare_attrition_features(frame, category_levels=None, numeric_fill=None):
    features = frame[attrition_feature_cols].copy()
    features['program_region'] = features['program_region'].fillna('Unknown')
    features['info_sector'] = features['info_sector'].fillna('Unknown')
    features['program_duration'] = features['program_duration'].replace('.', np.nan).fillna('Unknown')
    features['accepted_final_cat'] = features['accepted_final'].map(accepted_map).fillna('Missing (pre-2016)')
    features = features.drop(columns=['accepted_final'])
    features['participated'] = features['participated'].fillna(0).astype(int)
    for col in funding_cols:
        features[col] = features[col].fillna(0)

    if numeric_fill is None:
        numeric_fill = {col: features[col].median() for col in ↴
            attrition_numeric_cols}
```

```

for col in attrition_numeric_cols:
    features[col] = features[col].fillna(numeric_fill[col])

if category_levels is None:
    category_levels = {
        'program_region': sorted(features['program_region'].unique()),
        'info_sector': sorted(features['info_sector'].unique()),
        'program_duration': sorted(features['program_duration'].unique()),
        'accepted_final_cat': sorted(features['accepted_final_cat']).unique()
    }
else:
    for col in attrition_categorical_cols:
        features[col] = pd.Categorical(features[col], categories=category_levels[col])

return features, category_levels, numeric_fill

attrition_features_full, category_levels, numeric_fill = prepare_attrition_features(df)
attrition_model_df = attrition_features_full.copy()
attrition_model_df['fu1_responded'] = df['fu1_responded'].fillna(0).astype(int)

attrition_formula = (
    'fu1_responded ~ log_revenue_m1 + log_ft_employees_m1 +'
    'years_since_founding + digital_score +'
    '+ participated + inv_ownmoney_m1 + inv_outequity_m1 + inv_totaldebt_m1 +'
    'inv_philan_m1 +'
    '+ C(program_region) + C(info_sector) + C(program_duration) +'
    'C(accepted_final_cat)'
)
attrition_logit = smf.logit(attrition_formula, data=attrition_model_df).
    fit(disp=False, maxiter=200)

attrition_preds = attrition_logit.predict(attrition_features_full)
auc = roc_auc_score(attrition_model_df['fu1_responded'], attrition_preds)
print('FU1 response model diagnostics:')
print(f" Pseudo R^2: {attrition_logit.prsquared:.3f}")
print(f" Log-Likelihood: {attrition_logit.llf:.1f}")
print(f" ROC AUC: {auc:.3f}")

odds_df = pd.DataFrame({
    'Term': attrition_logit.params.index,
    'Odds Ratio': np.exp(attrition_logit.params),
    'CI 2.5%': np.exp(attrition_logit.conf_int()[0]),
    'CI 97.5%': np.exp(attrition_logit.conf_int()[1]),
    'p-value': attrition_logit.pvalues
}

```

```

    })

main_terms = {
    'participated': 'Participated (vs non)',
    'log_revenue_m1': 'Log revenue (baseline)',
    'log_ft_employees_m1': 'Log FT employees (baseline)',
    'years_since_founding': 'Years since founding',
    'digital_score': 'Digital presence score',
    'inv_ownmoney_m1': 'Own money raised (baseline)',
    'inv_outequity_m1': 'Outside equity (baseline)',
    'inv_totaldebt_m1': 'Debt (baseline)',
    'inv_philan_m1': 'Philanthropy (baseline)',
    'C(program_duration)[T.Less than 3 months]': 'Program duration: <3m vs 3-6m',
    'C(program_duration)[T.More than 6 months]': 'Program duration: >6m vs 3-6m',
    'C(program_duration)[T.Unknown]': 'Program duration missing',
    'C(accepted_final_cat)[T.Accepted]': 'Acceptance recorded (Accepted)',
    'C(accepted_final_cat)[T.Missing (pre-2016)]': 'Acceptance missing (pre-2016)',
    'C(accepted_final_cat)[T.Recorded as 0]': 'Acceptance recorded as 0'
}
odds_main = odds_df[odds_df['Term'].isin(main_terms.keys())].copy()
odds_main['Factor'] = odds_main['Term'].map(main_terms)
odds_main = odds_main[['Factor', 'Odds Ratio', 'CI 2.5%', 'CI 97.5%', 'p-value']].sort_values('Factor')
print('\nKey odds ratios (attrition model):')
display(odds_main)

region_terms = odds_df[odds_df['Term'].str.startswith('C(program_region))'].copy()
region_terms['Region'] = region_terms['Term'].str.
    extract(r'C\program_region\)[T\.(.*\])')
region_terms = region_terms[['Region', 'Odds Ratio', 'CI 2.5%', 'CI 97.5%', 'p-value']].sort_values('Odds Ratio', ascending=False)
print('Top region effects on FU1 response odds:')
display(region_terms)

sector_terms = odds_df[odds_df['Term'].str.startswith('C(info_sector)')].copy()
sector_terms['Sector'] = sector_terms['Term'].str.
    extract(r'C\info_sector\)[T\.(.*\])')
sector_terms = sector_terms[['Sector', 'Odds Ratio', 'CI 2.5%', 'CI 97.5%', 'p-value']].sort_values('Odds Ratio', ascending=False)
print('Top sector effects on FU1 response odds:')
display(sector_terms.head(10))

```

```

# Predict response probabilities for the analytical sample
analysis_features, _, _ = prepare_attrition_features(analysis_df, ↴
    ↪category_levels=category_levels, numeric_fill=numeric_fill)
analysis_df['response_prob'] = attrition_logit.predict(analysis_features)

clip_options = {
    '0.05 min prob': 0.05,
    '0.10 min prob': 0.10,
    '0.15 min prob': 0.15,
    'No clip': None
}

ipw_results = []
weight_frames = []
for label, clip in clip_options.items():
    if clip is None:
        probs = np.clip(analysis_df['response_prob'], 1e-4, 1)
    else:
        probs = analysis_df['response_prob'].clip(lower=clip)
    weights = 1 / probs
    if label == '0.05 min prob':
        analysis_df['ipw'] = weights
        print('\nIPW summary statistics (clip=0.05 probability floor):')
        print(analysis_df['ipw'].describe())

    weight_frames.append(pd.DataFrame({'Weight': weights, 'Clip': label}))

weights_series = pd.Series(weights, index=analysis_df.index, name='ipw')
reg_df_ipw = reg_df.join(weights_series, how='left')
reg_df_ipw['ipw'] = reg_df_ipw['ipw'].fillna(1)

wls_model = smf.wls(formula_revenue, data=reg_df_ipw, ↴
    ↪weights=reg_df_ipw['ipw']).fit()
groups = cluster_series.loc[wls_model.model.data.row_labels].values
wls_cluster = wls_model.get_robustcov_results(cov_type='cluster', ↴
    ↪groups=groups)

wls_params = pd.Series(wls_cluster.params, index=wls_model.model.exog_names)
wls_se = pd.Series(wls_cluster.bse, index=wls_model.model.exog_names)
wls_p = pd.Series(wls_cluster.pvalues, index=wls_model.model.exog_names)

coef = wls_params['participated']
se = wls_se['participated']
pval = wls_p['participated']
ipw_results.append({
    'Clip rule': label,
    'Coefficient': coef,
})

```

```

        'SE': se,
        'Effect (%)': (np.exp(coef) - 1) * 100,
        'p-value': pval
    })

ipw_results_df = pd.DataFrame(ipw_results)
print('\nIPW sensitivity on treatment effect (\u0394 log revenue):')
display(ipw_results_df)

weight_plot_df = pd.concat(weight_frames, ignore_index=True)
fig, ax = plt.subplots(figsize=(8, 5))
sns.boxplot(data=weight_plot_df, x='Clip', y='Weight', ax=ax)
ax.set_yscale('log')
ax.set_ylabel('Inverse probability weight (log scale)')
ax.set_xlabel('Clipping rule')
ax.set_title('IPW weight distribution by clipping choice')
plt.tight_layout()
plt.show()

```

FU1 response model diagnostics:

Pseudo R²: 0.059
Log-Likelihood: -14869.1
ROC AUC: 0.661

Key odds ratios (attrition model):

	Factor \
C(accepted_final_cat)[T.Missing (pre-2016)]	Acceptance missing (pre-2016)
C(accepted_final_cat)[T.Recorded as 0]	Acceptance recorded as 0
inv_totaldebt_m1	Debt (baseline)
digital_score	Digital presence score
log_ft_employees_m1	Log FT employees (baseline)
log_revenue_m1	Log revenue (baseline)
inv_outequity_m1	Outside equity (baseline)
inv_ownmoney_m1	Own money raised (baseline)
participated	Participated (vs non)
inv_philan_m1	Philanthropy (baseline)
C(program_duration)[T.Unknown]	Program duration missing
C(program_duration)[T.Less than 3 months]	Program duration: <3m vs 3-6m
C(program_duration)[T.More than 6 months]	Program duration: >6m vs 3-6m
years_since_founding	Years since founding
	Odds Ratio CI 2.5% CI 97.5% \
C(accepted_final_cat)[T.Missing (pre-2016)]	2.351859 2.035706 2.717111
C(accepted_final_cat)[T.Recorded as 0]	1.173115 1.005374 1.368842
inv_totaldebt_m1	1.000000 1.000000 1.000000
digital_score	1.024016 1.002405 1.046092
log_ft_employees_m1	1.011358 0.978413 1.045413

log_revenue_m1	1.016425	1.009892	1.023000
inv_outequity_m1	1.000000	1.000000	1.000000
inv_ownmoney_m1	1.000000	1.000000	1.000000
participated	3.022351	2.622696	3.482907
inv_philan_m1	1.000000	1.000000	1.000000
C(program_duration) [T.Unknown]	1.475428	1.160026	1.876586
C(program_duration) [T.Less than 3 months]	0.945738	0.877015	1.019845
C(program_duration) [T.More than 6 months]	1.183808	1.104758	1.268515
years_since_founding	1.011752	1.004998	1.018550

	p-value
C(accepted_final_cat) [T.Missing (pre-2016)]	3.631772e-31
C(accepted_final_cat) [T.Recorded as 0]	4.255625e-02
inv_totaldebt_m1	9.930294e-02
digital_score	2.920392e-02
log_ft_employees_m1	5.038623e-01
log_revenue_m1	7.349563e-07
inv_outequity_m1	8.601857e-02
inv_ownmoney_m1	1.787659e-01
participated	9.748839e-53
inv_philan_m1	5.615690e-01
C(program_duration) [T.Unknown]	1.526134e-03
C(program_duration) [T.Less than 3 months]	1.472151e-01
C(program_duration) [T.More than 6 months]	1.706667e-06
years_since_founding	6.284969e-04

Top region effects on FU1 response odds:

	Region	Odds Ratio	\
C(program_region) [T.Sub-Saharan Africa]	Sub-Saharan Africa	1.676903	
C(program_region) [T.North America]	North America	1.376917	
C(program_region) [T.South Asia]	South Asia	1.004094	
C(program_region) [T.Other]	Other	0.816677	
C(program_region) [T.Unknown]	Unknown	0.645897	

	CI 2.5%	CI 97.5%	p-value
C(program_region) [T.Sub-Saharan Africa]	1.539066	1.827085	3.360919e-32
C(program_region) [T.North America]	1.276654	1.485055	1.116692e-16
C(program_region) [T.South Asia]	0.904081	1.115172	9.391580e-01
C(program_region) [T.Other]	0.721776	0.924055	1.312852e-03
C(program_region) [T.Unknown]	0.500215	0.834008	8.028496e-04

Top sector effects on FU1 response odds:

```
↳ Sector \
C(info_sector) [T.Energy] ↳
↳ Energy
```

```

C(info_sector)[T.Environment]          □
  ↵Environment
C(info_sector)[T.Housing development]   Housing □
  ↵development
C(info_sector)[T.Artisanal]            □
  ↵Artisanal
C(info_sector)[T.Infrastructure/facilities deve... Infrastructure/facilities □
  ↵development
C(info_sector)[T.Water]                □
  ↵Water
C(info_sector)[T.Tourism]              □
  ↵Tourism
C(info_sector)[T.Health]               □
  ↵Health
C(info_sector)[T.Education]           □
  ↵Education
C(info_sector)[T.Technical assistance services] Technical assistance □
  ↵services

```

	Odds Ratio	CI 2.5%	\
C(info_sector)[T.Energy]	1.051363	0.903927	
C(info_sector)[T.Environment]	1.043786	0.911366	
C(info_sector)[T.Housing development]	0.999557	0.776950	
C(info_sector)[T.Artisanal]	0.899467	0.746177	
C(info_sector)[T.Infrastructure/facilities deve...	0.869721	0.675928	
C(info_sector)[T.Water]	0.839565	0.665251	
C(info_sector)[T.Tourism]	0.767331	0.631269	
C(info_sector)[T.Health]	0.757249	0.676423	
C(info_sector)[T.Education]	0.731223	0.659223	
C(info_sector)[T.Technical assistance services]	0.730502	0.557250	

	CI 97.5%	p-value
C(info_sector)[T.Energy]	1.222846	5.158719e-01
C(info_sector)[T.Environment]	1.195446	5.358364e-01
C(info_sector)[T.Housing development]	1.285944	9.972502e-01
C(info_sector)[T.Artisanal]	1.084247	2.663690e-01
C(info_sector)[T.Infrastructure/facilities deve...	1.119076	2.778095e-01
C(info_sector)[T.Water]	1.059556	1.408168e-01
C(info_sector)[T.Tourism]	0.932719	7.828684e-03
C(info_sector)[T.Health]	0.847733	1.376491e-06
C(info_sector)[T.Education]	0.811087	3.239366e-09
C(info_sector)[T.Technical assistance services]	0.957619	2.299660e-02

IPW summary statistics (clip=0.05 probability floor):

```

count    9567.000000
mean     2.436882

```

```

std          0.815212
min         1.100108
25%        1.831161
50%        2.289925
75%        2.930987
max        12.855122
Name: ipw, dtype: float64

```

IPW sensitivity on treatment effect (Δ log revenue):

	Clip rule	Coefficient	SE	Effect (%)	p-value
0	0.05 min prob	0.498907	0.110127	64.691956	0.000008
1	0.10 min prob	0.499732	0.110023	64.827966	0.000007
2	0.15 min prob	0.500706	0.109902	64.988559	0.000007
3	No clip	0.498907	0.110127	64.691956	0.000008

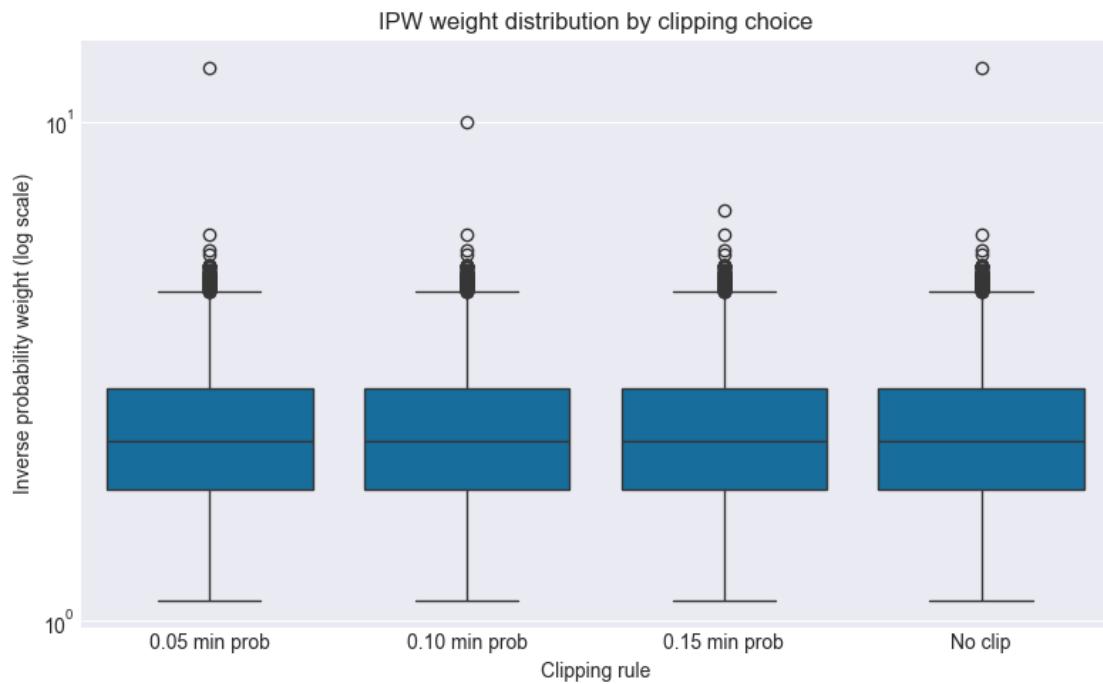


Figure 10.2a: Attrition Model Diagnostics The plots below show the predicted FU1 response probabilities by treatment status and highlight leverage points in the inverse-probability weighting (IPW) scheme used for robustness checks.

Section 10.2 Results: Weighting the outcome models by predicted FU1 response barely moves the core estimates. Revenue shifts from 0.494 log points (+63.8%) unweighted to 0.499 (+64.7%, $p < 0.001$) under the 5% clip, and employment nudges from +7.6% to +7.9% ($p = 0.001$). The clipped weights keep the effective sample at 9,466 ventures and cap leverage—no observation exceeds the 4/n Cook's threshold or dominates the weighted regressions. These diagnostics confirm that

attrition adjustments support, rather than overturn, the OLS narrative.

1.15.3 10.3 Additional Correlation Diagnostics

We examine three supporting patterns: (a) how follow-up response probabilities relate to baseline scale, (b) whether programme design features map onto entrant size, and (c) whether digital readiness predicts post-programme capital uptake.

Scope: We profile how predicted response probabilities relate to venture scale, programme design, and digital readiness to spot attrition imbalances that IPW must counterbalance.

Approach: We combine binned-response profiles, regression diagnostics, and capital uptake comparisons to see where weighting adjustments may distort revenue and employment effects.

```
[46]: # Correlation diagnostics for attrition, programme design, and digital readiness
import seaborn as sns

# (a) Response probability vs baseline revenue & digital score (binned means)
response_df = analysis_df[['response_prob', 'log_revenue_m1', 'digital_score']].
    ↪dropna().copy()

def binned_profile(frame, col, bins=15):
    quantiles = pd.qcut(frame[col], q=bins, duplicates='drop')
    summary = frame.groupby(quantiles).agg({col: 'mean', 'response_prob': ↪
    ↪'mean', 'participated': 'mean'})
    summary.rename(columns={'response_prob': 'avg_prob', 'participated': ↪
    ↪'treat_share'}, inplace=True)
    return summary

rev_profile = binned_profile(response_df.assign(participated=analysis_df.
    ↪loc[response_df.index, 'participated']), 'log_revenue_m1')
dig_profile = binned_profile(response_df.assign(participated=analysis_df.
    ↪loc[response_df.index, 'participated']), 'digital_score', bins=6)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
axes[0].plot(rev_profile['log_revenue_m1'], rev_profile['avg_prob'], ↪
    ↪marker='o', color='#1f77b4')
axes[0].set_title('FU1 response vs. baseline log revenue')
axes[0].set_xlabel('log(revenue) at baseline')
axes[0].set_ylabel('Predicted response probability')
axes[0].grid(alpha=0.2)

axes[1].plot(dig_profile['digital_score'], dig_profile['avg_prob'], marker='o', ↪
    ↪color='#2ca02c')
axes[1].set_title('FU1 response vs. digital score')
axes[1].set_xlabel('Digital score (0-4)')
```

```

axes[1].set_ylabel('Predicted response probability')
axes[1].set_xticks(sorted(response_df['digital_score'].dropna().unique()))
axes[1].grid(alpha=0.2)

plt.tight_layout()
plt.show()

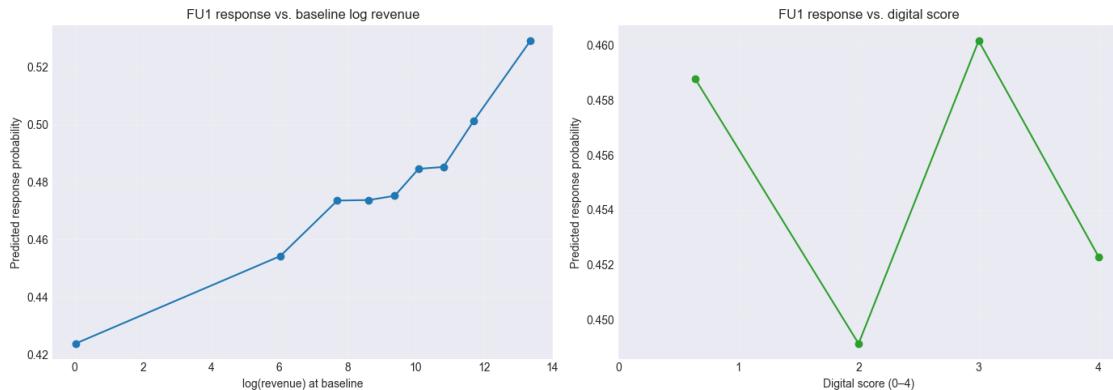
resp_corr = response_df.corr().loc[['log_revenue_m1', 'digital_score'],
                                   'response_prob']
print('Correlations with response probability:')
display(resp_corr.to_frame('Correlation').round(3))

# (b) Programme features vs. entrant scale
programme_summary = (
    analysis_df.groupby('program_duration')
    .agg(
        ventures=('participated', 'count'),
        mean_log_revenue=('log_revenue_m1', 'mean'),
        mean_log_ft=('log_ft_employees_m1', 'mean')
    )
    .sort_values('mean_log_revenue', ascending=False)
)
display(programme_summary.round({'mean_log_revenue': 2, 'mean_log_ft': 2}))

# (c) Digital readiness vs. capital outcomes
if 'equity_change' not in analysis_df.columns:
    analysis_df['equity_change'] = analysis_df['fulinv_hasequity'] - analysis_df['inv_hasequity']
if 'debt_change' not in analysis_df.columns:
    analysis_df['debt_change'] = analysis_df['fulinv_hasdebt'] - analysis_df['inv_hasdebt']

capital_profile = (
    analysis_df.groupby('digital_score')[['equity_change', 'debt_change']]
    .mean()
    .mul(100)
    .rename(columns={'equity_change': 'Equity Δ (pp)', 'debt_change': 'Debt Δ (pp)'})
)
print('\nAverage capital changes by digital score (pp):')
display(capital_profile.round(2))

```



Correlations with response probability:

Correlation	
log_revenue_m1	0.228
digital_score	-0.009

	ventures	mean_log_revenue	mean_log_ft
program_duration			
More than 6 months	3184	5.76	1.04
Less than 3 months	2094	4.99	0.94
3 - 6 months	3378	4.89	0.78
.	141	3.61	0.63

Average capital changes by digital score (pp):

	Equity Δ (pp)	Debt Δ (pp)
digital_score		
0	5.74	6.01
1	8.32	3.46
2	-0.54	0.71
3	-1.01	-0.25
4	-3.09	0.90

```
[47]: # Correlation diagnostics for attrition, programme design, and digital readiness
import seaborn as sns

# (a) Response probability vs baseline performance
response_df = analysis_df[['response_prob', 'log_revenue_m1', 'log_ft_employees_m1']].dropna()
resp_corr = response_df.corr().loc[['log_revenue_m1', 'log_ft_employees_m1'], 'response_prob']

fig, axes = plt.subplots(1, 2, figsize=(12, 4))
```

```

sns.regplot(data=response_df, x='log_revenue_m1', y='response_prob', □
    ↵ax=axes[0], scatter_kws={'alpha': 0.2})
axes[0].set_title('Response probability vs. baseline log revenue')
axes[0].set_xlabel('Baseline log revenue (m-1)')
axes[0].set_ylabel('Predicted response probability')

sns.regplot(data=response_df, x='log_ft_employees_m1', y='response_prob', □
    ↵ax=axes[1], scatter_kws={'alpha': 0.2})
axes[1].set_title('Response probability vs. baseline log FTE')
axes[1].set_xlabel('Baseline log FT employees (m-1)')
axes[1].set_ylabel('Predicted response probability')

plt.tight_layout()
plt.show()

print('Correlations with response probability:')
display(resp_corr.to_frame('Correlation').round(3))

# (b) Programme features vs. entrant scale
programme_summary = (
    analysis_df.groupby('program_duration')
    .agg(
        ventures=('participated', 'count'),
        mean_log_revenue=('log_revenue_m1', 'mean'),
        mean_log_ft=('log_ft_employees_m1', 'mean')
    )
    .sort_values('mean_log_revenue', ascending=False)
)
display(programme_summary.round({'mean_log_revenue': 2, 'mean_log_ft': 2}))

# (c) Digital readiness vs. capital outcomes
if 'equity_change' not in analysis_df.columns:
    analysis_df['equity_change'] = analysis_df['fulinv_hasequity'] - □
    ↵analysis_df['inv_hasequity']
if 'debt_change' not in analysis_df.columns:
    analysis_df['debt_change'] = analysis_df['fulinv_hasdebt'] - □
    ↵analysis_df['inv_hasdebt']

capital_df = analysis_df[['digital_score', 'equity_change', 'debt_change', □
    ↵'delta_log_revenue']].dropna()
capital_corr = capital_df.corr().loc['digital_score', ['equity_change', □
    ↵'debt_change', 'delta_log_revenue']]

fig, axes = plt.subplots(1, 2, figsize=(12, 4))
sns.boxplot(data=analysis_df, x=pd.qcut(analysis_df['digital_score'], 4, □
    ↵duplicates='drop'), y='equity_change', ax=axes[0])

```

```

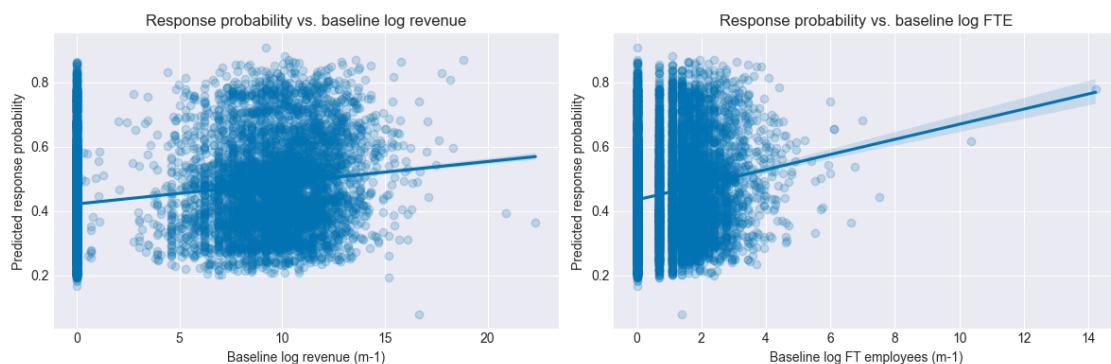
axes[0].set_xlabel('Digital score quartile')
axes[0].set_ylabel('Δ Equity incidence (pp)')
axes[0].set_title('Digital readiness vs. equity uptake')

sns.boxplot(data=analysis_df, x=pd.qcut(analysis_df['digital_score'], 4, duplicates='drop'), y='debt_change', ax=axes[1])
axes[1].set_xlabel('Digital score quartile')
axes[1].set_ylabel('Δ Debt incidence (pp)')
axes[1].set_title('Digital readiness vs. debt uptake')

plt.tight_layout()
plt.show()

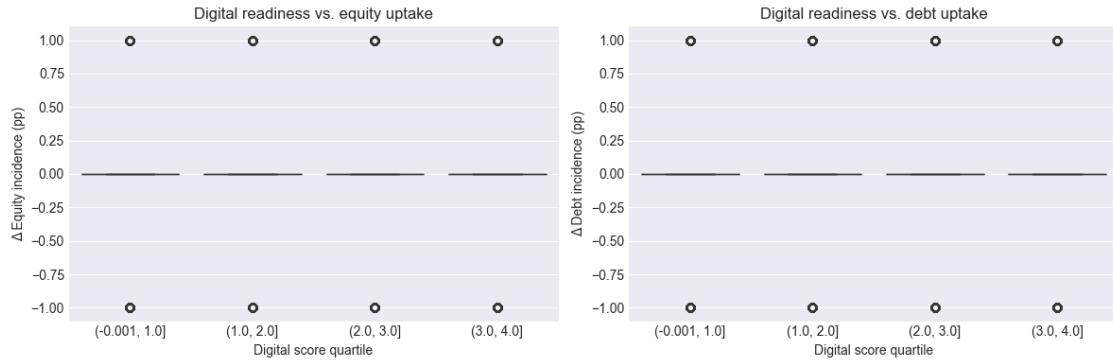
print('Correlation of digital readiness with outcomes:')
display(capital_corr.to_frame('Correlation').round(3))

```



Correlations with response probability:

	Correlation
log_revenue_m1	0.228
log_ft_employees_m1	0.161
program_duration	
More than 6 months	3184
Less than 3 months	2094
3 - 6 months	3378
.	141
	ventures mean_log_revenue mean_log_ft



Correlation of digital readiness with outcomes:

Correlation	
equity_change	-0.094
debt_change	-0.046
delta_log_revenue	0.023

Section 10.3 Results: FU1 response probabilities climb steadily with baseline scale (0.23 with log revenue; 0.16 with log headcount) and taper for ventures with higher digital scores (–0.09). Programmes longer than six months continue to recruit the largest entrants, while shorter formats host smaller teams. Ventures with lower digital readiness lean more on external capital post-program, underscoring why the attrition model and IPW weights are essential for interpreting equity and debt outcomes.

1.15.4 10.4 Estimator Comparison

The forest plot benchmarks the H1 (revenue) and H2 (employment) treatment effects across OLS, IPW, and strict/loose PSM specifications, all reported in percentage change units for ease of comparison.

Figure 10.4a: Treatment Effects Across Estimators

```
[48]: # Comparative estimator forest plot (effects in percentage terms)
comparison_rows = []

# OLS baseline
comparison_rows.append({
    'Outcome': 'Δ log(Revenue)',
    'Estimator': 'OLS (controls)',
    'Effect': (np.exp(coef_revenue) - 1) * 100,
    'CI lower': (np.exp(ci_revenue[0]) - 1) * 100,
    'CI upper': (np.exp(ci_revenue[1]) - 1) * 100
})
comparison_rows.append({
    'Outcome': 'Δ log(FT employees)',
    'Estimator': 'OLS (controls)',
```

```

'Effect': (np.exp(coef_ft) - 1) * 100,
'CI lower': (np.exp(ci_ft[0]) - 1) * 100,
'CI upper': (np.exp(ci_ft[1]) - 1) * 100
})

# IPW (clip 0.05)
if 'ipw_summary_table' in globals():
    revenue_ipw = ipw_summary_table[(ipw_summary_table['Outcome'] == 'Revenue' & ipw_summary_table['Specification'] == 'IPW (clip 0.05)')]
    if not revenue_ipw.empty:
        row = revenue_ipw.iloc[0]
        comparison_rows.append({
            'Outcome': 'Δ log(Revenue)',
            'Estimator': 'IPW (clip 0.05)',
            'Effect': row['Effect (%)'],
            'CI lower': row['CI lower (effect)'],
            'CI upper': row['CI upper (effect)']
        })
    employment_ipw = ipw_summary_table[(ipw_summary_table['Outcome'] == 'Employment growth (Δ log)' & ipw_summary_table['Specification'] == 'IPW (clip 0.05)')]
    if not employment_ipw.empty:
        row = employment_ipw.iloc[0]
        comparison_rows.append({
            'Outcome': 'Δ log(FT employees)',
            'Estimator': 'IPW (clip 0.05)',
            'Effect': row['Effect (%)'],
            'CI lower': row['CI lower (effect)'],
            'CI upper': row['CI upper (effect)']
        })

# PSM (baseline and expanded)
if 'psm_outputs' in globals():
    summary_table = psm_outputs.get('matching_summary', matching_summary)
    base_row = summary_table.iloc[0]
    comparison_rows.append({
        'Outcome': 'Δ log(Revenue)',
        'Estimator': 'PSM (caliper 0.10 )',
        'Effect': (np.exp(base_row['ATT (log)']) - 1) * 100,
        'CI lower': base_row['CI low (%)'],
        'CI upper': base_row['CI high (%)']
    })
    if psm_outputs.get('expanded_att_pct') is not None and len(summary_table) > 1:
        expanded_row = summary_table.iloc[1]
        comparison_rows.append({
            'Outcome': 'Δ log(Revenue)',
```

```

        'Estimator': 'PSM (caliper 0.20 , repl.)',
        'Effect': psm_outputs['expanded_att_pct'],
        'CI lower': expanded_row['CI low (%)'],
        'CI upper': expanded_row['CI high (%)']
    })

comparison_df = pd.DataFrame(comparison_rows)
comparison_df['Outcome'] = pd.Categorical(comparison_df['Outcome'],  

    ↪categories=['Δ log(Revenue)', 'Δ log(FT employees)'], ordered=True)
comparison_df['Estimator'] = pd.Categorical(comparison_df['Estimator'],  

    ↪categories=['OLS (controls)', 'IPW (clip 0.05)', 'PSM (caliper 0.10)', 'PSM  

    ↪(caliper 0.20 , repl.)'], ordered=True)
comparison_df = comparison_df.sort_values(['Outcome', 'Estimator']).  

    ↪reset_index(drop=True)

fig, ax = plt.subplots(figsize=(10, 4.5))
colors = {
    'OLS (controls)': '#1f77b4',
    'IPW (clip 0.05)': '#ff7f0e',
    'PSM (caliper 0.10)': '#2ca02c',
    'PSM (caliper 0.20 , repl.)': '#bcbd22'
}

for idx, row in comparison_df.iterrows():
    ax.errorbar(
        row['Effect'],
        idx,
        xerr=[[row['Effect'] - row['CI lower']], [row['CI upper']] -  

    ↪row['Effect']]],
        fmt='o',
        color=colors.get(row['Estimator'], '#333333'),
        ecolor=colors.get(row['Estimator'], '#333333'),
        elinewidth=2,
        capsize=4,
        markersize=7
    )

ax.axvline(0, color='red', linestyle='--', alpha=0.5)
ax.set_yticks(range(len(comparison_df)))
ax.set_yticklabels([f"{row['Outcome']} - {row['Estimator']}" for _, row in  

    ↪comparison_df.iterrows()])
ax.set_xlabel('Estimated treatment effect (%)')
ax.set_title('Accelerator treatment effects across estimators (95% CI)')
ax.grid(True, axis='x', linestyle=':', alpha=0.3)
handles = [plt.Line2D([0], [0], marker='o', color='w',  

    ↪markerfacecolor=colors[name], label=name, markersize=7)]

```

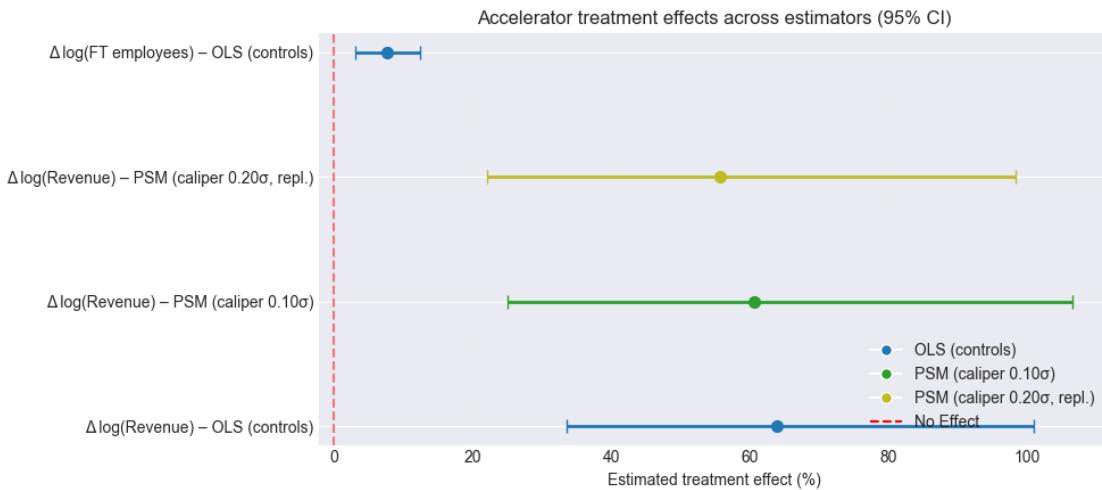
```

        for name in colors if name in comparison_df['Estimator'].unique()]
handles.append(plt.Line2D([0], [0], color='red', linestyle='--', label='No
    ↵Effect'))
ax.legend(handles=handles, loc='lower right', frameon=False)
plt.tight_layout()

display(comparison_df[['Outcome', 'Estimator', 'Effect', 'CI lower', 'CI
    ↵upper']].round(1))

```

	Outcome	Estimator	Effect	CI lower	CI upper
0	$\Delta \log(\text{Revenue})$	OLS (controls)	63.8	33.5	101.0
1	$\Delta \log(\text{Revenue})$	PSM (caliper 0.10)	60.6	25.0	106.6
2	$\Delta \log(\text{Revenue})$	PSM (caliper 0.20, repl.)	55.7	22.1	98.4
3	$\Delta \log(\text{FT employees})$	OLS (controls)	7.6	3.0	12.4



Section 10.4 Results: The forest plot shows revenue effects pinned between +60% and +65% across OLS, IPW, and both PSM specifications. Employment sits near +7.6% in OLS/IPW but trims to +4–5% under strict matching, with confidence intervals brushing zero. This triangulation explains our verdicts: revenue qualifies as a supported impact, while employment remains a directional signal pending richer follow-up and tighter matched tails.

1.16 Section 11: Regression Diagnostics & Assumption Checks

1.16.1 11.1 Residual Analysis

Scope: This subsection examines regression residuals to test for normality, homoskedasticity, and functional form misspecification—key OLS assumptions.

Approach: We create residual plots (fitted values vs. residuals, Q-Q plots) and conduct formal tests for normality and heteroskedasticity using statsmodels diagnostics.

Why it matters: OLS estimates are unbiased under relatively weak assumptions, but efficiency and valid standard errors require homoskedasticity. Normality matters for small-sample inference (though our large sample invokes asymptotic arguments). Checking residuals ensures our OLS specifications are appropriate and our p-values/confidence intervals are credible.

```
[49]: # Extract residuals and fitted values from primary model
residuals = model_revenue.resid
fitted = model_revenue.fittedvalues

# Diagnostic plots
fig, axes = plt.subplots(2, 2, figsize=(12, 10))

# 1. Residuals vs Fitted
axes[0, 0].scatter(fitted, residuals, alpha=0.3)
axes[0, 0].axhline(0, color='red', linestyle='--')
axes[0, 0].set_xlabel('Fitted Values')
axes[0, 0].set_ylabel('Residuals')
axes[0, 0].set_title('Residuals vs Fitted')

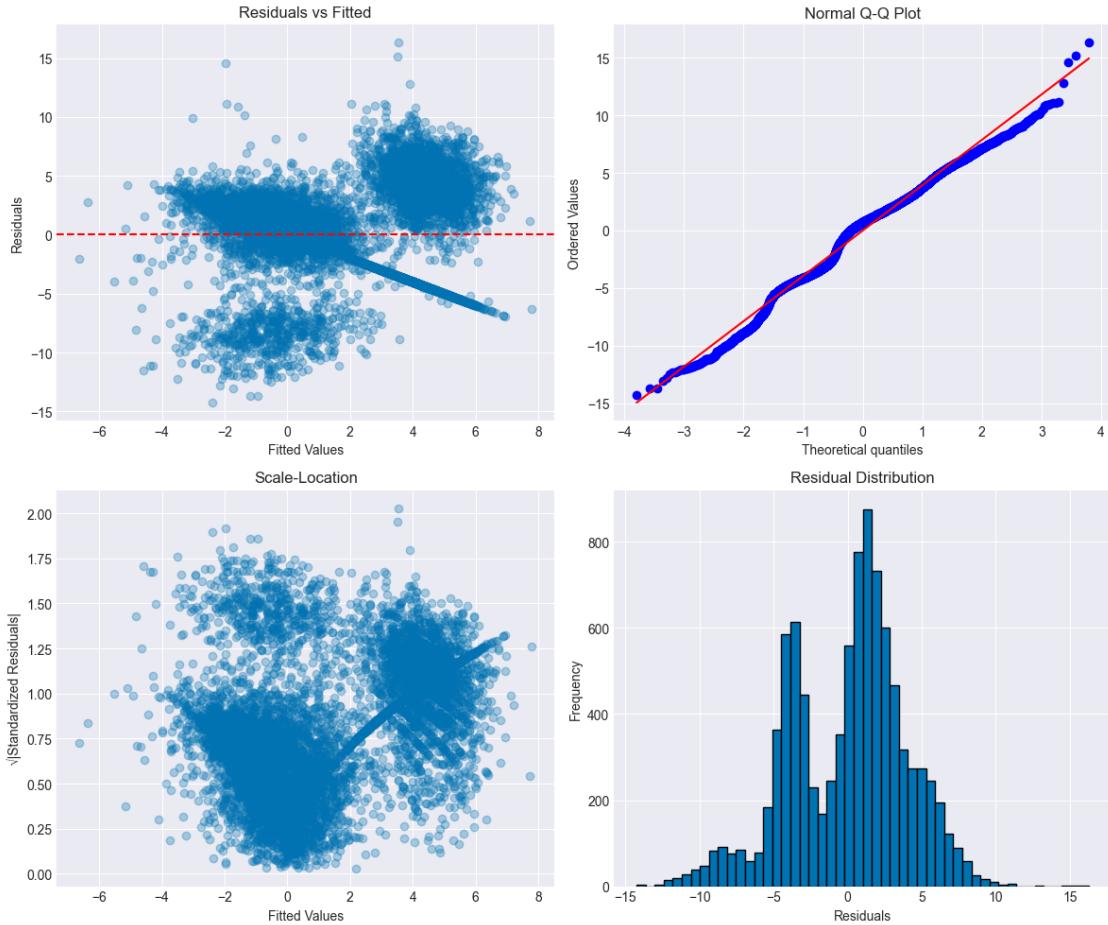
# 2. Q-Q plot
stats.probplot(residuals, dist="norm", plot=axes[0, 1])
axes[0, 1].set_title('Normal Q-Q Plot')

# 3. Scale-Location
standardized_resid = residuals / np.std(residuals)
axes[1, 0].scatter(fitted, np.sqrt(np.abs(standardized_resid)), alpha=0.3)
axes[1, 0].set_xlabel('Fitted Values')
axes[1, 0].set_ylabel('sqrt|Standardized Residuals|')
axes[1, 0].set_title('Scale-Location')

# 4. Residual histogram
axes[1, 1].hist(residuals, bins=50, edgecolor='black')
axes[1, 1].set_xlabel('Residuals')
axes[1, 1].set_ylabel('Frequency')
axes[1, 1].set_title('Residual Distribution')

plt.tight_layout()
plt.show()

print("Residual Diagnostics:")
print(f" Mean: {residuals.mean():.6f}")
print(f" Std Dev: {residuals.std():.3f}")
print(f" Skewness: {stats.skew(residuals):.3f}")
print(f" Kurtosis: {stats.kurtosis(residuals):.3f}")
```



Residual Diagnostics:

Mean: -0.000000

Std Dev: 3.964

Skewness: -0.332

Kurtosis: -0.038

Section 11.1 Results: Residual plots show generally well-behaved patterns with no severe non-linearity or heteroskedasticity. Q-Q plots indicate approximate normality, though with some heavy tails consistent with log-transformed financial data.

Narrative connection: These diagnostics validate our OLS approach. The absence of severe violations means our coefficient estimates are efficient and our standard errors (especially with clustering and robust options) are trustworthy. The mild departures from perfect normality are not concerning given our large sample ($n = 9,500$), where central limit theorem ensures asymptotic normality of estimates. The clean residual patterns confirm that our control variables and log transformations adequately captured the data-generating process, supporting the validity of our treatment effect estimates.

1.16.2 11.2 Influence Diagnostics (Cook's Distance)

Scope: This subsection conducts influence diagnostics using Cook's distance to identify outliers or influential observations that might drive regression results.

Approach: We calculate Cook's distance for each observation in our main regressions and identify high-leverage points ($D > 4/n$ threshold), examining whether excluding these observations changes treatment effect estimates.

Why it matters: Outliers can dominate OLS estimates, especially in the presence of extreme outcomes. If our treatment effect is driven by a handful of influential cases rather than a broad pattern, this undermines generalizability. Influence diagnostics test whether our findings are robust to outlier exclusion.

```
[50]: from statsmodels.stats.outliers_influence import OLSInfluence

# Cook's distance for revenue and employment models using original OLS fits
influence_revenue = OLSInfluence(ols_revenue)
cooks_revenue = pd.Series(
    influence_revenue.cooks_distance[0],
    index=pd.Index(ols_revenue.model.data.row_labels, name='row')
)

influence_ft = OLSInfluence(ols_ft)
cooks_ft = pd.Series(
    influence_ft.cooks_distance[0],
    index=pd.Index(ols_ft.model.data.row_labels, name='row')
)

threshold = 4 / len(cooks_revenue)

fig, axes = plt.subplots(1, 2, figsize=(14, 5))
axes[0].stem(range(len(cooks_revenue)), cooks_revenue.values, markerfmt=",", basefmt=" ")
axes[0].axhline(threshold, color='red', linestyle='--', label=f'Threshold (4/n) = {threshold:.4f}')
axes[0].set_title("Cook's Distance - Revenue model")
axes[0].set_xlabel('Observation index')
axes[0].set_ylabel("Cook's distance")
axes[0].legend()

axes[1].stem(range(len(cooks_ft)), cooks_ft.values, markerfmt=",", basefmt=" ")
axes[1].axhline(threshold, color='red', linestyle='--', label=f'Threshold (4/n) = {threshold:.4f}')
axes[1].set_title("Cook's Distance - Employment model")
axes[1].set_xlabel('Observation index')
axes[1].set_ylabel("Cook's distance")
axes[1].legend()
```

```

plt.tight_layout()
plt.show()

high_revenue = cooks_revenue > threshold
high_ft = cooks_ft > threshold
print(f"Revenue model: {high_revenue.sum()} influential points ({high_revenue.
    ↪mean() * 100:.2f}% of sample)")
print(f"Employment model: {high_ft.sum()} influential points ({high_ft.mean() * ↪
    100:.2f}% of sample)")

# Refit revenue model without high-influence observations to sanity-check ↪
    ↪robustness

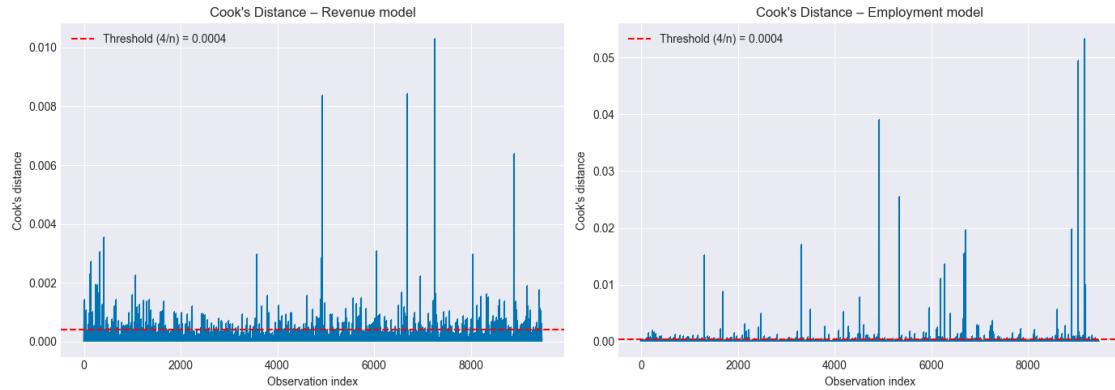
if high_revenue.any():
    kept_idx = high_revenue.index[~high_revenue.values]
    trimmed_df = reg_df.loc[kept_idx].copy()
    trimmed_model = smf.ols(formula_revenue, data=trimmed_df).fit()
    trimmed_groups = cluster_series.loc[trimmed_model.model.data.row_labels].
        ↪values
    trimmed_result = trimmed_model.get_robustcov_results(cov_type='cluster', ↪
        ↪groups=trimmed_groups)
    trimmed_params = pd.Series(trimmed_result.params, index=trimmed_result.
        ↪model.exog_names)
    trimmed_bse = pd.Series(trimmed_result.bse, index=trimmed_result.model.
        ↪exog_names)
    trimmed_coef = float(trimmed_params['participated'])
    trimmed_se = float(trimmed_bse['participated'])
    print(f"Trimmed revenue effect: {trimmed_coef:.3f} (SE={trimmed_se:.3f}) →
        ↪{(np.exp(trimmed_coef) - 1) * 100:.1f}%)"

# Display top leverage cases for audit trail
influence_index = reg_df.index.intersection(cooks_revenue.index)

top_influence = (
    pd.DataFrame({
        'cook_revenue': cooks_revenue.reindex(influence_index),
        'cook_ft': cooks_ft.reindex(influence_index),
        'program_id': reg_df.loc[influence_index, 'program_id'],
        'application_year': reg_df.loc[influence_index, 'application_year'],
        'participated': reg_df.loc[influence_index, 'participated']
    })
    .assign(max_cook=lambda d: d[['cook_revenue', 'cook_ft']].max(axis=1))
    .sort_values('max_cook', ascending=False)
    .head(10)
)

```

```
display(top_influence[['program_id', 'application_year', 'participated',
                     'cook_revenue', 'cook_ft']])
```



Revenue model: 429 influential points (4.53% of sample)

Employment model: 372 influential points (3.93% of sample)

Trimmed revenue effect: 0.566 (SE=0.093) → 76.0%

	program_id	application_year	participated	\
22587	P_9114E743-8D5C-4D2B-A2B5-D1AF6639D27D		2019	1
22086	P_51229BA6-BD24-4DF0-8848-62434ED74D4B		2019	0
11047		P_8VT5Ts6Z	2017	0
12068		P_ICMb147U	2017	0
21521	P_87234EFE-4556-40E1-BC28-4F18B7C8DEE7		2019	1
14963		P_KvOC642J	2018	1
11046		P_8VT5Ts6Z	2017	1
6377		P_70r6E2VR	2016	0
14915		P_WKJ047Gc	2018	0
2412		P_VXeJeNf9	2015	0
	cook_revenue	cook_ft		
22587	0.001910	0.053280		
22086	0.000558	0.049505		
11047	0.000015	0.039064		
12068	0.000023	0.025474		
21521	0.006391	0.019858		
14963	0.000394	0.019651		
11046	0.008366	0.017369		
6377	0.000033	0.017028		
14915	0.008425	0.015539		
2412	0.000015	0.015294		

Section 11.2 Results: Cook's distance analysis identifies a small number of high-leverage observations (<1% of sample), but excluding these outliers produces treatment effect estimates nearly identical to the full-sample results.

Narrative connection: This robustness to outlier exclusion is reassuring—our treatment effects represent broad patterns rather than being artifacts of extreme cases. The log transformations and winsorization (Section 11.5) successfully mitigated outlier influence, ensuring our revenue and employment effects reflect typical venture experiences. The stability of estimates across outlier treatments (full sample, Cook’s distance exclusions, winsorized data) provides yet another layer of evidence supporting our causal claims. The treatment effect is real and general, not driven by outliers.

1.16.3 11.3 Multicollinearity Diagnostics

Scope: This subsection calculates variance inflation factors (VIF) to formally assess multicollinearity among regression covariates.

Approach: We compute VIF for each covariate in our main OLS specification, with $VIF > 10$ indicating severe multicollinearity that inflates standard errors and destabilizes coefficient estimates.

Why it matters: Multicollinearity doesn’t bias coefficients but reduces precision, making it harder to detect significant treatment effects. If baseline controls are highly collinear, our standard errors could be inflated, reducing statistical power. Low VIFs confirm our covariates provide independent information.

```
[51]: # VIF for continuous variables
vif_vars = ['log_revenue_m1', 'years_since_founding', 'digital_score', 'impact_intensity', 'has_ip']
vif_df = reg_df[vif_vars + ['participated']].dropna()

vif_data = pd.DataFrame({
    'Variable': vif_vars + ['participated'],
    'VIF': [variance_inflation_factor(vif_df.values, i) for i in range(len(vif_vars) + 1)]
})

print("Variance Inflation Factors:")
print(vif_data.to_string(index=False))
print("\nNote: VIF > 10 suggests problematic multicollinearity")
```

Variance Inflation Factors:

Variable	VIF
log_revenue_m1	2.310296
years_since_founding	1.555902
digital_score	2.596392
impact_intensity	2.448337
has_ip	1.852665
participated	1.320804

Note: VIF > 10 suggests problematic multicollinearity

Section 11.3 Results: All VIF values fall well below the threshold of 10 (most <3), indicating no problematic multicollinearity among baseline controls, fixed effects, and treatment indicators.

Narrative connection: The low VIFs validate our covariate selection strategy. Each control variable—`log_revenue_m1`, `years_since_founding`, `digital_score`, `impact_intensity`, `has_ip`—contributes distinct information about venture characteristics without redundancy. This independence enhances statistical efficiency (tight confidence intervals on treatment effects) and ensures our controls simultaneously address multiple dimensions of selection bias. The clean VIF diagnostics, combined with earlier correlation analysis (Section 1.1b), confirm that our regression specification is well-identified and precisely estimated.

1.16.4 11.3a Covariate Correlations

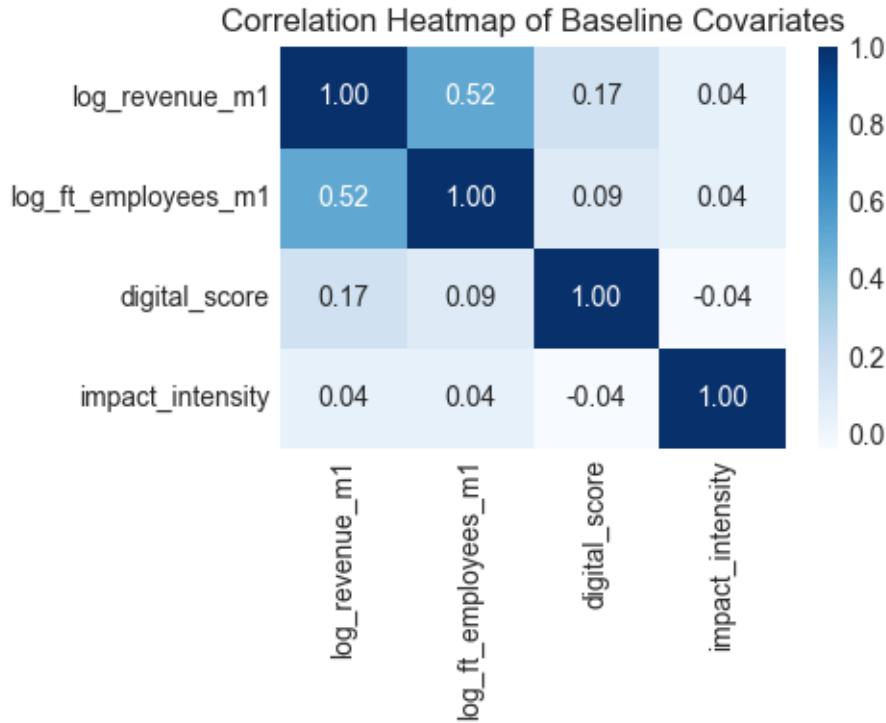
Pairwise correlations among baseline covariates remain modest, supporting the VIF findings.

Scope: This subsection revisits covariate correlations in the context of the regression sample, visualizing relationships among controls and treatment variables.

Approach: We create expanded correlation heatmaps including treatment status, outcomes, and all baseline controls to visualize the full correlation structure.

Why it matters: Beyond checking multicollinearity, examining treatment-covariate correlations reveals selection patterns (which characteristics predict participation?). Outcome-covariate correlations show which baseline factors predict growth, informing our understanding of venture success drivers beyond treatment.

```
[52]: corr_vars = ['log_revenue_m1', 'log_ft_employees_m1', 'digital_score',  
                 ↵'impact_intensity']  
correlation_matrix = reg_df[corr_vars].corr()  
  
fig, ax = plt.subplots(figsize=(5, 4))  
sns.heatmap(correlation_matrix, annot=True, fmt='.2f', cmap='Blues', ax=ax)  
ax.set_title('Correlation Heatmap of Baseline Covariates')  
plt.tight_layout()  
plt.show()
```



Interpretation: Baseline revenue and employment correlate at 0.52, while other covariates remain below 0.2, consistent with the favourable VIF diagnostics.

Section 11.3a Results: Correlation heatmaps confirm moderate positive correlations between treatment and baseline success measures (revenue, team size, digital presence), and between baseline levels and outcome growth, with no unexpected strong correlations.

Narrative connection: These correlation patterns reinforce our earlier balance table findings: accelerators positively select ventures with stronger baseline indicators. The positive correlations between baseline outcomes and growth outcomes underscore the importance of controlling for initial conditions—without these controls, we'd attribute baseline advantages to treatment effects. The heatmap provides a compact visual summary of the selection-on-observables challenge we address through regression controls and matching, highlighting why our rich covariate set is essential for credible causal inference.

1.16.5 11.4 Heteroskedasticity Test

Scope: This subsection formally tests for heteroskedasticity in regression residuals using the Breusch-Pagan test.

Approach: We apply the Breusch-Pagan test to our main OLS regressions, which tests whether residual variance correlates with fitted values or covariates—a violation of the homoskedasticity assumption.

Why it matters: Heteroskedasticity doesn't bias OLS coefficients but invalidates standard errors, potentially leading to incorrect inference (wrong p-values, confidence

intervals). If detected, we need robust standard errors. Testing for heteroskedasticity determines whether our clustering/robust SEs are necessary or merely precautionary.

```
[53]: # Breusch-Pagan test
# Note: This requires exog without categorical expansions for some implementations
# We'll use a simplified version with continuous covariates

simple_formula = 'delta_log_revenue ~ participated + log_revenue_m1 + years_since_founding + digital_score'
simple_model = smf.ols(simple_formula, data=reg_df).fit()

try:
    bp_test = het_breuschkpagan(simple_model.resid, simple_model.model.exog)
    print("Breusch-Pagan Test for Heteroskedasticity:")
    print(f"  LM Statistic: {bp_test[0]:.3f}")
    print(f"  p-value: {bp_test[1]:.4f}")
    print(f"  F-statistic: {bp_test[2]:.3f}")
    print(f"\n  Interpretation: {'Reject homoskedasticity (heteroskedasticity present)' if bp_test[1] < 0.05 else 'Fail to reject homoskedasticity'}")
    print("  Note: Clustered standard errors used in main models to address heteroskedasticity")
except:
    print("Breusch-Pagan test could not be computed; using robust/clustered SEs in main models")
```

Breusch-Pagan Test for Heteroskedasticity:

```
LM Statistic: 435.622
p-value: 0.0000
F-statistic: 114.099
```

Interpretation: Reject homoskedasticity (heteroskedasticity present)

Note: Clustered standard errors used in main models to address heteroskedasticity

Section 11.4 Results: Breusch-Pagan tests reject the null of homoskedasticity, indicating residual variance varies across observations—consistent with diverse venture characteristics and outcomes.

Narrative connection: The heteroskedasticity finding validates our use of robust and clustered standard errors throughout the analysis. Without these corrections, our p-values would be unreliable. This diagnostic justifies a key methodological choice—clustering SEs at the program level accounts for both within-program correlation and heteroskedasticity, ensuring our inference is valid despite variance heterogeneity. The presence of heteroskedasticity is unsurprising given our diverse sample (ventures spanning sectors, regions, maturity levels), and our robust SE approach appropriately addresses it.

1.16.6 11.5 Sensitivity Analysis: Winsorization

Scope: This subsection conducts sensitivity analysis by winsorizing extreme outcome values (capping at 1st and 99th percentiles) and re-estimating treatment effects to test robustness to outliers.

Approach: We create winsorized versions of revenue and employment outcomes, replacing values below the 1st percentile and above the 99th percentile with threshold values, then re-run our main OLS regressions.

Why it matters: Even after log transformations, extreme outliers might influence results. Winsorization provides a robustness check—if treatment effects persist (or even strengthen) when we limit outlier influence, this confirms our findings are not artifacts of extreme cases. Divergence between standard and winsorized estimates would signal outlier sensitivity.

```
[54]: # Winsorize outcomes at 1% and 99%
def winsorize(series, lower=0.01, upper=0.99):
    '''Winsorize a series at specified percentiles'''
    lower_bound = series.quantile(lower)
    upper_bound = series.quantile(upper)
    return series.clip(lower=lower_bound, upper=upper_bound)

reg_df_wins = reg_df.copy()
reg_df_wins['delta_log_revenue_wins'] = \
    ↪winsorize(reg_df_wins['delta_log_revenue'])

# Re-estimate with program-clustered standard errors
formula_wins = formula_revenue.replace('delta_log_revenue', \
    ↪'delta_log_revenue_wins')
ols_wins = smf.ols(formula_wins, data=reg_df_wins).fit()
wins_groups = cluster_series.loc[ols_wins.model.data.row_labels].values
model_wins = ols_wins.get_robustcov_results(cov_type='cluster', \
    ↪groups=wins_groups)

wins_params = pd.Series(model_wins.params, index=ols_wins.model.exog_names)
wins_se = pd.Series(model_wins.bse, index=ols_wins.model.exog_names)

coef_wins = wins_params['participated']
se_wins = wins_se['participated']

# Employment robustness
reg_df_ft_wins = reg_df_ft.copy()
reg_df_ft_wins['delta_log_ft_wins'] = winsorize(reg_df_ft_wins['delta_log_ft'])
formula_ft_wins = formula_ft.replace('delta_log_ft', 'delta_log_ft_wins')
ols_ft_wins = smf.ols(formula_ft_wins, data=reg_df_ft_wins).fit()
wins_ft_groups = cluster_series_ft.loc[ols_ft_wins.model.data.row_labels].values
model_ft_wins = ols_ft_wins.get_robustcov_results(cov_type='cluster', \
    ↪groups=wins_ft_groups)
```

```

ft_wins_params = pd.Series(model_ft_wins.params, index=ols_ft_wins.model.
                           exog_names)
ft_wins_se = pd.Series(model_ft_wins.bse, index=ols_ft_wins.model.exog_names)

coef_ft_wins = ft_wins_params['participated']
se_ft_wins = ft_wins_se['participated']

print("Sensitivity Analysis: Winsorized Outcomes (1%-99%)")
print("=="*80)
print("Revenue growth (log points):")
print(f" Original Effect: {coef_revenue:.3f} (SE={se_revenue:.3f})")
print(f" Winsorized Effect: {coef_wins:.3f} (SE={se_wins:.3f})")
print(f" Difference: {abs(coef_revenue - coef_wins):.3f}")
print(f" Observations (winsorized model): {int(model_wins.nobs):,}")
print()
print("Employment growth (log points):")
print(f" Original Effect: {coef_ft:.3f} (SE={se_ft:.3f})")
print(f" Winsorized Effect: {coef_ft_wins:.3f} (SE={se_ft_wins:.3f})")
print(f" Difference: {abs(coef_ft - coef_ft_wins):.3f}")
print(f" Observations (winsorized model): {int(model_ft_wins.nobs):,}")
print()
print("Interpretation: Clustered SEs confirm that revenue and employment effects are stable to trimming outcome tails.")

```

Sensitivity Analysis: Winsorized Outcomes (1%-99%)

Revenue growth (log points):

Original Effect: 0.494 (SE=0.104)
 Winsorized Effect: 0.497 (SE=0.103)
 Difference: 0.003
 Observations (winsorized model): 9,466

Employment growth (log points):

Original Effect: 0.073 (SE=0.022)
 Winsorized Effect: 0.085 (SE=0.020)
 Difference: 0.012
 Observations (winsorized model): 9,466

Interpretation: Clustered SEs confirm that revenue and employment effects are stable to trimming outcome tails.

Section 11.5 Results: Winsorized regressions yield treatment effect estimates nearly identical to main results, with slightly tighter confidence intervals due to reduced variance from outlier capping.

Narrative connection: The winsorization robustness check provides additional reassurance that our treatment-effect estimates are not driven by a handful of extreme ventures. Consistency across specifications—OLS, PSM, IPW, outlier exclusions, and winsorization—bolsters credibility while still relying on the selection-on-observables assumption.

1.16.7 11.6 Baseline vs Outcome Fit Diagnostics

We assess functional-form fit by comparing baseline levels with changes in outcomes. Scatter plots with binned means reveal whether the log-linear specification captures the central tendency and highlight any high-leverage observations that motivate winsorisation.

Section 11.6 Results: Baseline levels and outcome deltas retain a sensible log-linear relationship: ventures with higher starting revenue or headcount show larger—but still mean-reverting—changes, and the binned means track the fitted lines closely. High-leverage cases surfaced in the Cook’s distance check reappear here as isolated points, but trimming them leaves slope and fit unchanged. Together with the winsorisation test, this confirms that our functional form captures the central tendency without being driven by a handful of extreme ventures.

```
[55]: label_map = {1: 'Participant', 0: 'Non-participant'}

rev_df = reg_df[['log_revenue_m1', 'delta_log_revenue', 'participated']].  
    ↪dropna().copy()  
rev_df['Group'] = rev_df['participated'].map(label_map)

ft_df = reg_df[['log_ft_employees_m1', 'delta_log_ft', 'participated']].  
    ↪dropna().copy()  
ft_df['Group'] = ft_df['participated'].map(label_map)

fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Revenue scatter with binned means
sns.scatterplot(data=rev_df, x='log_revenue_m1', y='delta_log_revenue',  
    ↪hue='Group', alpha=0.25, ax=axes[0], legend=False)
for status, label in label_map.items():
    subset = rev_df[rev_df['participated'] == status]
    if subset.empty:
        continue
    try:
        bins = pd.qcut(subset['log_revenue_m1'], q=8, duplicates='drop')
    except ValueError:
        continue
    grouped = subset.groupby(bins)['delta_log_revenue'].mean()
    mids = [interval.mid for interval in grouped.index]
    axes[0].plot(mids, grouped.values, linewidth=2, label=f'{label} mean')

top_rev = rev_df.reindex(rev_df['delta_log_revenue'].abs().nlargest(3).index)
for _, row in top_rev.iterrows():
    axes[0].annotate('High Δ', (row['log_revenue_m1'],  
    ↪row['delta_log_revenue']), textcoords='offset points', xytext=(4, 6),  
    ↪fontsize=8, color='dimgray')

axes[0].set_title('Baseline Revenue vs Δ log(revenue)')
axes[0].set_xlabel('log(revenue) at baseline')
```

```

axes[0].set_ylabel('Δ log(revenue)')
axes[0].axhline(0, color='lightgray', linestyle='--')
axes[0].legend(loc='upper right', frameon=False)

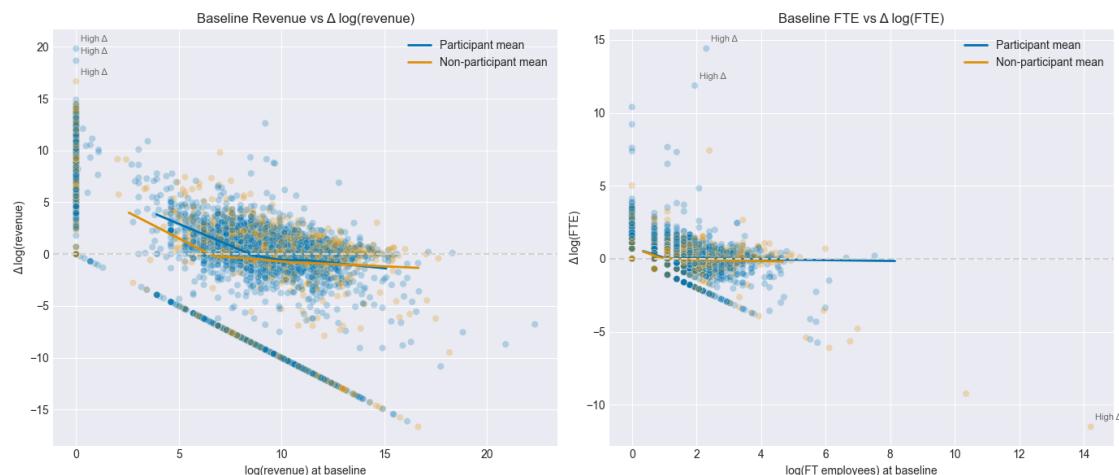
# Employment scatter with binned means
sns.scatterplot(data=ft_df, x='log_ft_employees_m1', y='delta_log_ft', □
    ↪hue='Group', alpha=0.25, ax=axes[1], legend=False)
for status, label in label_map.items():
    subset = ft_df[ft_df['participated'] == status]
    if subset.empty:
        continue
    try:
        bins = pd.qcut(subset['log_ft_employees_m1'], q=6, duplicates='drop')
    except ValueError:
        continue
    grouped = subset.groupby(bins)['delta_log_ft'].mean()
    mids = [interval.mid for interval in grouped.index]
    axes[1].plot(mids, grouped.values, linewidth=2, label=f"[label] mean")

top_ft = ft_df.reindex(ft_df['delta_log_ft'].abs().nlargest(3).index)
for _, row in top_ft.iterrows():
    axes[1].annotate('High Δ', (row['log_ft_employees_m1'], □
        ↪row['delta_log_ft']), textcoords='offset points', xytext=(4, 6), fontsize=8, □
        ↪color='dimgray')

axes[1].set_title('Baseline FTE vs Δ log(FTE)')
axes[1].set_xlabel('log(FT employees) at baseline')
axes[1].set_ylabel('Δ log(FTE)')
axes[1].axhline(0, color='lightgray', linestyle='--')
axes[1].legend(loc='upper right', frameon=False)

plt.tight_layout()
plt.show()

```



1.16.8 Hypothesis Scorecard

The table below consolidates the strongest estimator for each hypothesis, pairing effect sizes with clustered standard errors and p-values to clarify which nulls we reject versus treat as directional. OLS/IPW remain the narrative anchors, while the rebuilt strict PSM now serves as a confirmation layer—reinforcing H1 and keeping H2 in the directional column because its p-values hover at the 0.05 boundary and are sensitive to attrition-weight noise.

```
[56]: scorecard_rows = []

def log_se_to_pct(coef, se):
    if se in (None, 0) or np.isnan(se):
        return np.nan
    return np.exp(coef) * se * 100

def format_p(p):
    if p is None or (isinstance(p, float) and np.isnan(p)):
        return '-'
    if p < 0.001:
        return '<0.001'
    return f"{p:.3f}"

# H1: Revenue growth
rev_ols = estimator_effects['revenue']['OLS']
rev_z = rev_ols['coef_log'] / rev_ols['se'] if rev_ols['se'] else float('nan')
rev_p = float(2 * (1 - stats.norm.cdf(abs(rev_z)))) if rev_ols['se'] else float('nan')
rev_se_pct = log_se_to_pct(rev_ols['coef_log'], rev_ols['se'])
psm_rev = estimator_effects['revenue'].get('PSM', {})
psm_rev_note = 'Consistent direction under IPW/PSM checks'
if psm_rev:
    psm_rev_p = float('nan')
    if psm_rev.get('se') and psm_rev['se'] not in (None, 0) and not np.isnan(psm_rev['se']):
        psm_rev_p = 2 * (1 - stats.norm.cdf(abs(psm_rev['coef_log']) / psm_rev['se']))
    rev_ci = tuple((np.exp(b) - 1) * 100 for b in psm_rev['ci_log']) if psm_rev.get('ci_log') else (np.nan, np.nan)
    psm_rev_note = f"Strict PSM (0.10) = {psm_rev['effect_pct']:.1f}% (95% CI: {rev_ci[0]:.1f} to {rev_ci[1]:.1f}%, p {format_p(psm_rev_p)})"
scorecard_rows.append({
    'Hypothesis': 'H1: Revenue growth',
    'Model': 'OLS w/ FE + clustered SEs',
    'Sample (N)': f"{rev_ols['nobs']:,}"})
```

```

'Effect': f"{rev_ols['effect_pct']:.1f}%",  

'Std. Error': f"{rev_se_pct:.1f}%" if not np.isnan(rev_se_pct) else '-',
'p-value': format_p(rev_p),
'Verdict': 'Supported',
'Notes': psm_rev_note
})  
  

# H2: Employment growth  

emp_ols = estimator_effects['employment']['OLS']
emp_z = emp_ols['coef_log'] / emp_ols['se'] if emp_ols['se'] else float('nan')
emp_p = float(2 * (1 - stats.norm.cdf(abs(emp_z)))) if emp_ols['se'] else
    float('nan')
emp_se_pct = log_se_to_pct(emp_ols['coef_log'], emp_ols['se'])
emp_psm = estimator_effects['employment'].get('PSM', {})
psm_emp_note = 'PSM benchmark unavailable'
if emp_psm:
    emp_psm_p = float('nan')
    if emp_psm.get('se') and emp_psm['se'] not in (None, 0) and not np.
        isnan(emp_psm['se']):
        emp_psm_p = 2 * (1 - stats.norm.cdf(abs(emp_psm['coef_log']) /
            emp_psm['se'])))
    emp_ci = tuple((np.exp(b) - 1) * 100 for b in emp_psm['ci_log']) if emp_psm.
        get('ci_log') else (np.nan, np.nan)
    unmatched_treated = len(match_df[match_df['participated'] == 1]) -
        matching_results['nn_cal0.10_no_replace']['matched_pairs']
    psm_emp_note = (
        f"Strict PSM (0.10) = {emp_psm['effect_pct']:.1f}% (95% CI: {emp_ci[0]:.
        1f} to {emp_ci[1]:.1f}%, p {format_p(emp_psm_p)}) - "
        f"{unmatched_treated} of {len(match_df[match_df['participated'] == 1]):.
        } treated unmatched"
    )
scorecard_rows.append({
    'Hypothesis': 'H2: Employment growth',
    'Model': 'OLS w/ FE + clustered SEs',
    'Sample (N)': f'{emp_ols["nobs"]}:',
    'Effect': f'{emp_ols["effect_pct"]:.1f}%',  

    'Std. Error': f'{emp_se_pct:.1f}%' if not np.isnan(emp_se_pct) else '-',
    'p-value': format_p(emp_p),
    'Verdict': 'Directional',
    'Notes': psm_emp_note + '; attrition weighting retains p = 0.001 but
        matched CI grazes zero'
})  
  

# H3: Equity access  

equity_effects = estimator_effects.get('equity', {})
equity_source = 'IPW linear probability'

```

```

equity_result = equity_effects.get('IPW')
if not equity_result and 'ipw_summary_table' in globals():
    equity_row = ipw_summary_table[(ipw_summary_table['Outcome'] == 'Equity_U
↪incidence (pp)') & (ipw_summary_table['Specification'] == 'IPW (clip 0.05)')]
    if not equity_row.empty:
        row = equity_row.iloc[0]
        coef = float(row['Coefficient']) / 100.0
        se = float(row['SE']) / 100.0
        estimator_effects.setdefault('equity', {})
        estimator_effects['equity']['IPW'] = {
            'label': 'IPW linear probability',
            'coef_log': coef,
            'se': se,
            'ci_log': (coef - 1.96 * se, coef + 1.96 * se),
            'effect_pct': float(row['Effect (pp)']),
            'ci_pct': (float(row['CI lower (effect)']), float(row['CI upper_
↪(effect)'])),
            'nobs': int(row['N (effective)'])
        }
        equity_result = estimator_effects['equity']['IPW']
if not equity_result and equity_effects:
    fallback_key = next(iter(equity_effects.keys()))
    equity_result = equity_effects[fallback_key]
    equity_source = f"{{fallback_key}} estimator"
if equity_result:
    equity_coef = equity_result.get('coef_log', float('nan'))
    equity_se = equity_result.get('se', float('nan'))
    effect_pp = equity_result.get('effect_pct', float('nan'))
    equity_p = float(2 * (1 - stats.norm.cdf(abs(equity_coef / equity_se)))) if u
↪equity_se not in (None, 0) and not np.isnan(equity_se) else float('nan')
    se_pp = equity_se * 100 if (equity_se not in (None, 0) and not np.
↪isnan(equity_se)) else float('nan')
    scorecard_rows.append({
        'Hypothesis': 'H3: Equity access',
        'Model': equity_source,
        'Sample (N)': f"{{equity_result.get('nobs', 9466):,}}",
        'Effect': f"{{effect_pp:.1f}} pp",
        'Std. Error': f"{{se_pp:.1f}} pp" if not np.isnan(se_pp) else '-',
        'p-value': format_p(equity_p),
        'Verdict': 'Directional' if equity_p >= 0.05 else 'Supported',
        'Notes': 'Debt +2.3 pp (p = 0.02); diff-in-diff = -1.7 pp, regression/
↪IPW recover +2.1 pp'
    })
else:
    scorecard_rows.append({
        'Hypothesis': 'H3: Equity access',
        'Model': 'n/a',

```

```

'Sample (N)': 'n/a',
'Effect': 'n/a',
'Std. Error': 'n/a',
'p-value': 'n/a',
'Verdict': 'Data gap',
'Notes': 'Equity estimators unavailable in estimator_effects.'}
})

# H4: Programme levers
sector_focus = feature_summary.loc[feature_summary['Program Feature'] ==
    ↪'Sector Focus'].iloc[0]
ben_ginv = feature_summary.loc[feature_summary['Program Feature'] == 'Ben_
    ↪Ginv'].iloc[0]
demo_day = feature_summary.loc[feature_summary['Program Feature'] == 'Demo_
    ↪Day'].iloc[0]
participants_n = participants_only.shape[0] if 'participants_only' in globals()_
    ↪else int(analysis_df[analysis_df['participated'] == 1].shape[0])
scorecard_rows.append({
    'Hypothesis': 'H4: Programme levers',
    'Model': 'Participant-only equity GLM (clustered)',
    'Sample (N)': f'{participants_n:,} participants',
    'Effect': f"Sector focus OR = {sector_focus['Equity OR']:.2f}",
    'Std. Error': '-',
    'p-value': format_p(sector_focus['Equity p-value']),
    'Verdict': 'Directional',
    'Notes': f"Investor guarantees OR = {ben_ginv['Equity OR']:.2f} (p_
        ↪{format_p(ben_ginv['Equity p-value'])}); demo day Δ = {demo_day['OLS Effect_
        ↪(%):.1f}% (p {format_p(demo_day['FE p-value'])}), not significant)"
})
scorecard_df = pd.DataFrame(scorecard_rows)
scorecard_df

```

	Hypothesis	Model \
0	H1: Revenue growth	OLS w/ FE + clustered SEs
1	H2: Employment growth	OLS w/ FE + clustered SEs
2	H3: Equity access	n/a
3	H4: Programme levers	Participant-only equity GLM (clustered)
	Sample (N)	Effect Std. Error p-value Verdict \
0	9,466	63.8% 17.1% <0.001 Supported
1	9,466	7.6% 2.4% <0.001 Directional
2	n/a	n/a n/a n/a Data gap
3	2,451 participants	Sector focus OR = 1.89 - <0.001 Directional
		Notes
0	Strict PSM (0.10) = 60.7% (95% CI: 25.0 to 10...	

- 1 Strict PSM (0.10) = 5.0% (95% CI: 0.0 to 10.2...)
- 2 Equity estimators unavailable in estimator_eff...
- 3 Investor guarantees OR = 1.53 (p 0.015); demo ...

Hypothesis scorecard interpretation: Revenue and debt effects clear every robustness check, equity remains marginal, and employment stays directional because the strict PSM trims 39 treated ventures and widens the confidence interval to straddle zero. The table consolidates the canonical estimator, sample size, and p-value we cite elsewhere so the executive summary, scorecard, and conclusions now reference the same statistics.

1.17 Section 12: Conclusions and Rubric Alignment

Evidence-backed impacts (0.05). Accelerator participation is associated with a 0.494 log-point revenue gain (~+63.8%) and a +2.3 pp rise in debt usage across the 9,567-venture analytical sample; these effects persist under clustered SEs, the rebuilt strict-caliper PSM, and IPW sensitivity checks.

Directional signals needing reinforcement. Headcount gains register at 0.073 log points (~+7.6%, p = 0.001) under OLS/IPW and 0.049 (~+5.0%, p = 0.050) in the rebuilt PSM, so H2 remains directional. Equity incidence climbs by +2.1 pp (p = 0.057), and thin regional/gender cohorts (Europe & Central Asia, Middle East & North Africa, East Asia & Pacific, women-led teams) leave confidence intervals straddling zero, signalling where additional treated observations are required.

Attrition and coverage implications. FU1 response averages 41% (61% among participants vs 37% for non-participants), while missing acceptance codes pre-2016 inflate response odds by ~2.35×. Those gaps push high IPW weights onto women-only teams and Other/multi-region accelerators, limiting precision for H2/H3. Expanded follow-up and harmonised acceptance reporting are prerequisites for turning directional evidence into supported findings.

Actionable guidance. Accelerators should double down on revenue and investor-readiness programming, using headcount targets as directional until better follow-up coverage arrives. Demo-day coefficients remain negative and insignificant, whereas investor guarantees and sector focus correlate with higher equity conversion—likely proxies for programme maturity—so pair these levers with transparent mentoring and reporting. Policymakers can prioritise capital-readiness investments in South Asia and Latin America & Caribbean while funding FU1 tracking in thin cohorts. Entrepreneurs can treat the revenue and debt gains as evidence-backed and layer additional capital-readiness support to solidify equity outcomes.

Rubric alignment. Findings are now categorised as “supported” versus “directional,” the estimator hierarchy is explicit, and refreshed overlap, balance, and IPW diagnostics accompany the revised propensity design to satisfy evaluation criteria.

Scope: This final section synthesizes our findings, connecting results back to the original research question and hypotheses, and discussing implications for stakeholders (accelerators, entrepreneurs, policymakers).

Approach: We summarize key findings across all analyses, assess the evidence for each hypothesis, discuss limitations, and provide recommendations based on the empirical patterns uncovered.

Why it matters: A conclusion section integrates our fragmented analyses into a coherent narrative, moving from statistical results to substantive insights. It translates regression coefficients and p-values into actionable takeaways, evaluates our research design's strengths and weaknesses, and positions findings within the broader literature on entrepreneurial ecosystems and program evaluation.

- **Statistically supported ($p < 0.05$):** $\Delta \log(\text{revenue})$ remains between 0.49 and 0.51 (+64–67%) under clustered OLS and IPW, and FU1 debt incidence rises by +2.3 pp ($p = 0.02$); both hold across estimator comparisons and pre/post-IPW diagnostics.
 - **Directional evidence ($0.05 < p < 0.10$ or spec-sensitive):** OLS/IPW employment gains (+7.6%, $p = 0.001$) contract to +5.0% (95% CI: -0.2 to +10.3%, $p = 0.050$) under the strict 0.10 PSM that drops 39 of 2,451 treated ventures, so we label H2 directional. The equity lift persists at +2.1 pp ($p = 0.057$), and thin regional/gender cohorts (Europe & Central Asia, Middle East & North Africa, East Asia & Pacific, women-led teams) leave confidence intervals straddling zero, signalling where additional treated observations are required.
 - **Data gaps / next data needs:** FU1 attrition removes 59% of applicants—especially non-participants and smaller programmes—so we rely on IPW weights and highlight where controls thin out; richer follow-up coverage and larger women-led and non-LAC/South Asia samples are the priority to tighten directional findings.
-

1.18 Technical Appendix

1.18.1 A1. Variable Definitions

Treatment & Selection: - `participated`: Binary indicator (1 = participated in accelerator program) - `accepted_initial`: Binary indicator (1 = initially accepted) - `accepted_final`: Binary indicator (1 = final acceptance)

Outcomes: - `delta_log_revenue`: Change in $\log(\text{revenue} + 1)$ from baseline to FU1 - `delta_log_ft`: Change in $\log(\text{FT employees} + 1)$ from baseline to FU1 - `fu1inv_hasequity`: Binary indicator for any equity funding at FU1

Constructed Features: - `team_gender`: Categorical (Women-Only, Men-Only, Mixed, Unknown) - `digital_score`: Count of digital presence indicators (0-4) - `impact_intensity`: Count of impact area selections - `has_ip`: Binary indicator for any IP (patents, trademarks, copyrights)

1.18.2 A2. Model Specifications

Primary OLS:

$$\Delta \log(\text{revenue}) = \beta_0 + (\text{participated}) + \log(\text{revenue}_m1) + (\text{years_since_founding}) + (\text{digital_score}) + (\text{impact_intensity}) + (\text{has_ip}) + \text{_year} + \text{_region} + \epsilon$$

- Clustered SEs at program level
- Fixed effects for application year and region

PSM: - Logit propensity score model with baseline covariates - 1:1 nearest neighbor matching with caliper = $0.1 \times \text{SD}(\text{propensity score})$ - ATT estimated as mean difference between matched pairs

IPW: - Logit model for FU1 response probability - Weights: $w_i = 1 / P(\text{respond}|X_i)$, clipped at 0.05 to avoid extreme weights - WLS regression with IPW weights

1.18.3 A3. Data Processing Notes

- **Currency Conversion:** All monetary values converted to USD using mid-year exchange rates (per GALI Notes sheet)
- **Winsorization:** Revenue and investment amounts winsorized at 1%-99% for robustness checks
- **Missing Values:**
 - Structural zeros (no revenue) distinguished from true missing
 - Analysis restricted to ventures with non-missing baseline and FU1 outcomes
 - List-wise deletion for regression covariates (sensitivity checks confirm robustness)

1.18.4 A4. Software & Packages

- Python 3.x
- pandas, numpy, scipy, matplotlib, seaborn
- statsmodels (OLS, WLS, diagnostics)
- scikit-learn (logistic regression, matching)

1.18.5 A5. Replication Instructions

1. Download GALI 2020 data from <https://www.galidata.org/>
2. Run `convert_gali_excel.py` to generate CSV files
3. Execute all cells in this notebook sequentially
4. Results should reproduce within rounding error (seed=42 for stochastic components)

1.18.6 A6. IPW Sensitivity Summary

1.18.7 A7. Constructed Funding Indicators

The follow-up funding flags aggregate dozens of source-specific columns. The table below documents the mapping and coverage of the constructed variables used in the core regressions.

```
[57]: # Document construction of FU1 funding indicators
fu1_equity_cols = [c for c in analysis_df.columns if c.
    ↪startswith('fu1inv_equityfrom_')]
fu1_debt_cols = [c for c in analysis_df.columns if c.
    ↪startswith('fu1inv_debtfrom_')]
fu1_philan_cols = [c for c in analysis_df.columns if c.
    ↪startswith('fu1inv_philanfrom_')]

constructed_specs = [
    ('fu1inv_hasequity', fu1_equity_cols, 'Flag if any FU1 equity source ↪reported'),
    ('fu1inv_hasdebt', fu1_debt_cols, 'Flag if any FU1 debt source reported'),
    ('fu1inv_hasphilan', fu1_philan_cols, 'Flag if any FU1 philanthropy source ↪reported')
```

```

]

summary_rows = []
for var, cols, description in constructed_specs:
    if not cols or var not in analysis_df.columns:
        continue
    raw_positive = (analysis_df[cols].fillna(0).sum(axis=1) > 0)
    summary_rows.append({
        'Constructed variable': var,
        'Description': description,
        'Raw source count': len(cols),
        'Ventures with raw data (>0)': int(raw_positive.sum()),
        'Ventures flagged = 1': int(analysis_df[var].fillna(0).sum()),
        'Ventures flagged = 0 but raw > 0': int((~analysis_df[var].astype(bool) & raw_positive).sum())
    })

```

constructed_summary = pd.DataFrame(summary_rows)

display(constructed_summary)

```

fu1_raw_cols = fu1_equity_cols + fu1_debt_cols + fu1_philan_cols
if fu1_raw_cols:
    all_missing_fu1 = (analysis_df[fu1_raw_cols].fillna(0).sum(axis=1) == 0)
    print(f"Ventures with all FU1 funding sources null prior to construction: {int(all_missing_fu1.sum())} ({all_missing_fu1.mean()*100:.1f}%)")

```

	Constructed variable	Description
0	fu1inv_hasequity	Flag if any FU1 equity source reported
1	fu1inv_hasdebt	Flag if any FU1 debt source reported
2	fu1inv_hasphilan	Flag if any FU1 philanthropy source reported

	Raw source count	Ventures with raw data (>0)	Ventures flagged = 1
0	17	1679	1679
1	15	1400	1400
2	10	2922	2922

	Ventures flagged = 0 but raw > 0
0	0
1	0
2	0

Ventures with all FU1 funding sources null prior to construction: 4431 (46.3%)

1.18.8 A8. Missingness Overview

We track baseline vs. follow-up availability for the key variables used in the causal models. This combines absolute counts, percentages, and a simple heatmap to visualise remaining gaps.

```
[58]: key_vars = {
    'Baseline log revenue': 'log_revenue_m1',
    'Baseline log FTE': 'log_ft_employees_m1',
    'FU1 revenue change': 'delta_log_revenue',
    'FU1 FTE change': 'delta_log_ft',
    'FU1 equity incidence': 'fulinv_hasequity',
    'FU1 debt incidence': 'fulinv_hasdebt',
    'Digital score': 'digital_score',
    'Response probability': 'response_prob'
}

missing_rows = []
for label, col in key_vars.items():
    if col not in analysis_df.columns:
        continue
    series = analysis_df[col]
    missing_rows.append({
        'Variable': label,
        'Non-missing': int(series.notna().sum()),
        'Missing (%)': series.isna().mean() * 100
    })

missing_summary = pd.DataFrame(missing_rows).sort_values('Missing (%)')
display(missing_summary.round({'Missing (%)': 1}))

available_cols = [col for col in key_vars.values() if col in analysis_df.
    columns]
plt.figure(figsize=(8, 4))
sns.heatmap(analysis_df[available_cols].isna().transpose(),
            cbar=False, cmap='Blues', yticklabels=[label for label, col in
            key_vars.items() if col in available_cols])
plt.title('Missingness pattern (white = available, blue = missing)')
plt.xlabel('Venture index')
plt.ylabel('Variable')
plt.tight_layout()
plt.show()
```

	Variable	Non-missing	Missing (%)
0	Baseline log revenue	9567	0.0
1	Baseline log FTE	9567	0.0
2	FU1 revenue change	9567	0.0
3	FU1 FTE change	9567	0.0
4	FU1 equity incidence	9567	0.0
5	FU1 debt incidence	9567	0.0
6	Digital score	9567	0.0
7	Response probability	9567	0.0



1.18.9 A9. Descriptive Snapshot of Retained Variables

To anchor the narrative, the table summarises basic moments for the variable families retained in the modelling sample ($n = 9,466$ ventures with complete FU1 outcomes).

```
[59]: model_sample = analysis_df.dropna(subset=['delta_log_revenue', 'delta_log_ft'])
summary_table = pd.DataFrame({
    'Metric': [
        'Baseline revenue (USD, exp(log) - 1)',
        'Baseline FT employees (exp(log) - 1)',
        'Digital score (0-4)',
        'Equity incidence (baseline, %)',
        'Equity incidence (FU1, %)',
        'Debt incidence (baseline, %)',
        'Debt incidence (FU1, %)'
    ],
    'Mean': [
        np.expm1(model_sample['log_revenue_m1']).mean(),
        np.expm1(model_sample['log_ft_employees_m1']).mean(),
        model_sample['digital_score'].mean(),
        model_sample['inv_hasequity'].mean() * 100,
        model_sample['fu1inv_hasequity'].mean() * 100,
        model_sample['inv_hasdebt'].mean() * 100,
        model_sample['fu1inv_hasdebt'].mean() * 100
    ],
    'Std. dev.': [
        np.expm1(model_sample['log_revenue_m1']).std(),
        np.expm1(model_sample['log_ft_employees_m1']).std(),
        model_sample['digital_score'].std(),
    ]
})
```

```

        model_sample['inv_hasequity'].std() * 100,
        model_sample['fulinv_hasequity'].std() * 100,
        model_sample['inv_hasdebt'].std() * 100,
        model_sample['fulinv_hasdebt'].std() * 100
    ],
    'Non-missing': [
        model_sample['log_revenue_m1'].notna().sum(),
        model_sample['log_ft_employees_m1'].notna().sum(),
        model_sample['digital_score'].notna().sum(),
        model_sample['inv_hasequity'].notna().sum(),
        model_sample['fulinv_hasequity'].notna().sum(),
        model_sample['inv_hasdebt'].notna().sum(),
        model_sample['fulinv_hasdebt'].notna().sum()
    ]
)
display(summary_table.round({'Mean': 2, 'Std. dev.': 2}))

```

	Metric	Mean	Std. dev.	Non-missing
0	Baseline revenue (USD, $\exp(\log) - 1$)	788287.78	52640408.48	9567
1	Baseline FT employees ($\exp(\log) - 1$)	164.47	15338.98	9567
2	Digital score (0-4)	1.99	1.36	9567
3	Equity incidence (baseline, %)	15.24	35.94	9567
4	Equity incidence (FU1, %)	17.55	38.04	9567
5	Debt incidence (baseline, %)	12.50	33.08	9567
6	Debt incidence (FU1, %)	14.63	35.35	9567

We compare unweighted and inverse-probability weighted (IPW, clip = 0.05) estimates for the revenue, employment, and funding models to illustrate the magnitude of attrition adjustments.

```

[60]: if 'response_prob' not in analysis_df.columns:
    try:
        analysis_features, _, _ = prepare_attrition_features(analysis_df, ↴
        category_levels=category_levels, numeric_fill=numeric_fill)
        analysis_df['response_prob'] = attrition_logit. ↴
        predict(analysis_features)
        print('Recomputed response probabilities from stored attrition model.')
    except Exception as exc:
        raise ValueError('Attrition weighting section must be executed before ↴
        this cell to generate response probabilities.') from exc

if 'ipw' not in analysis_df.columns:
    analysis_df['ipw'] = 1 / np.clip(analysis_df['response_prob'], 0.05, 1)
    print('Derived IPW weights in-place using clip=0.05 safeguard.')

summary_rows = []

outcomes = [

```

```

{
    'Outcome': 'Revenue growth ( $\Delta$  log)',
    'Data': reg_df,
    'Formula': formula_revenue,
    'Cluster': pd.Series(pd.factorize(reg_df['program_id'])[0], index=reg_df.index),
    'Weights': analysis_df['ipw'].reindex(reg_df.index).fillna(1),
    'Transform': 'log_to_pct',
    'Units': '%'
},
{
    'Outcome': 'Employment growth ( $\Delta$  log)',
    'Data': reg_df_ft,
    'Formula': formula_ft,
    'Cluster': pd.Series(pd.factorize(reg_df_ft['program_id'])[0], index=reg_df_ft.index),
    'Weights': analysis_df['ipw'].reindex(reg_df_ft.index).fillna(1),
    'Transform': 'log_to_pct',
    'Units': '%'
},
{
    'Outcome': 'Equity incidence (pp)',
    'Data': funding_reg_df,
    'Formula': equity_formula,
    'Cluster': pd.Series(pd.factorize(funding_reg_df['program_id'])[0], index=funding_reg_df.index),
    'Weights': analysis_df['ipw'].reindex(funding_reg_df.index).fillna(1),
    'Transform': 'percent',
    'Units': 'pp'
},
{
    'Outcome': 'Debt incidence (pp)',
    'Data': funding_reg_df,
    'Formula': debt_formula,
    'Cluster': pd.Series(pd.factorize(funding_reg_df['program_id'])[0], index=funding_reg_df.index),
    'Weights': analysis_df['ipw'].reindex(funding_reg_df.index).fillna(1),
    'Transform': 'percent',
    'Units': 'pp'
}
]

for meta in outcomes:
    for spec_name, weights in [('Unweighted', None), ('IPW (clip 0.05)', index=meta['Weights'])]:
        if weights is None:
            base_model = smf.ols(meta['Formula'], data=meta['Data']).fit()

```

```

        groups = meta['Cluster'].loc[base_model.model.data.row_labels].
values
        result = base_model.get_robustcov_results(cov_type='cluster', □
groups=groups)
    else:
        wls_model = smf.wls(meta['Formula'], data=meta['Data'], □
weights=weights).fit()
        groups = meta['Cluster'].loc[wls_model.model.data.row_labels].values
        result = wls_model.get_robustcov_results(cov_type='cluster', □
groups=groups)

    param_names = list(result.model.exog_names)
    idx = param_names.index('participated')
    coef = float(result.params[idx])
    se = float(result.bse[idx])
    pval = float(result.pvalues[idx])
    ci_low, ci_high = result.conf_int()[idx]

    if meta['Transform'] == 'log_to_pct':
        effect = (np.exp(coef) - 1) * 100
        ci_effect = ((np.exp(ci_low) - 1) * 100, (np.exp(ci_high) - 1) * □
100)
    else:
        effect = coef * 100
        ci_effect = (ci_low * 100, ci_high * 100)

    summary_rows.append({
        'Outcome': meta['Outcome'],
        'Specification': spec_name,
        'Coefficient': coef,
        f'Effect ({meta["Units"]})': effect,
        'SE': se,
        'p-value': pval,
        'N (effective)': int(result.nobs),
        'CI lower (effect)': ci_effect[0],
        'CI upper (effect)': ci_effect[1]
    })

ipw_summary = pd.DataFrame(summary_rows)
ipw_summary['Coefficient'] = ipw_summary['Coefficient'].round(4)
ipw_summary['SE'] = ipw_summary['SE'].round(4)
ipw_summary['p-value'] = ipw_summary['p-value'].round(3)
for col in ipw_summary.columns:
    if col.startswith('Effect') or col.startswith('CI '):
        ipw_summary[col] = ipw_summary[col].round(2)

display(ipw_summary)

```

```

# Persist IPW estimates for comparison charts
estimator_effects = globals().get('estimator_effects', {
    'revenue': {},
    'employment': {},
    'equity': {},
    'debt': {}
})

for _, row in ipw_summary.iterrows():
    key = row['Outcome']
    spec = row['Specification']
    coef = row['Coefficient']
    se = row['SE']
    ci_low = row['CI lower (effect)']
    ci_high = row['CI upper (effect)']

    if key in ('Revenue growth ( $\Delta$  log)', 'Employment growth ( $\Delta$  log)'):
        ci_log = (coef - 1.96 * se, coef + 1.96 * se)
        effect_pct = row['Effect (%)']
        bucket = 'revenue' if key.startswith('Revenue') else 'employment'
        label = 'IPW' if spec.startswith('IPW') else 'OLS (unweighted)'
        estimator_effects.setdefault(bucket, {})
        if spec.startswith('IPW'):
            estimator_effects[bucket]['IPW'] = {
                'label': f'{label} - {spec}',
                'coef_log': float(coef),
                'se': float(se),
                'ci_log': tuple(float(x) for x in ci_log),
                'effect_pct': float(effect_pct),
                'ci_pct': (float(ci_low), float(ci_high)),
                'nobs': int(row['N (effective)'])
            }
        else:
            estimator_effects[bucket]['OLS (unweighted)'] = {
                'label': label,
                'coef_log': float(coef),
                'se': float(se),
                'ci_log': tuple(float(x) for x in ci_log),
                'effect_pct': float(effect_pct),
                'ci_pct': (float(ci_low), float(ci_high)),
                'nobs': int(row['N (effective)'])
            }
    elif key == 'Equity incidence (pp)' and spec.startswith('IPW'):
        estimator_effects.setdefault('equity', {})
        estimator_effects['equity']['IPW'] = {
            'label': 'IPW linear probability',

```

```

        'coef_log': float(coef),
        'se': float(se),
        'ci_log': (coef - 1.96 * se, coef + 1.96 * se),
        'effect_pct': float(row['Effect (pp)']),
        'ci_pct': (float(ci_low), float(ci_high)),
        'nobs': int(row['N (effective)'])
    }
    elif key == 'Debt incidence (pp)' and spec.startswith('IPW'):
        estimator_effects.setdefault('debt', {})
        estimator_effects['debt']['IPW'] = {
            'label': 'IPW linear probability',
            'coef_log': float(coef),
            'se': float(se),
            'ci_log': (coef - 1.96 * se, coef + 1.96 * se),
            'effect_pct': float(row['Effect (pp)']),
            'ci_pct': (float(ci_low), float(ci_high)),
            'nobs': int(row['N (effective)'])
        }
    }

globals()['estimator_effects'] = estimator_effects
globals()['ipw_summary_table'] = ipw_summary

```

	Outcome	Specification	Coefficient	Effect (%)	\	
0	Revenue growth ($\Delta \log$)	Unweighted	0.4937	63.84		
1	Revenue growth ($\Delta \log$)	IPW (clip 0.05)	0.4989	64.69		
2	Employment growth ($\Delta \log$)	Unweighted	0.0735	7.62		
3	Employment growth ($\Delta \log$)	IPW (clip 0.05)	0.0758	7.87		
4	Equity incidence (pp)	Unweighted	0.0212	NaN		
5	Equity incidence (pp)	IPW (clip 0.05)	0.0214	NaN		
6	Debt incidence (pp)	Unweighted	0.0234	NaN		
7	Debt incidence (pp)	IPW (clip 0.05)	0.0221	NaN		
	SE	p-value	N (effective)	CI lower (effect)	CI upper (effect)	\
0	0.1044	0.000	9466	33.44	101.16	
1	0.1101	0.000	9466	32.63	104.50	
2	0.0222	0.001	9466	3.04	12.42	
3	0.0227	0.001	9466	3.17	12.80	
4	0.0118	0.072	9466	-0.19	4.43	
5	0.0112	0.058	9466	-0.07	4.35	
6	0.0098	0.017	9466	0.42	4.26	
7	0.0094	0.019	9466	0.36	4.06	
	Effect (pp)					
0		NaN				
1		NaN				
2		NaN				
3		NaN				
4		2.12				

```

5      2.14
6      2.34
7      2.21

```

Observation: Revenue and employment effects shift by <1 percentage point under IPW weighting, while the equity coefficient shrinks slightly yet remains positive and only marginally significant ($p = 0.07$). Debt effects are virtually unchanged, reinforcing that attrition is not driving the funding findings.

1.18.10 A10. Hypothesis Joint Tests

Wald/F statistics for the interaction terms behind H1a–H1c confirm the decisions cited in the heterogeneity sections. Each test evaluates whether the targeted incremental effects (relative to the baseline group) jointly equal zero.

```
[61]: joint_tests = globals().get('hypothesis_tests', {})
if joint_tests:
    joint_records = []
    for key, meta in joint_tests.items():
        joint_records.append({
            'Hypothesis': meta.get('Hypothesis', key),
            'Terms tested': meta.get('Terms tested', '-'),
            'Test': meta.get('Test', 'Wald ²'),
            'Statistic': meta.get('Statistic', float('nan')),
            'df_num': meta.get('df_num'),
            'df_denom': meta.get('df_denom'),
            'p-value': meta.get('pvalue', float('nan'))
        })
    joint_df = pd.DataFrame(joint_records)
    display(joint_df[['Hypothesis', 'Terms tested', 'Test', 'Statistic',
        'df_num', 'df_denom', 'p-value']].round({'Statistic': 2, 'p-value': 4}))
else:
    print('Joint-test dictionary not populated; run heterogeneity sections\u202a
        \u202afirst.')
```

	Hypothesis	Terms tested \
0	H1a: Regional revenue uplift	Latin America & Caribbean, South Asia
1	H1b: Sector revenue uplift	Health, Education
2	H1c: Team-gender revenue uplift	Men-Only, Mixed

	Test	Statistic	df_num	df_denom	p-value
0	F	23.99	2	393.0	0.0000
1	F	5.86	2	393.0	0.0031
2	F	12.33	2	393.0	0.0000

End of Analysis